

# ***Tutorial on Software-Defined Networking***

**Maurício Tsugawa**  
**University of Florida**

PRAGMA Cloud Computing and Software-Defined  
Networking (SDN) Technology Workshop

March 20, 2013

# Outline

---

- Introduction to Software-Defined Networking
- Traditional Networking vs. SDN
  - Case study with VLAN
- ViNe IP overlays
  - UF research project
  - How it relates to SDN

# What is Software-Defined Networking?

---

- Broad Definition

- Open Network Foundation: “an architecture that enables direct programmability of networks”
- Internet Engineering Task Force: “an approach that enables applications to converse with and manipulate the control software of network devices and resources” – *Internet Draft, Sep. 2011 by T. Nadeau*

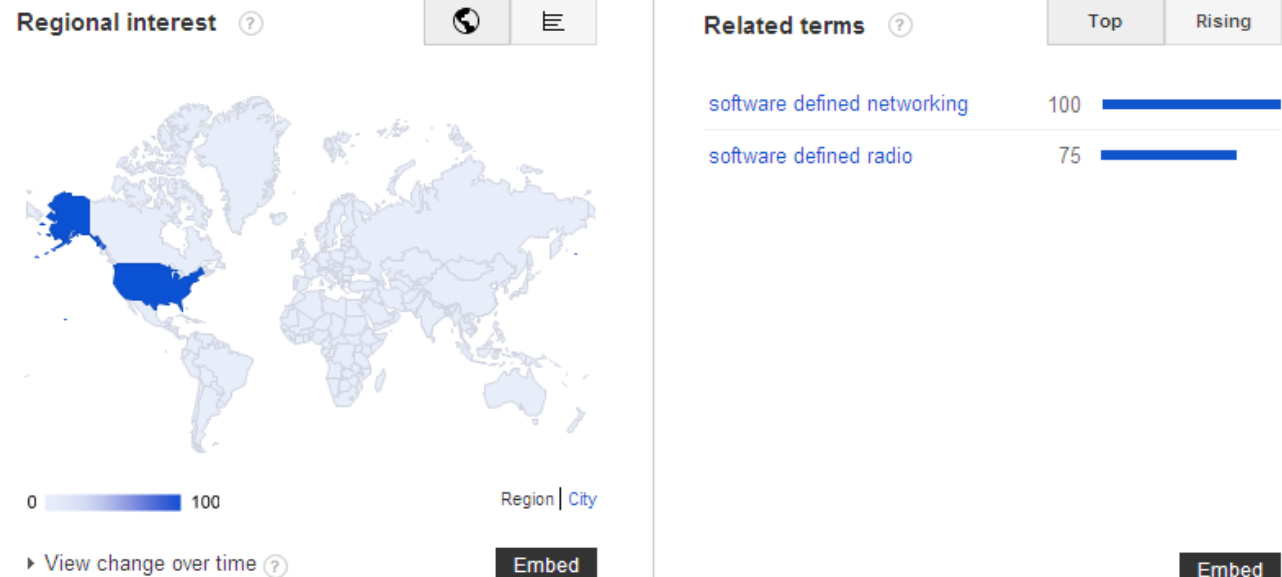

# What SDN is not

---

- OpenFlow
  - An approach to SDN with physical separation between control and data planes
  - Provides open interfaces (APIs)
  - Existing mechanisms can be implemented using OpenFlow, but existing mechanisms cannot implement Openflow
- SDN does not solve all networking problems
  - Offers a clean architecture to apply software engineering to networking
  - Enable buying network control system separately from hardware (switches)

# Software-Defined Networking

- Google Trends
  - Software-Defined Networking
- 2008: OpenFlow interface and NOX
- 2011: Open Networking Foundation



# OpenFlow

- Google Trends
  - OpenFlow

## Interest over time ?


The number 100 represents the peak search volume

☒ News headlines ☐ Forecast ?



Embed

## Regional interest ?



0 100

Region | City

► View change over time ?

Embed

## Related terms ?

Top Rising

openflow switch	100	<div></div>
sdn openflow	95	<div></div>
openflow cisco	75	<div></div>
nox openflow	60	<div></div>
openflow controller	55	<div></div>
openflow tutorial	50	<div></div>
openflow nec	50	<div></div>
openflow networking	45	<div></div>
google openflow	45	<div></div>
open flow	40	<div></div>

Embed

# Need for SDN

---

- Network infrastructure “ossification”
  - Large base of equipments and protocols
  - Networking experiments cannot compete with production traffic
  - No practical way to test new network protocols in realistic settings
- Closed systems
  - Vendor lock-in
  - Proprietary management interfaces – lack of standard or open interfaces
  - Hard to establish collaborations

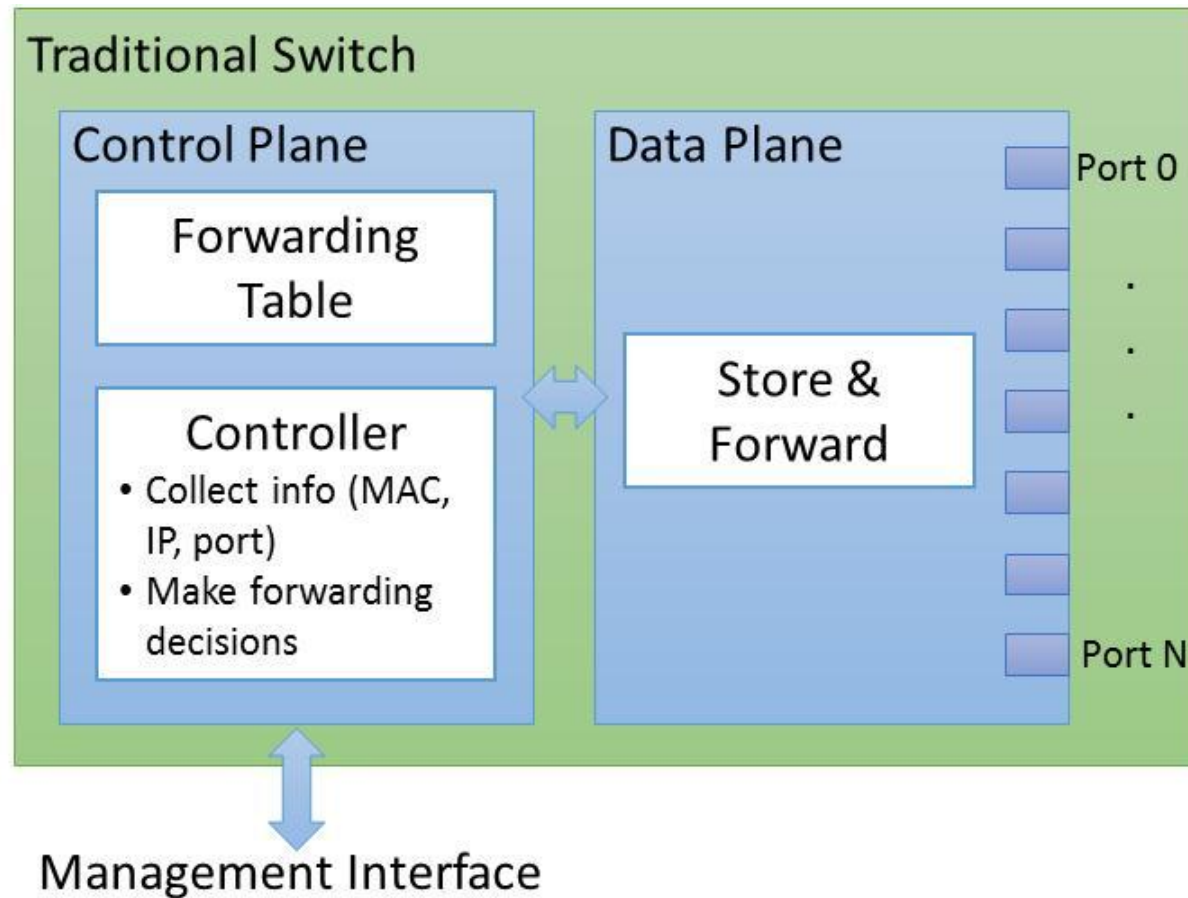
# Complexity in Networking

- Every problem requires an end-to-end solution
  - Involve all routers/switches (configure each one of them)
  - Consider the entire network topology
  - Interface with heterogeneous low level hardware/software
- Too many solutions/mechanisms
  - Routing, VLAN, VPN, MPLS, ACL, firewall
  - RFCs/IEEE standards
  - No software re-use



# Traditional Switch/Router

- Control Plane and Data Plane on the same box





# VLAN/802.1q

- Enables the co-existence of multiple broadcast domains (LANs) on the same infrastructure
- Ethernet frame extended to accommodate VLAN information


Ethernet Frame



802.1q Frame

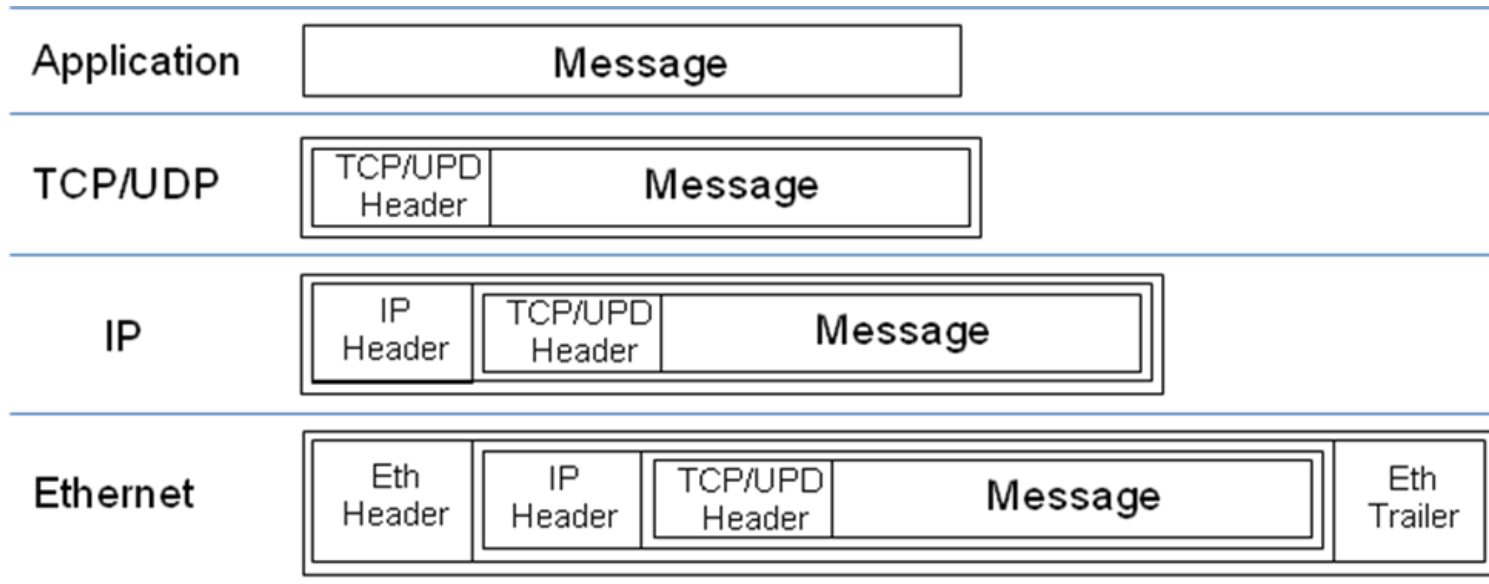


# VLAN/802.1q




# Tunneling/Encapsulation

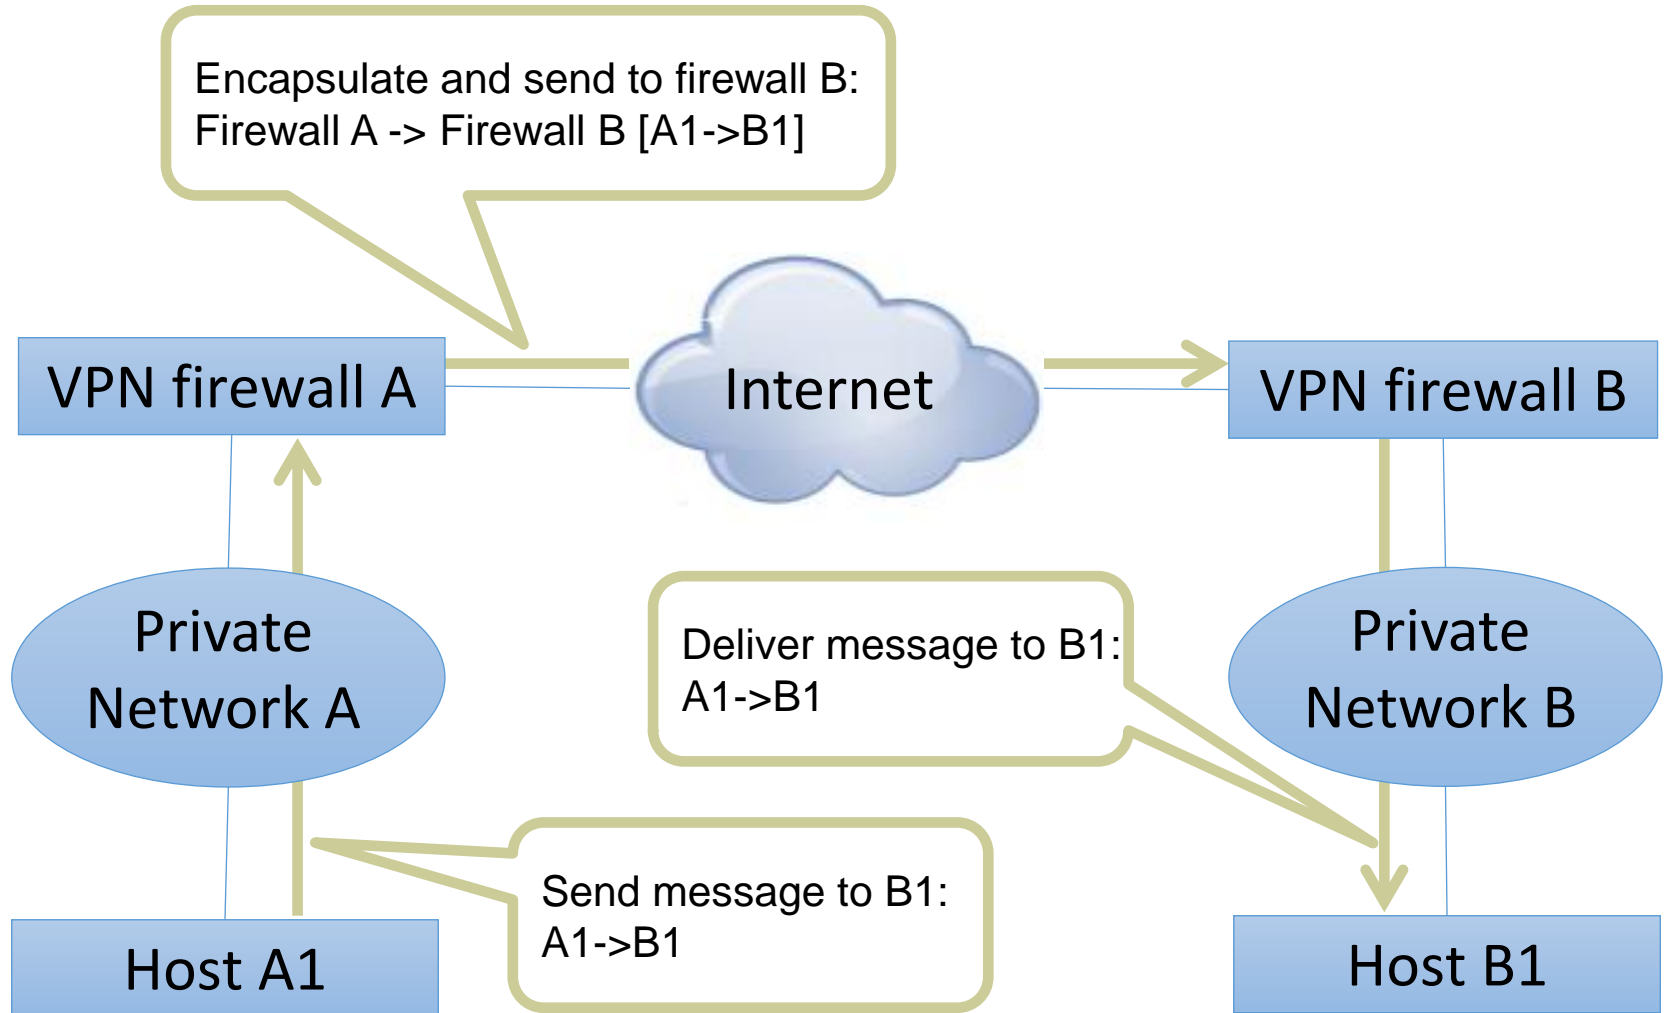
- Headers are added as messages are processed in each layer of network stack



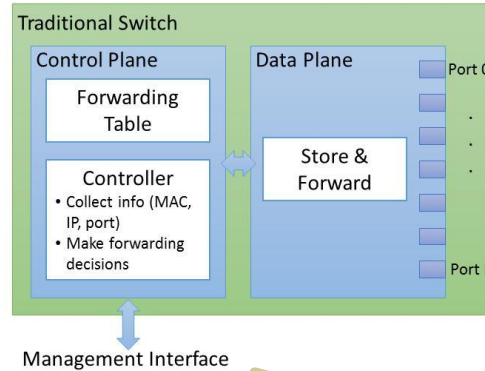
- IP packet encapsulation: additional header is added



# Tunneling



# VLAN configuration



```
# enable
# configure terminal
# vlan 300 Test_VLAN
# interface gigabit 0/1
# switch port mode trunk
# switch port trunk encapsulation dot1q
# switch port trunk native vlan 1
# switch port trunk allowed vlan 1,300
# exit
# interface gigabit 0/2
# switch port access access vlan 5
```

The screenshot shows the Linksys PoE Power Settings web interface. The top navigation bar includes 'Port Management', 'Setup', 'Port Management', 'VLAN Management', 'Statistics', 'ACL', 'Security', 'QoS', 'Spanning Tree', 'Multicast', and 'More >>'. The 'PoE Power Settings' section is active, showing 'Global Setting' and 'Power Allocation(37-180)' with a value of 180 watts. Below this is a 'Port Setting' table with columns for Port, Admin Status, Priority, Power Allocation (3000-15400mWatts), Mode, and Power Consumption (mWatts). The table lists ports g1 through g12, all with 'Enabled' status and 'low' priority. Power allocation is 15400mW for most ports, and 1400mW for g5. Power consumption is 0mW for most ports, and 3600mW for g5. A 'PoE Power Settings' sidebar on the right explains the power allocation process.

LINKSYS  
A Division of Cisco Systems, Inc.

24-port 10/100/1000 + 2-port mini-Gigabit PoE Switch SWV2824P

Port Management

Setup Port Management VLAN Management Statistics ACL Security QoS Spanning Tree Multicast More >>

Port Settings Link Aggregation LACP PoE Power Settings

PoE Power Settings

Global Setting

Power Allocation(37-180) 180 watts

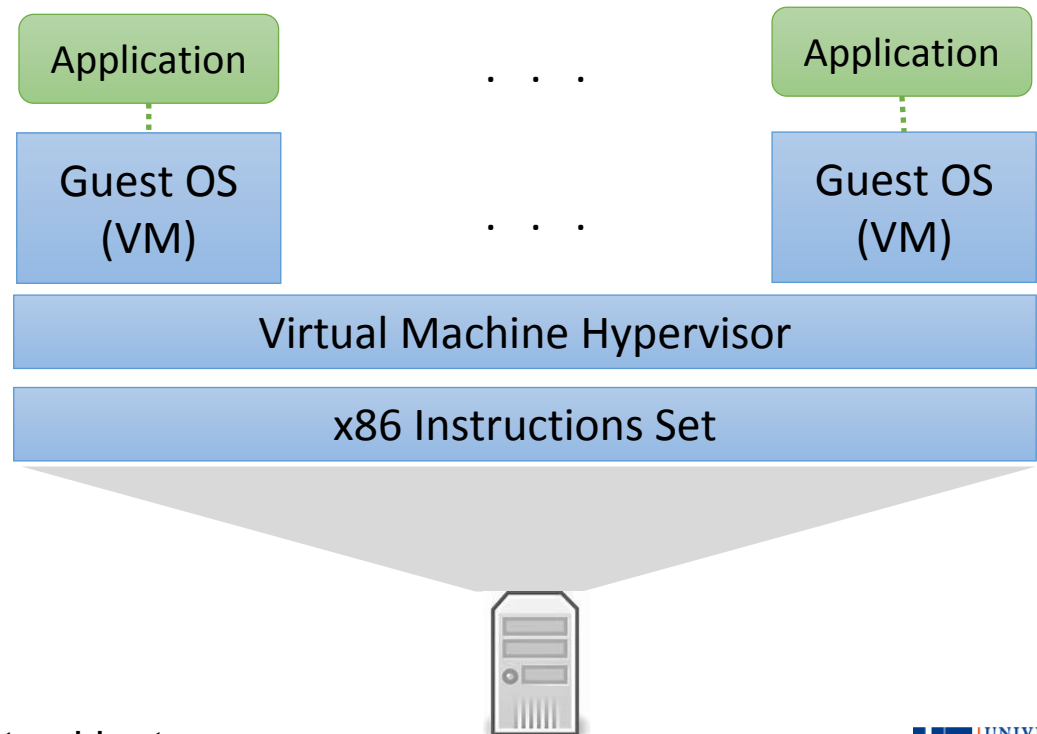
Port Setting

Port	Admin Status	Priority	Power Allocation (3000-15400mWatts)	Mode	Power Consumption (mWatts)
g1	<input checked="" type="checkbox"/> Enabled	low	15400	off	0
g2	<input checked="" type="checkbox"/> Enabled	low	15400	off	0
g3	<input checked="" type="checkbox"/> Enabled	low	15400	off	0
g4	<input checked="" type="checkbox"/> Enabled	low	15400	off	0
g5	<input checked="" type="checkbox"/> Enabled	low	15400	on	1400
g6	<input checked="" type="checkbox"/> Enabled	low	15400	on	3600
g7	<input checked="" type="checkbox"/> Enabled	low	15400	off	0
g8	<input checked="" type="checkbox"/> Enabled	low	15400	off	0
g9	<input checked="" type="checkbox"/> Enabled	low	15400	off	0
g10	<input checked="" type="checkbox"/> Enabled	low	15400	off	0
g11	<input checked="" type="checkbox"/> Enabled	low	15400	off	0
g12	<input checked="" type="checkbox"/> Enabled	low	15400	off	0

PoE Power Settings:  
The switch can provide DC power to a wide range of connected devices based on the IEEE 802.3at (802.3af) Power-over-Ethernet (PoE) standard. Once configured to supply power, an automatic detection process is initialized by the switch that is authenticated by a PoE signature from the connected device. Detection and authentication prevent damage to non-802.3af compliant devices.

# Software Engineering

- Abstractions and interfaces
  - Decompose large problems into smaller (manageable pieces)
  - Separation of concerns
  - Code reuse



# Networking Planes

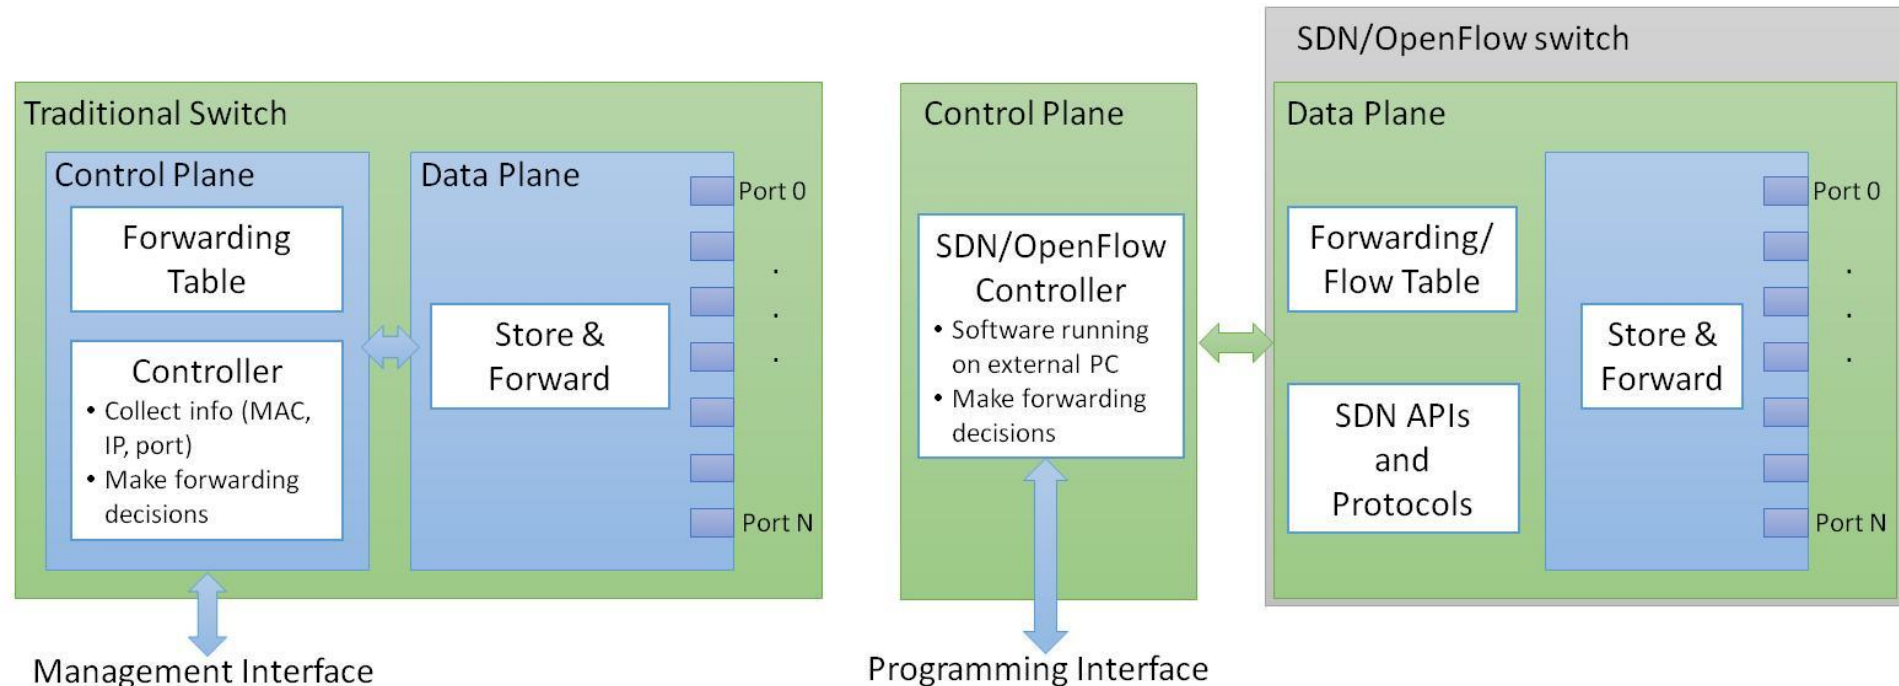
---

- Data Plane
  - Process messages/packets/frames according to local forwarding state
  - Implemented/optimized in hardware
- Control Plane
  - Adjust forwarding state
  - Distributed protocols/algorithms
  - Manual configuration and scripting
- SDN advocates full separation of control and data planes

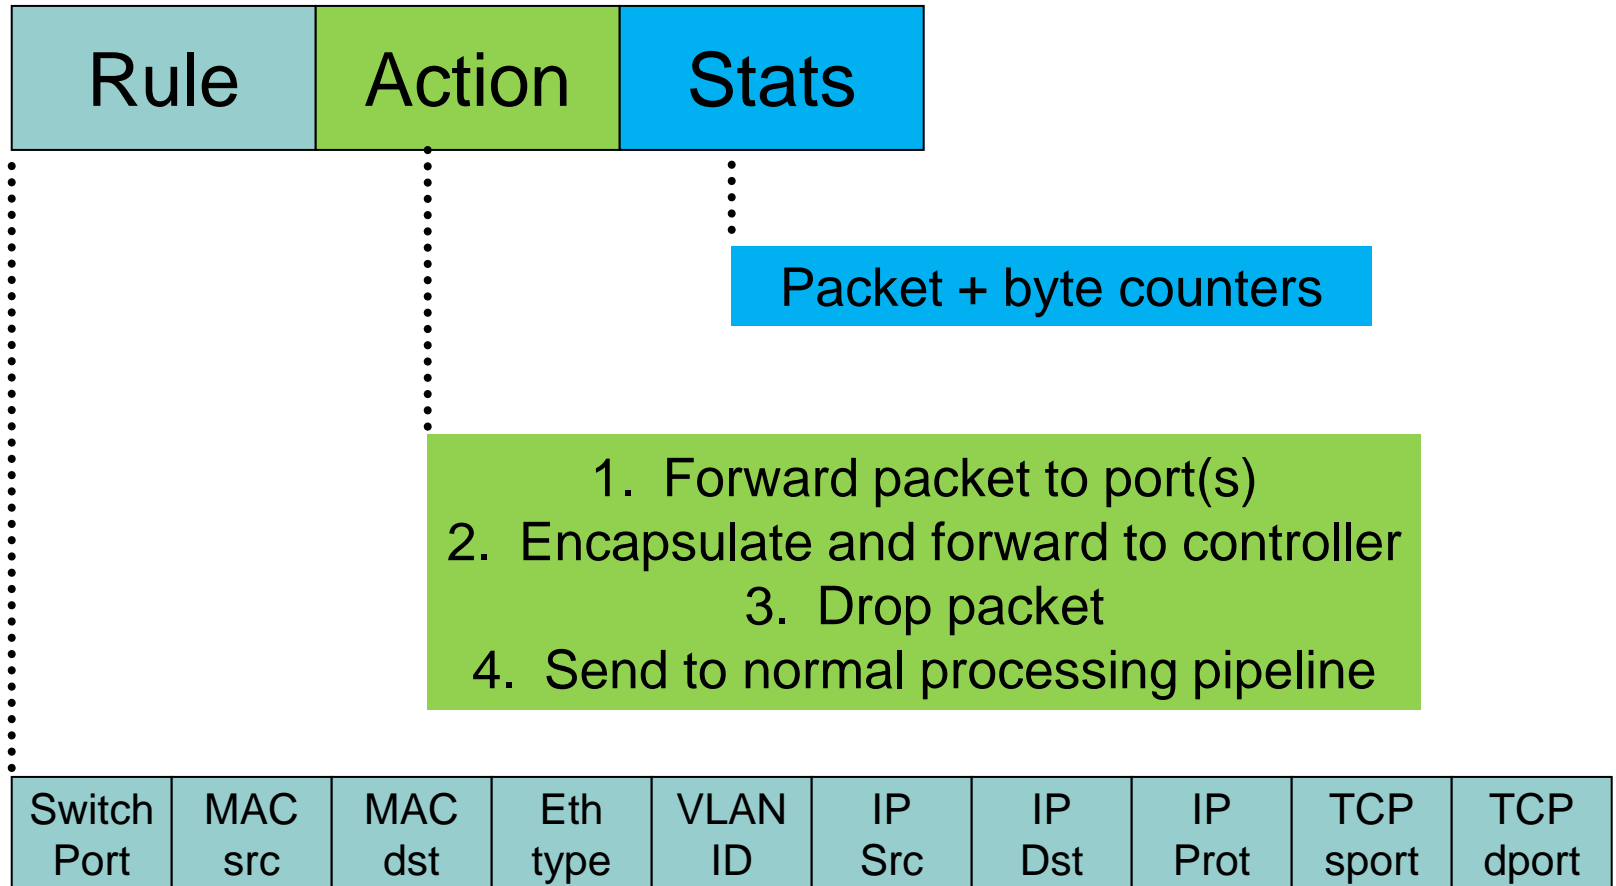


# OpenFlow Architecture

- Separate control plane and data plane
  - Run control plane software on general purpose hardware
  - Programmable data plane

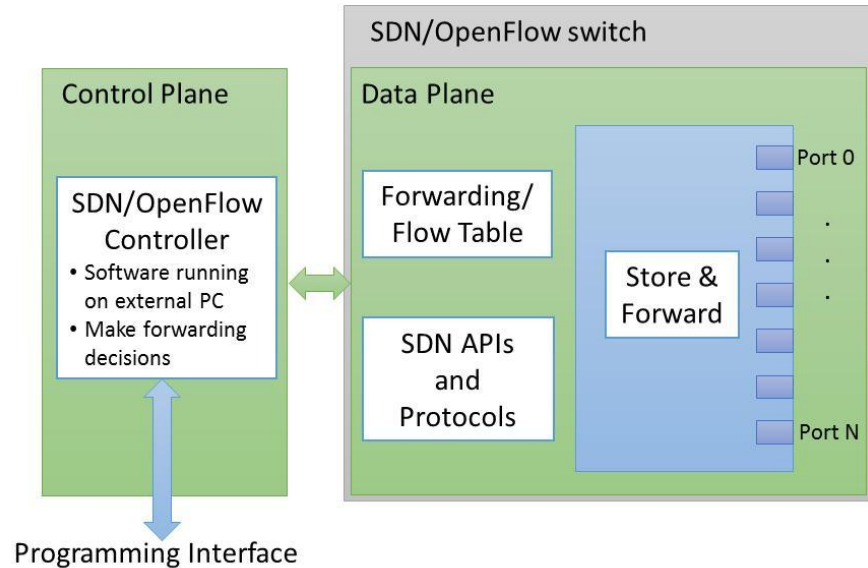


# OpenFlow Flow Table Entry



Source: Nick McKeown, "Why Can't I Innovate in My Wiring Closet?", MIT CSAIL Colloquium, April 2008

# VLAN Configuration with SDN




MAC Addresses of VLAN A

MAC Addresses of VLAN B


If src and dst MAC addresses belong to the same VLAN then deliver packet/frame

# Security-Functionality-Usability

- OpenFlow and current SDN movement started as an academic project
  - Originally intended for campus networks
  - Now, many wish to apply it everywhere (including WAN)
  - Main focus on functionality (programmability of networks)



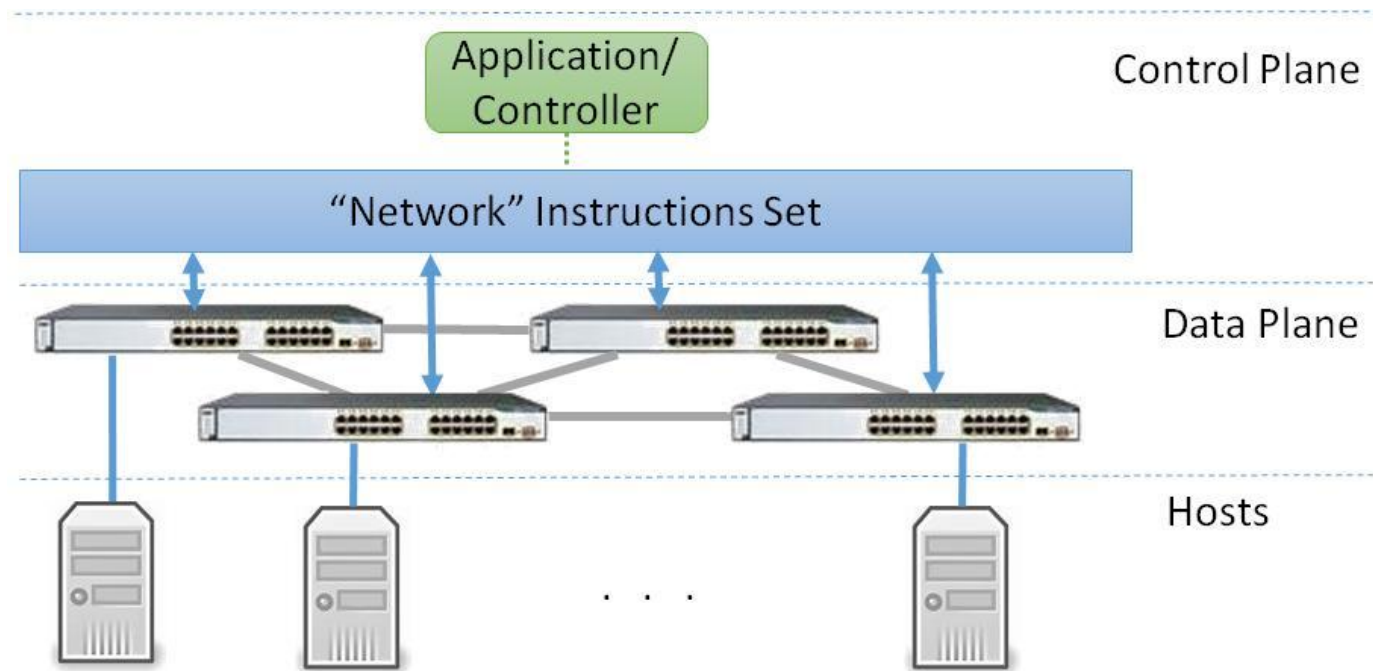
Source: Andrew Waite. InfoSec Triads: Security/Functionality/Ease-of-Use. June 12, 2010.



© Scott Adams, Inc./Dist. by UFS, Inc.

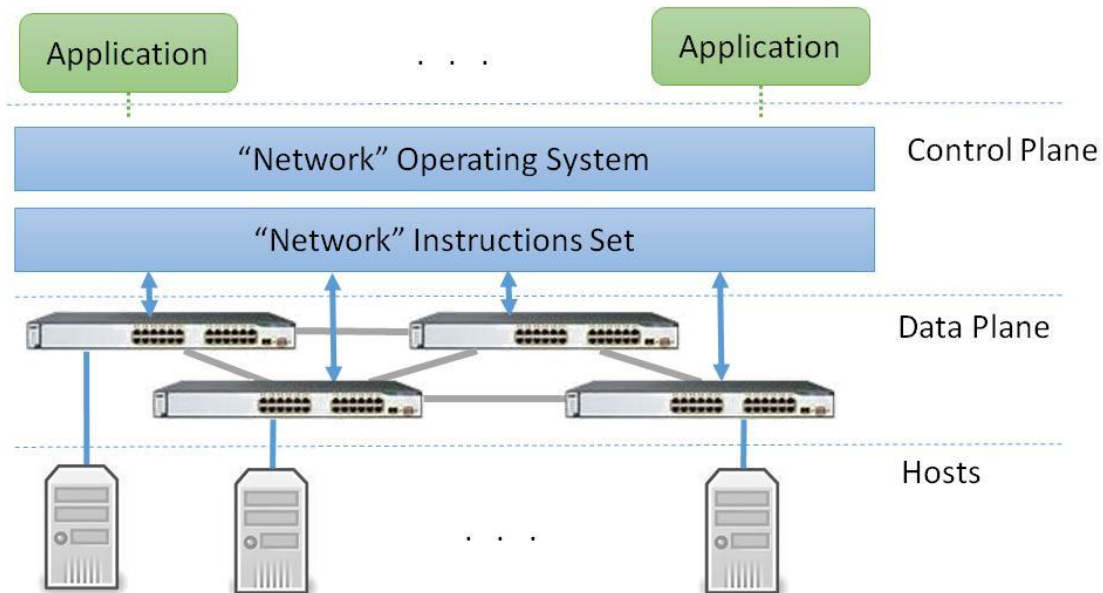
# Stand-alone deployment

- Single controller
  - A user/administrator has full control of the network environment



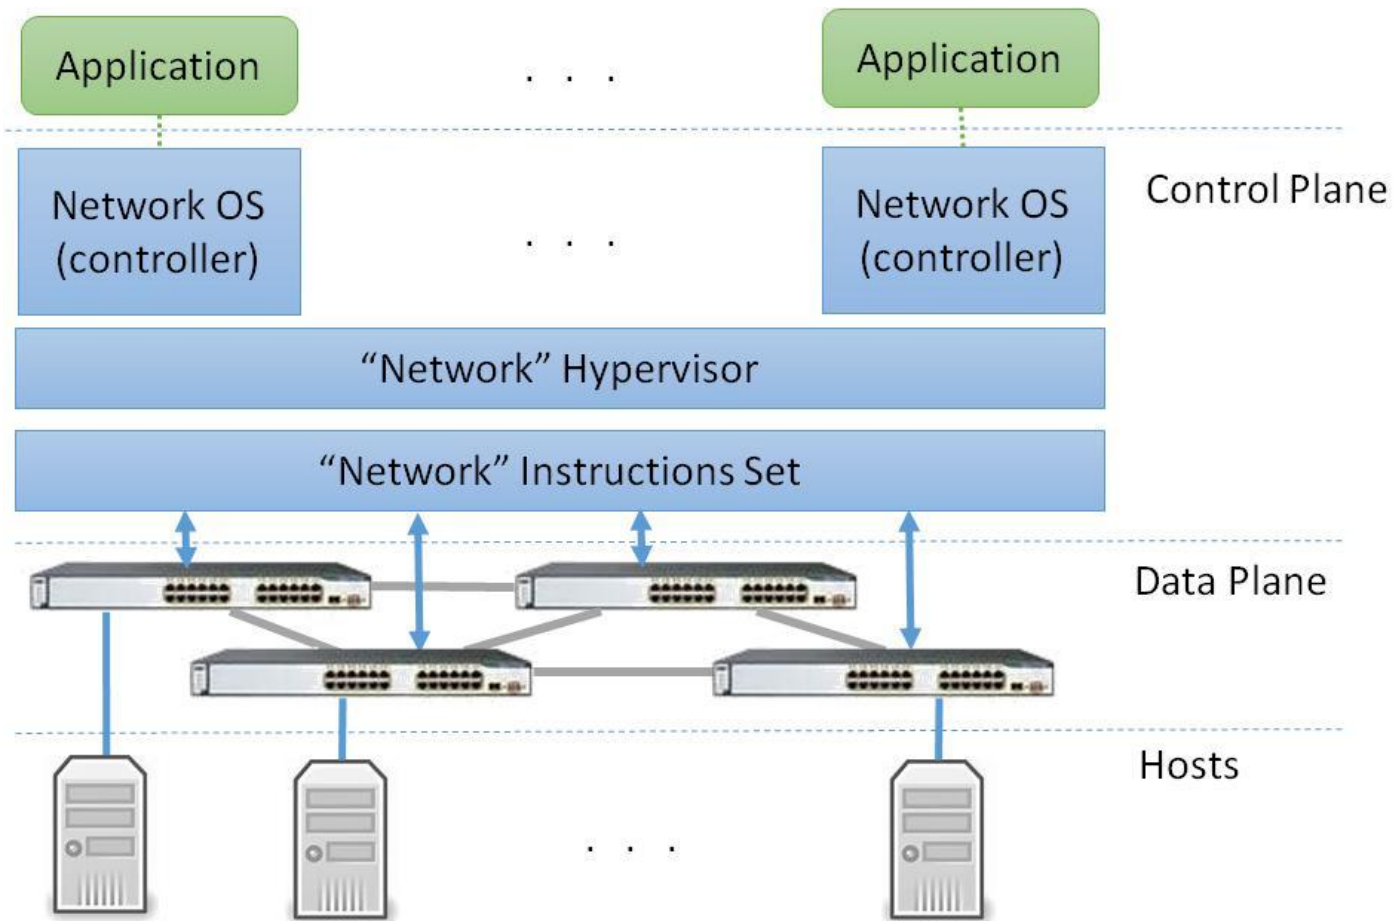
# Multiple applications

- Network Operating System
  - Coordinate data plane resources
    - Each application needs to be identified
    - Needs AAA
  - “Network system call”-like interface needed
    - Accommodate conflicting requests
    - Potential vulnerabilities



# Fully virtualized networking

- SDN offers functionality to implement network virtualization services



# Many Projects/Implementations

---

- Software Switch
  - Open vSwitch
- Network Operating Systems
  - NOX, Trema, FloodLight, Maestro
- Hypervisor
  - FlowVisor
- Routing
  - RouteFlow
- Many others



# Hands-on tutorials

---

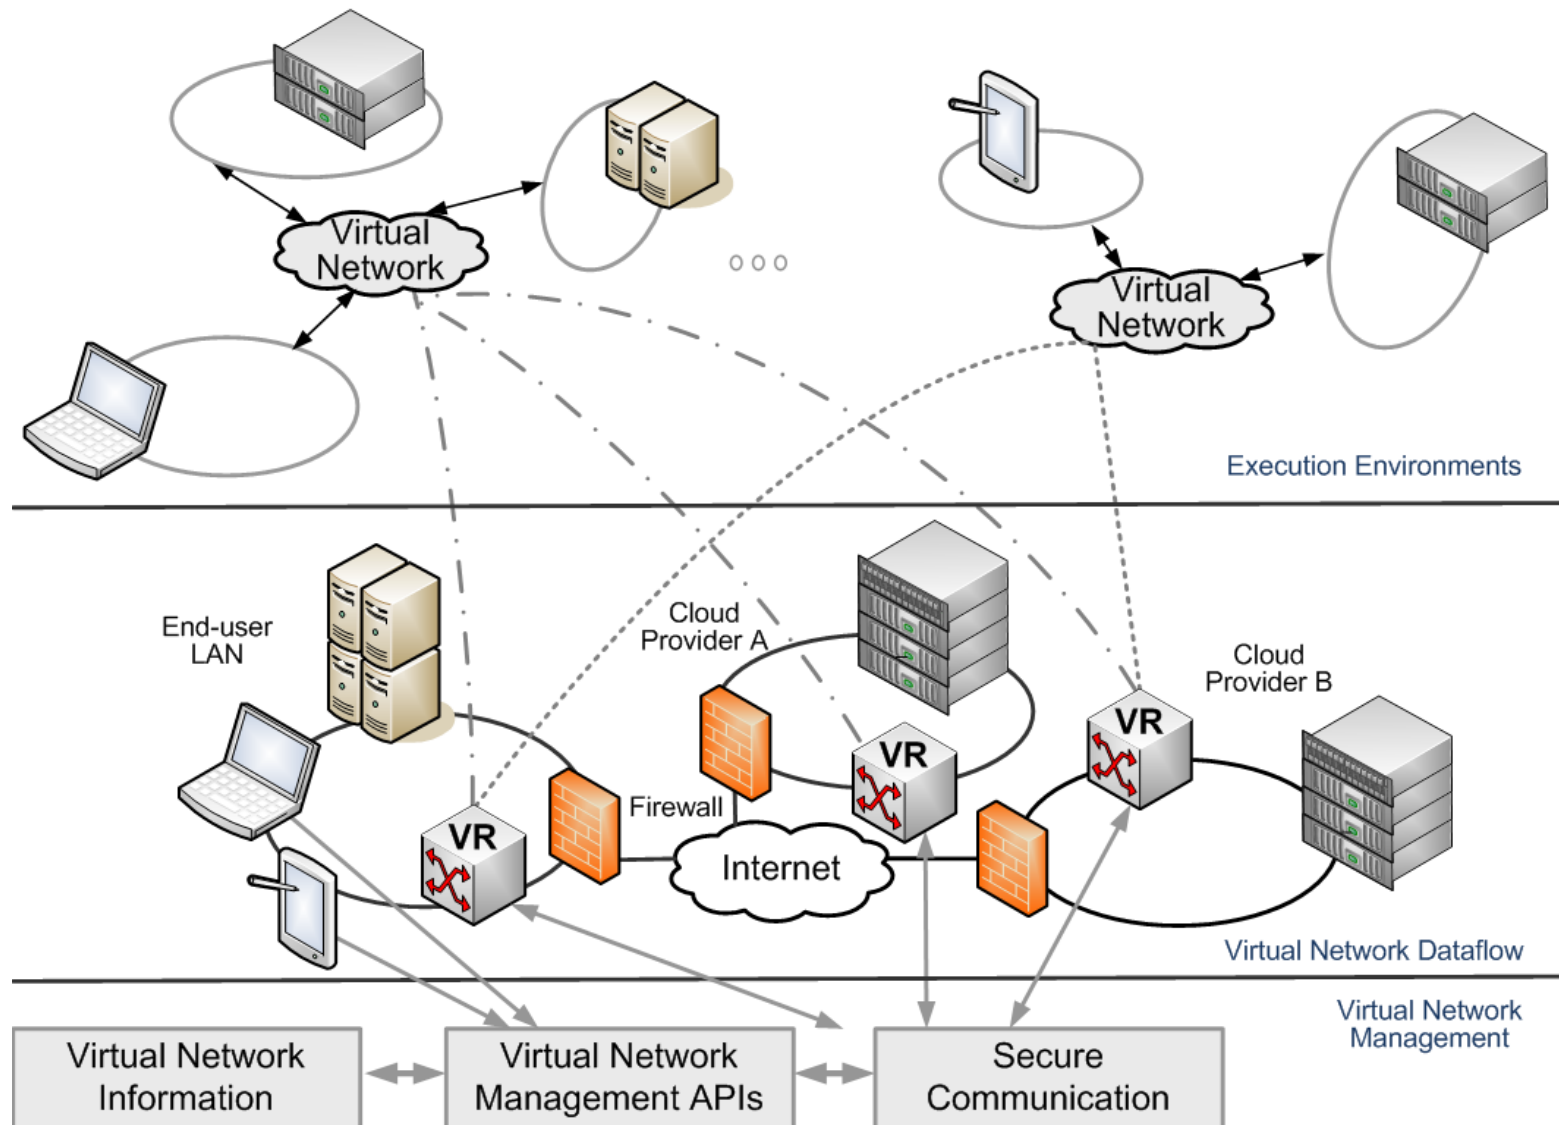
- [http://www.openflow.org/wk/index.php/OpenFlow\\_Tutorial](http://www.openflow.org/wk/index.php/OpenFlow_Tutorial)
  - Use tools pre-packaged in a VM
  - Modify a OpenFlow hub to a learning switch
  - Many controller/platform options
- <http://trema-tutorial.heroku.com/>
  - OpenFlow controller development using tremas

# Quick ViNe Overview

---

- ViNe implements routing and other communication mechanisms needed to deploy user-level virtual networks across WAN
- ViNe offers:
  - Full connectivity among machines (physical and virtual) on public and private networks – built-in firewall traversal
  - Multiple isolated overlays
  - Management APIs

# ViNe Architecture



# ViNe and SDN/OpenFlow

---

- Investigating mechanisms to integrate ViNe technology with SDN/OpenFlow
  - Help incremental deployment
  - Connect edges (servers running OVS) or OpenFlow islands
  - No immediate need for SDN in core networking
- Extend ViNe to talk OpenFlow
  - Implement OpenFlow controller features
  - Accept OpenFlow commands

# Command-Line Interface design

---

- Completely anew
  - Define set of commands
  - Define what each command will do when executed
  - Use compiler techniques to implement CLI
    - Lexical analysis
    - Parsing
    - Regular expression
  - Implement the commands behavior

# CLI design for ViNe

## ● Completely anew

- Define set of commands - want extensible set of commands
- Define what each command will do when executed – link to existing APIs
- Use compiler techniques to implement CLI
  - Lexical analysis
  - Parsing
  - Regular expression
- Implement the commands behavior – most are already implemented

Difficulties due to:

- Undefined set of commands
- Requirement of no recompilation

# Reconfigurable CLI

---

- Map commands to API
  - show Indt → call API that lists ViNe router's local routing table
- Reconfigurable commands set approach
  - Commands definition file describes how commands are mapped to APIs
    - Commands can be changed without recompilation
    - Commands behavior can be changed without recompilation

# Commands definition file

---

- ; Each line is separated by a #
- ; The left part denotes what a user can enter as a command, and the
- ; right part shows what class + method is associated with that command.
- ; Words prefixed by an asterisk (\*) are variables passed to the method.
- ; All variables are assumed to be Strings.
- ; Words prefixed by an dollar (\$) are constant strings passed to the method.
- ; Any lines above a command prefixed by a '#' are descriptions of that command.

```
testCat          # vine.mgt.cli.CmdTest method1
testCat *ip      # vine.mgt.cli.CmdTest method2
test             # vine.mgt.cli.CmdAnotherTest.c printMe
```

- At startup commands definition table are loaded (can be reloaded at run-time)
- Commands are interpreted through table lookup



# ViNe CLI demo

---

