

# **Deployment of Virtual Clusters for Molecular Docking Experiments on the PRAGMA Cloud**

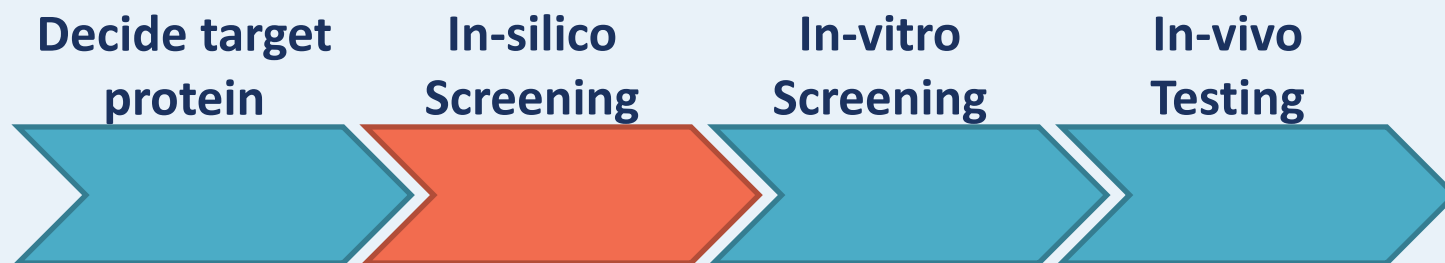
Kohei Ichikawa (NAIST/Osaka U),  
Kevin Lam, (Prime2013)  
Karen Rodriguez, (Prime2013)  
Wen-Wai Yim, (Prime2009)  
Jason Haga (UCSD)

# Objective

- Create virtual clusters for use in molecular docking experiments
  - **Consistent** reproducible, environment for docking studies
  - **Simplify** the learning curve, scripting, and data management
  - **Future potential** clinical treatment for disease

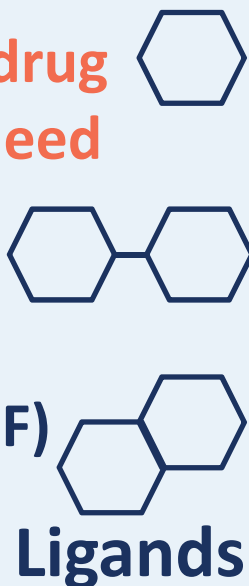
# Background: Docking Simulation

## Drug discovery workflow

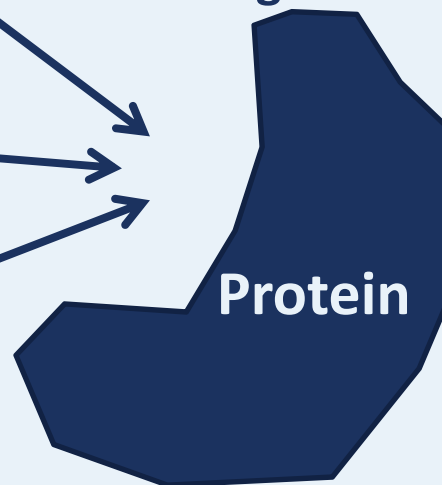


Tens of millions of drug candidate ligands need to be tested

**DOCK** (UCSF)



Check binding

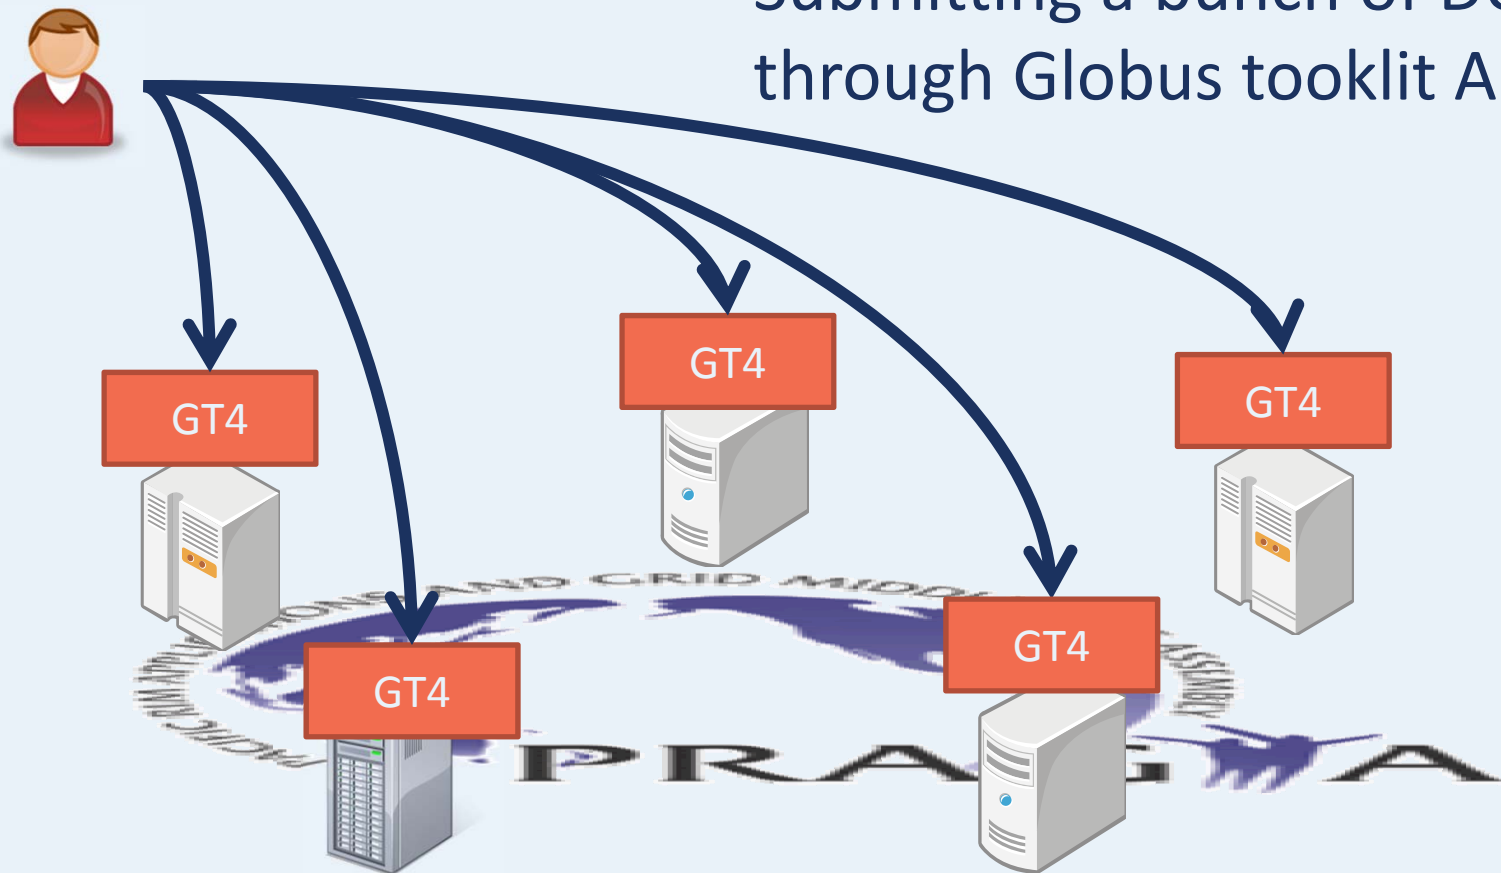


### Score Ranking

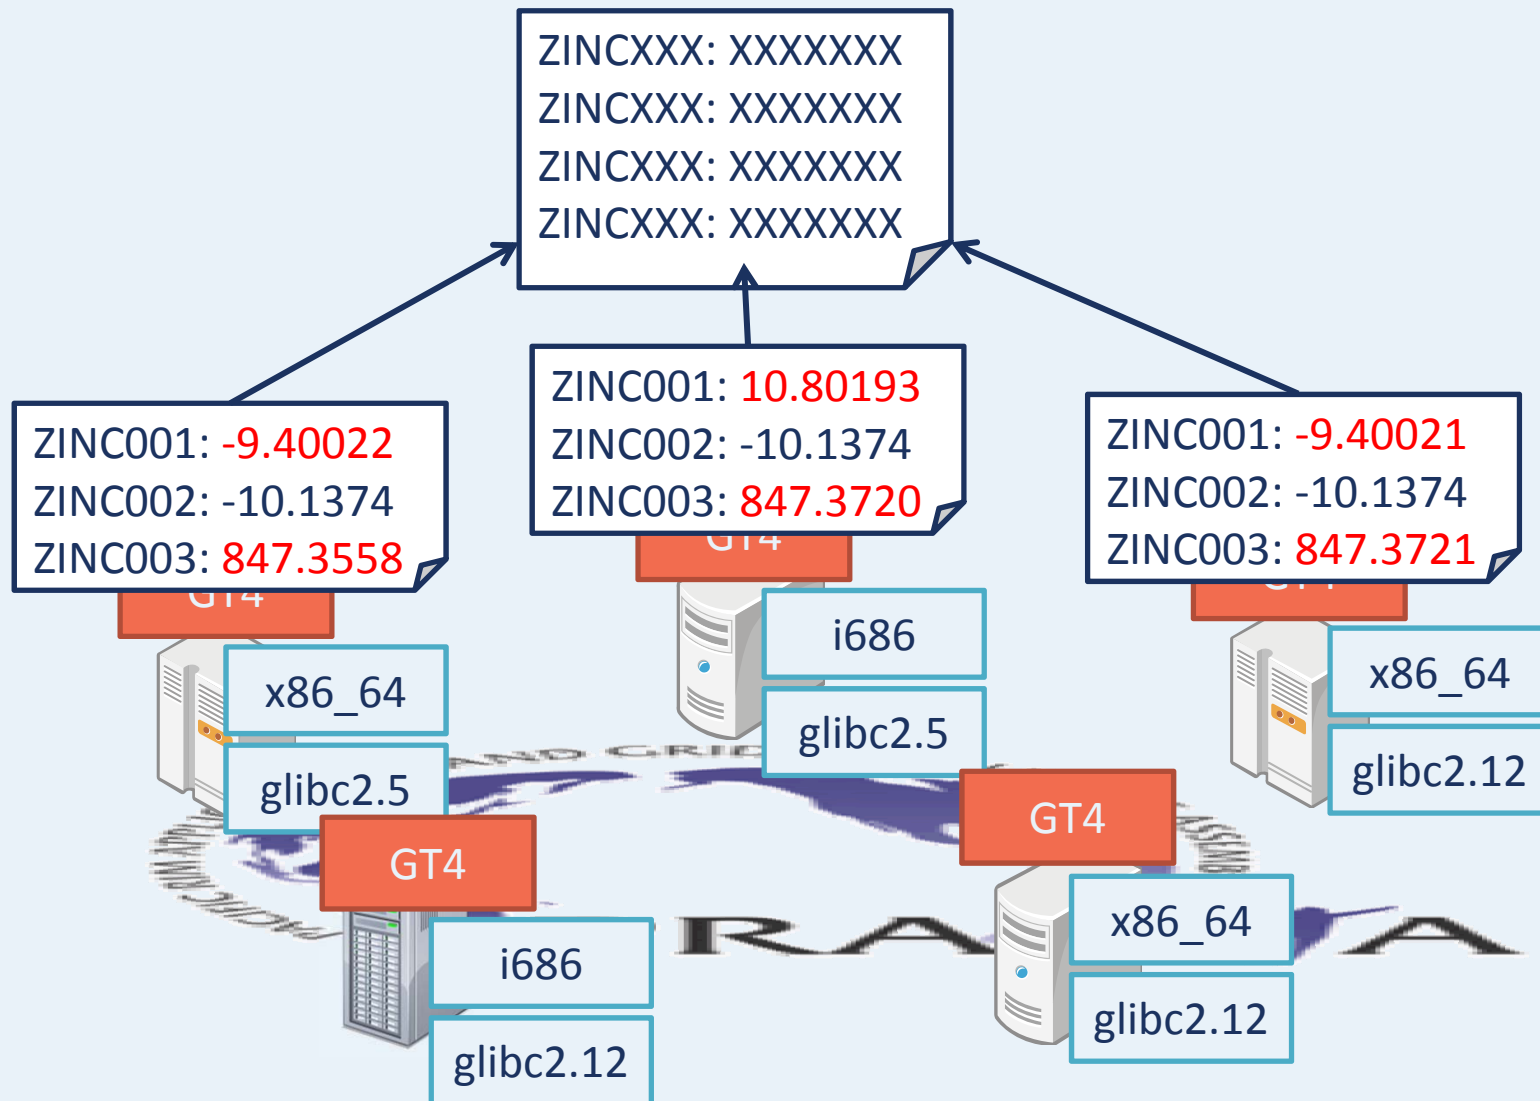
ZINC711: -15.5221  
ZINC715: -10.1374  
ZINC356: -9.40021  
ZINC746: -8.70667  
ZINC875: -1.05811  
ZINC226: 19.05393  
.....

# Grid computing environment

Submitting a bunch of DOCK jobs  
through Globus toolkit API



# Problem: **Heterogeneous Grid** platforms yield inconsistent results

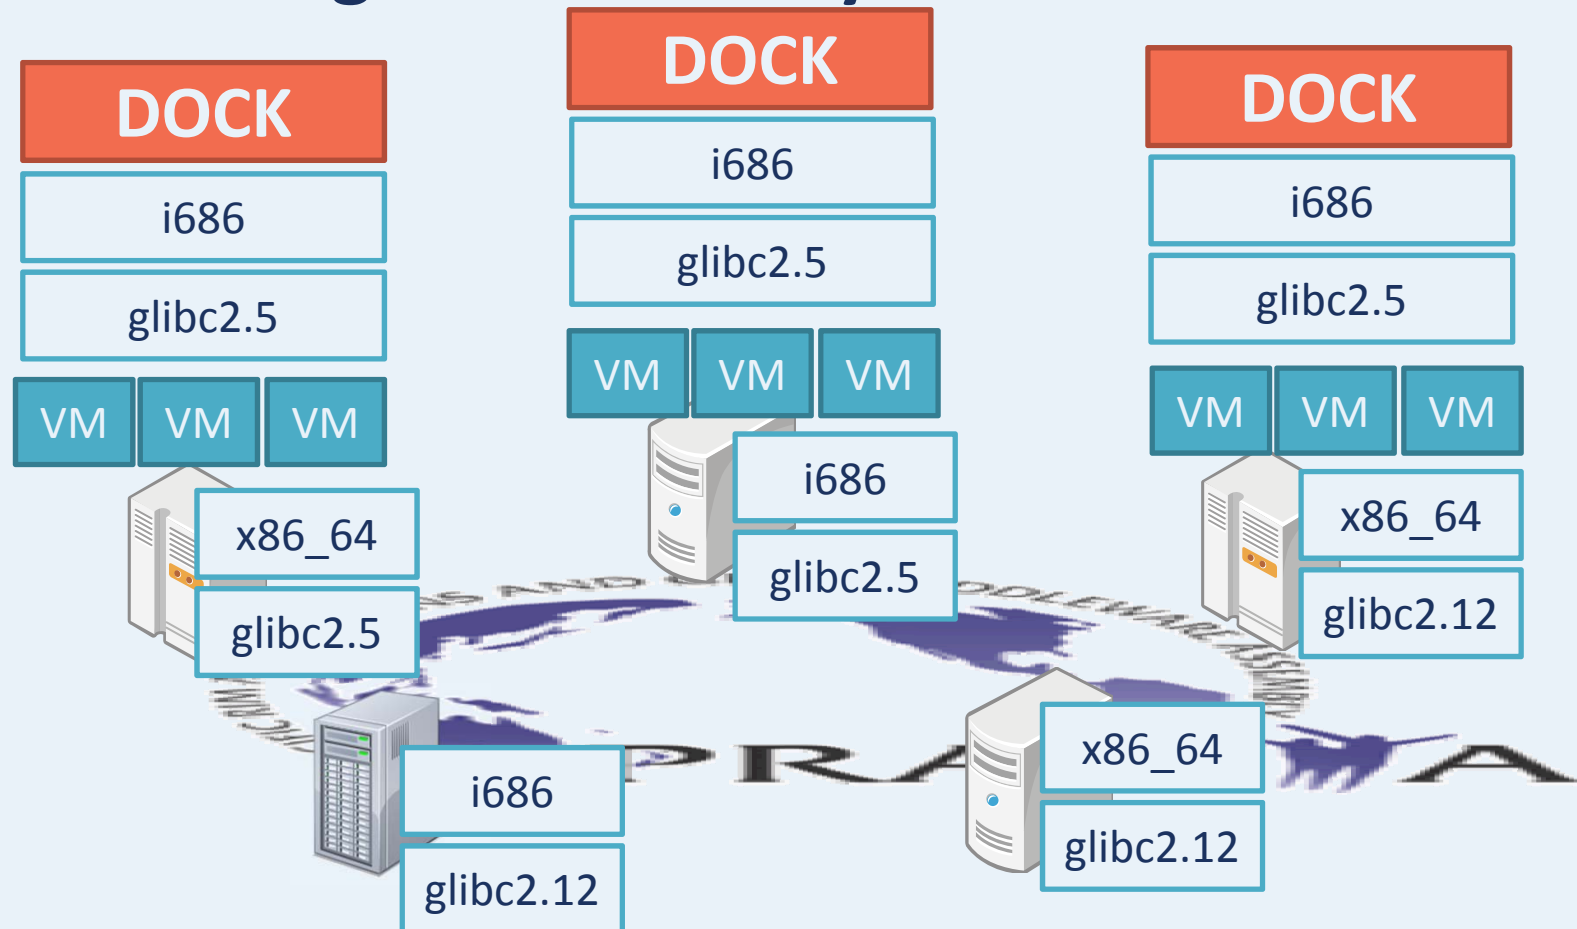


# Score comparison between UCSF DOCK Developer Results and Grid results

	<i>manufacturer's</i>	<i>nacona</i>	<i>komolongma</i>	<i>ocikbpra</i>	<i>aurora</i>	<i>cafe01</i>	<i>tea01</i>
<i>OS-bit</i>	64	32	32	32	32	32	32
<i>machine bit</i>	64	32	32	32	32	32	32
<i>gcc version</i>	3.4.6	4.1.2	3.4.6	3.4.6	3.4.6	3.4.6	3.4.6
<u>ID</u>	<u>Scores</u>						
<b>ZINC00013564</b>	<b>10.801939</b>	-9.40022	10.801939	-9.40021	-9.40021	-9.40021	-9.40021
<b>ZINC00150863</b>	<b>21.139143</b>	-12.4672	21.139143	21.13914	21.13914	21.13914	21.13914
<b>ZINC00152265</b>	<b>30.238361</b>	30.23803	30.238361	19.05393	19.05393	19.05393	19.05393
<b>ZINC00157111</b>	<b>-12.916615</b>	-12.9166	-12.916615	-15.5221	-15.5221	-15.5221	-15.5221
<b>ZINC00157152</b>	<b>-10.137384</b>	-10.1374	-10.137384	-10.1374	-10.1374	-10.1374	-10.1374
<b>ZINC00157402</b>	<b>168513.625</b>	168506.4	168513.625	145519.8	145519.8	145519.8	145519.8
<b>ZINC00157467</b>	<b>-8.706671</b>	-8.70678	-8.706671	-8.70667	-8.70667	-8.70667	-8.70667
<b>ZINC00157960</b>	<b>847.37207</b>	847.3558	847.37207	847.3721	847.3721	847.3721	847.3721
<b>ZINC00158442</b>	<b>52.56588</b>	52.56579	52.56588	52.56588	52.56588	52.56588	52.56588
<b>ZINC00158751</b>	<b>-1.058107</b>	-1.05816	-1.058107	-1.05811	-1.05811	-1.05811	-1.05811
<b>ZINC01555236</b>	<b>503.249725</b>	5.02E+08	503.249725				

# Solution: Using **virtual machine** as a layer of abstraction

Wen-wai Yim (Prime2009) has been working on testing the feasibility.



# Determined the best configuration for DOCK

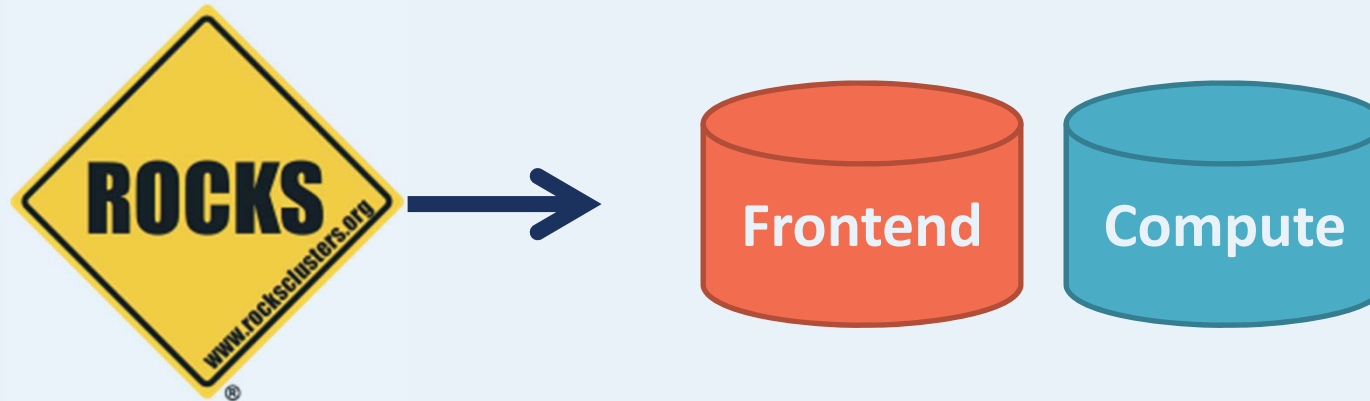
Kevin Lam and Karen Rodriguez (Prime2013) have been working on preparing VC images.

VM Tests							
Machine bit		32	64	64	64		64
OS bit		32	32	64	64		64
CentOS		5.9	5.2	5.9	6.4		6.4
gcc		4.1.2	4.1.2	4.1.2	4.7.7		4.4.7
Dock Version		6.2	6.2	6.2	6.2		6.6
	Developer's 6.2	Masternode Slavenode1	Master	Sailboat Kayak	CentOS1	Developer's 6.6	CentOS5
ZINC00158751	138.113892	138.11386	138.113861	138.114029	138.114029	-13.890112	-13.890112
ZINC00157960	21535.20898	21535.209	21535.20898	21535.30859	21535.30859	-17.022007	-17.022007
ZINC00158442	52.565872	52.56588	52.56588	52.565788	52.565788	-0.554128	-0.554128
ZINC00013564	-9.400209	10.801939	10.801939	-9.400218	-9.400218	-15.287382	-15.287382
ZINC01555236	-	503.24973	503.249725	1800544512	1800544512	-	-
ZINC00150863	21.139147	21.139143	21.139143	-12.467216	-12.467216	-13.396927	-13.396927
ZINC00152265	19.053925	30.238361	30.238361	30.238028	30.238028	-15.236704	-15.236704
ZINC00157111	-15.522114	-12.916615	-12.916615	-12.916644	-12.916644	-14.073079	-14.073079
ZINC00157152	-10.13739	-10.137384	-10.137384	-10.137392	-10.137392	-15.55584	-15.55584
ZINC00157402	168513.7031	168513.63	168513.625	168506.4063	168506.4063	-15.105305	-15.105305
ZINC00157467	-8.706671	-8.706671	-8.706671	-8.706783	-8.706783	-14.412565	-14.412565



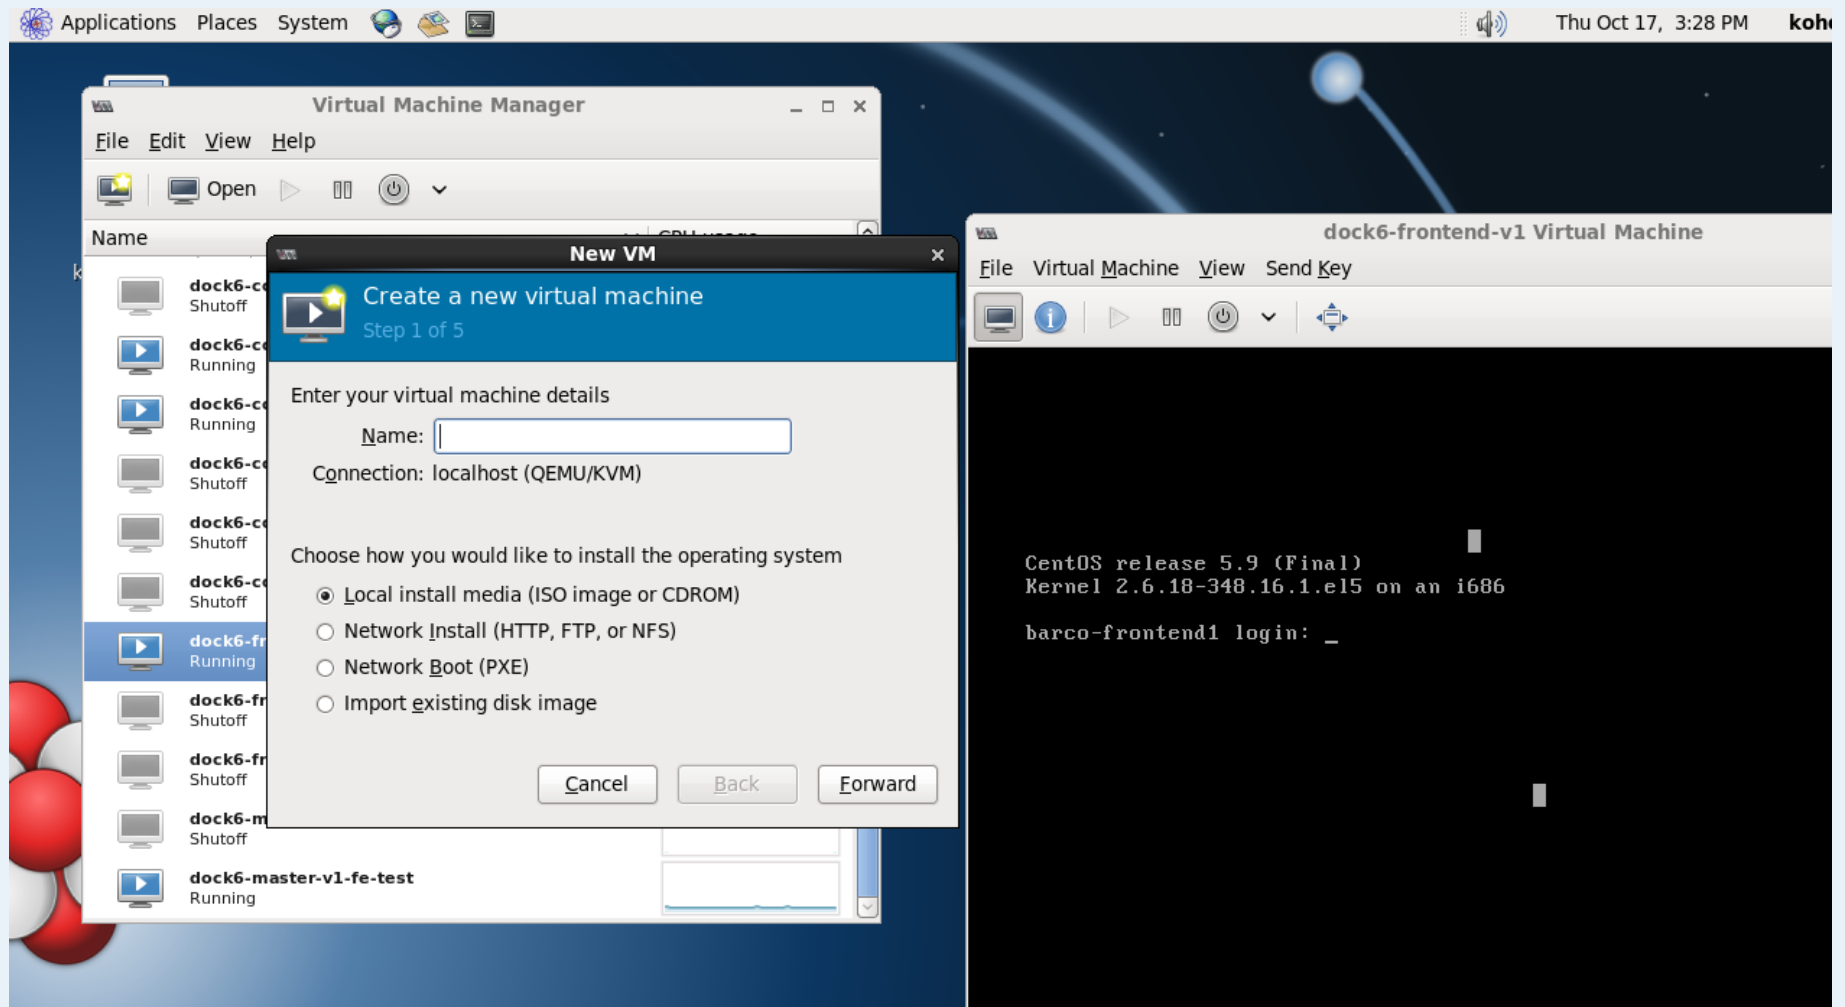
# Creating virtual cluster images

- Using Rocks:



- Easy to create a template of VC images
  - Not easy to customize OS architecture/libraries
- Creating from scratch
  - Using KVM & libvirt on Linux
  - Install OSs on the VMs manually

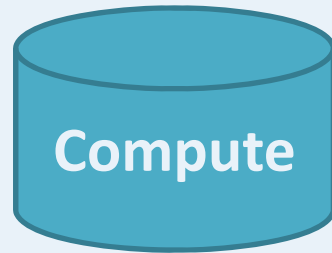
# Creating virtual cluster images from scratch



# Deployment of VC using pragma\_boot



Frontend



Compute



vc-in.xml

```
# pragma_boot -vcname dock6 -num_cpus 4
```



VM  
frontend

vc-out.xml



VM  
compute1

vc-out.xml



VM  
compute2

vc-out.xml



VM  
compute3

vc-out.xml



VM  
compute4

vc-out.xml



# Preparing vc-in.xml

```
<vc version='0.1'>
  <virtualization engine='kvm' type='hvm'
arch='x86_64'/>
  <driver>rocks6_client</driver>
  <frontend>
</frontend>
  <compute>
    <boot_dependency parent='frontend'>
      <wait type='clock' value='300'/>
    </boot_dependency>
  </compute>
  <networks>
    <network name='private'>
      <ipaddress>10.1.1.1</ipaddress>
      <netmask>255.255.0.0</netmask>
    </network>
    <frontend>
      <public>eth1</public>
    </frontend>
  </networks>
</vc>
```

/etc/libvirt/qemu/dock6-frontend.xml

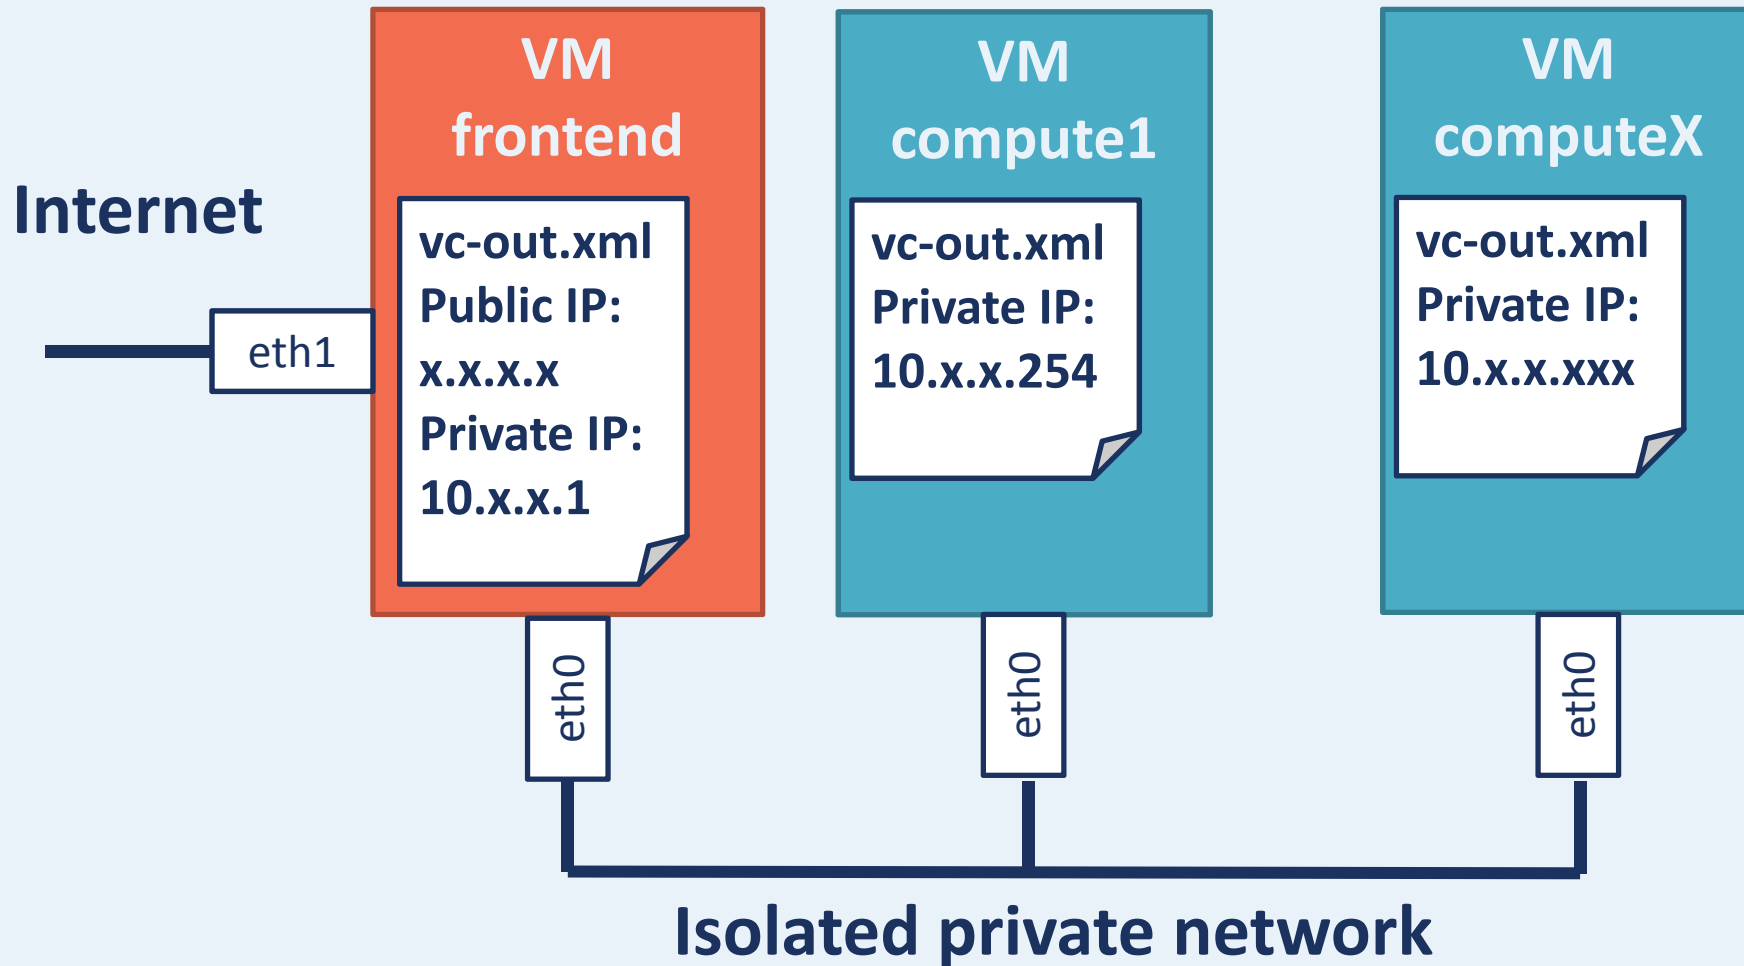
```
<domain type='kvm'>
  <name>
    dock6-frontend
  </name>
  .....
</domain>
```

/etc/libvirt/qemu/dock6-compute.xml

```
<domain type='kvm'>
  <name>
    dock6-compute
  </name>
  .....
</domain>
```

Copy & Paste

# Parsing vc-out.xml & configuring network

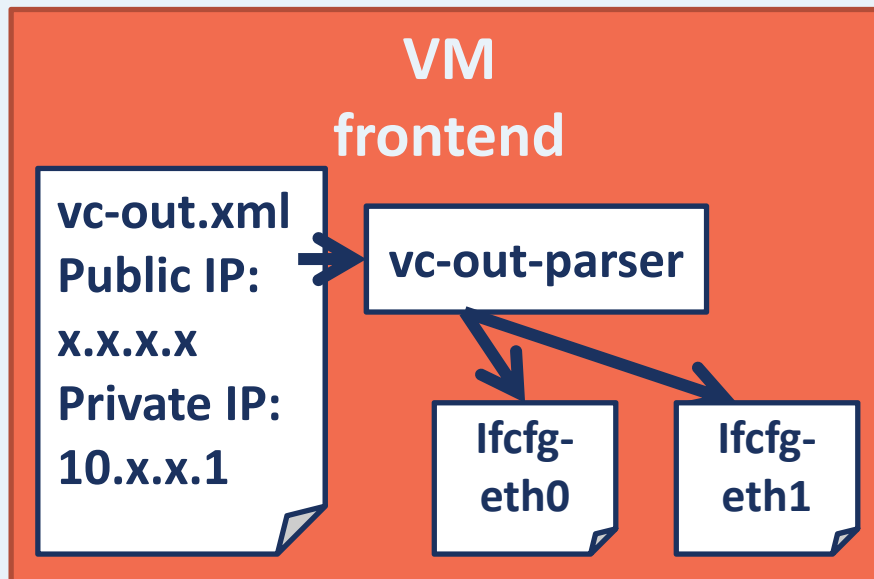


# vc-out parser

- <https://github.com/pragmagrid/vc-out-parser>

Developed by Luca

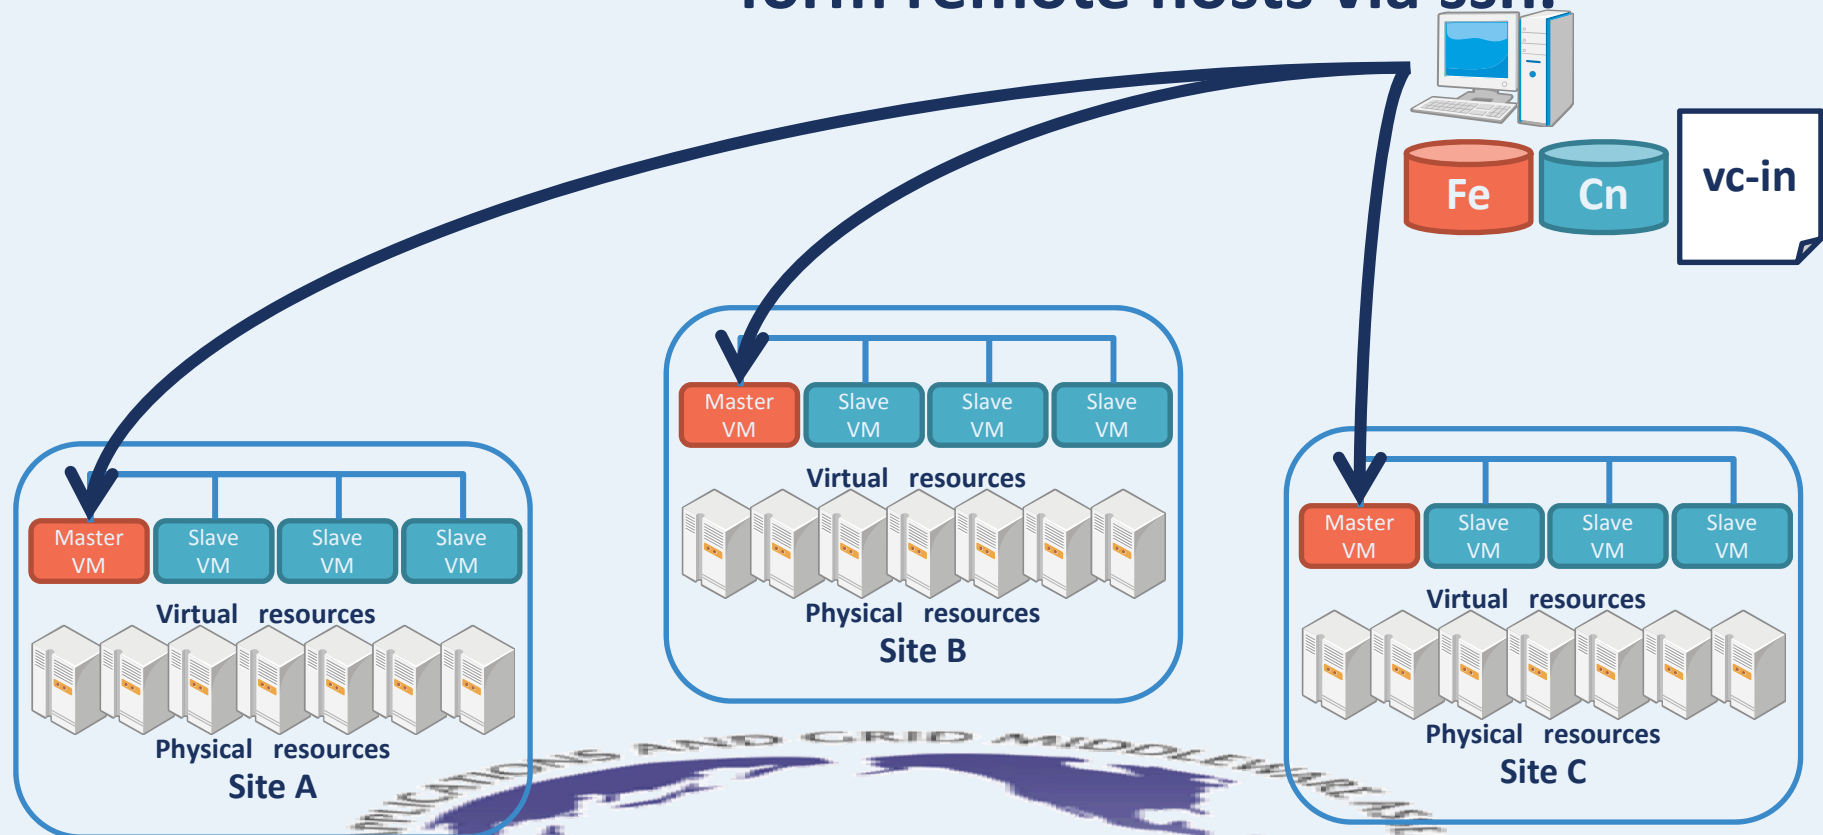
- Subset of configuration tools for Rocks VC (dynip)
- Generate basic network configuration for Redhat based Linux by parsing vc-out.xml



vc-out-paser generates network configurations in an early stage of the booting of the VM

# Virtual clusters on Pragma Cloud

Virtual clusters can be accessible from remote hosts via ssh.



# Demo



# Booting DOCK VC

```
# pragma_boot -vcname dock6 -num_cpus 8
```

OpenNebula  
Sunstone

Dashboard

System

Virtual Resources

Virtual Machines

Templates

Images

Files & Kernels

Infrastructure

Marketplace

OneFlow

Virtual Machines

9 TOTAL 9 ACTIVE 0 OFF 0 PENDING 0 FAILED

oneadmin

Refresh

Create

Filter

Play

Pause

Stop

Refresh

Delete

Grid

Search

ID	Owner	Group	Name	Status	Host	IPs	VNC
123	oneadmin	oneadmin	frontend_1(service_92)	RUNNING	asccmp050	10.2.0.1 163.220.57.212	
124	oneadmin	oneadmin	compute_0(service_92)	RUNNING	asccmp049	10.2.0.2	
125	oneadmin	oneadmin	compute_1(service_92)	RUNNING	asccmp048	10.2.0.3	
126	oneadmin	oneadmin	compute_2(service_92)	RUNNING	asccmp047	10.2.0.4	
127	oneadmin	oneadmin	compute_3(service_92)	RUNNING	asccmp046	10.2.0.5	
128	oneadmin	oneadmin	compute_4(service_92)	RUNNING	asccmp045	10.2.0.6	
129	oneadmin	oneadmin	compute_5(service_92)	RUNNING	asccmp044	10.2.0.7	
130	oneadmin	oneadmin	compute_6(service_92)	RUNNING	asccmp043	10.2.0.8	
131	oneadmin	oneadmin	compute_7(service_92)	RUNNING	asccmp042	10.2.0.9	

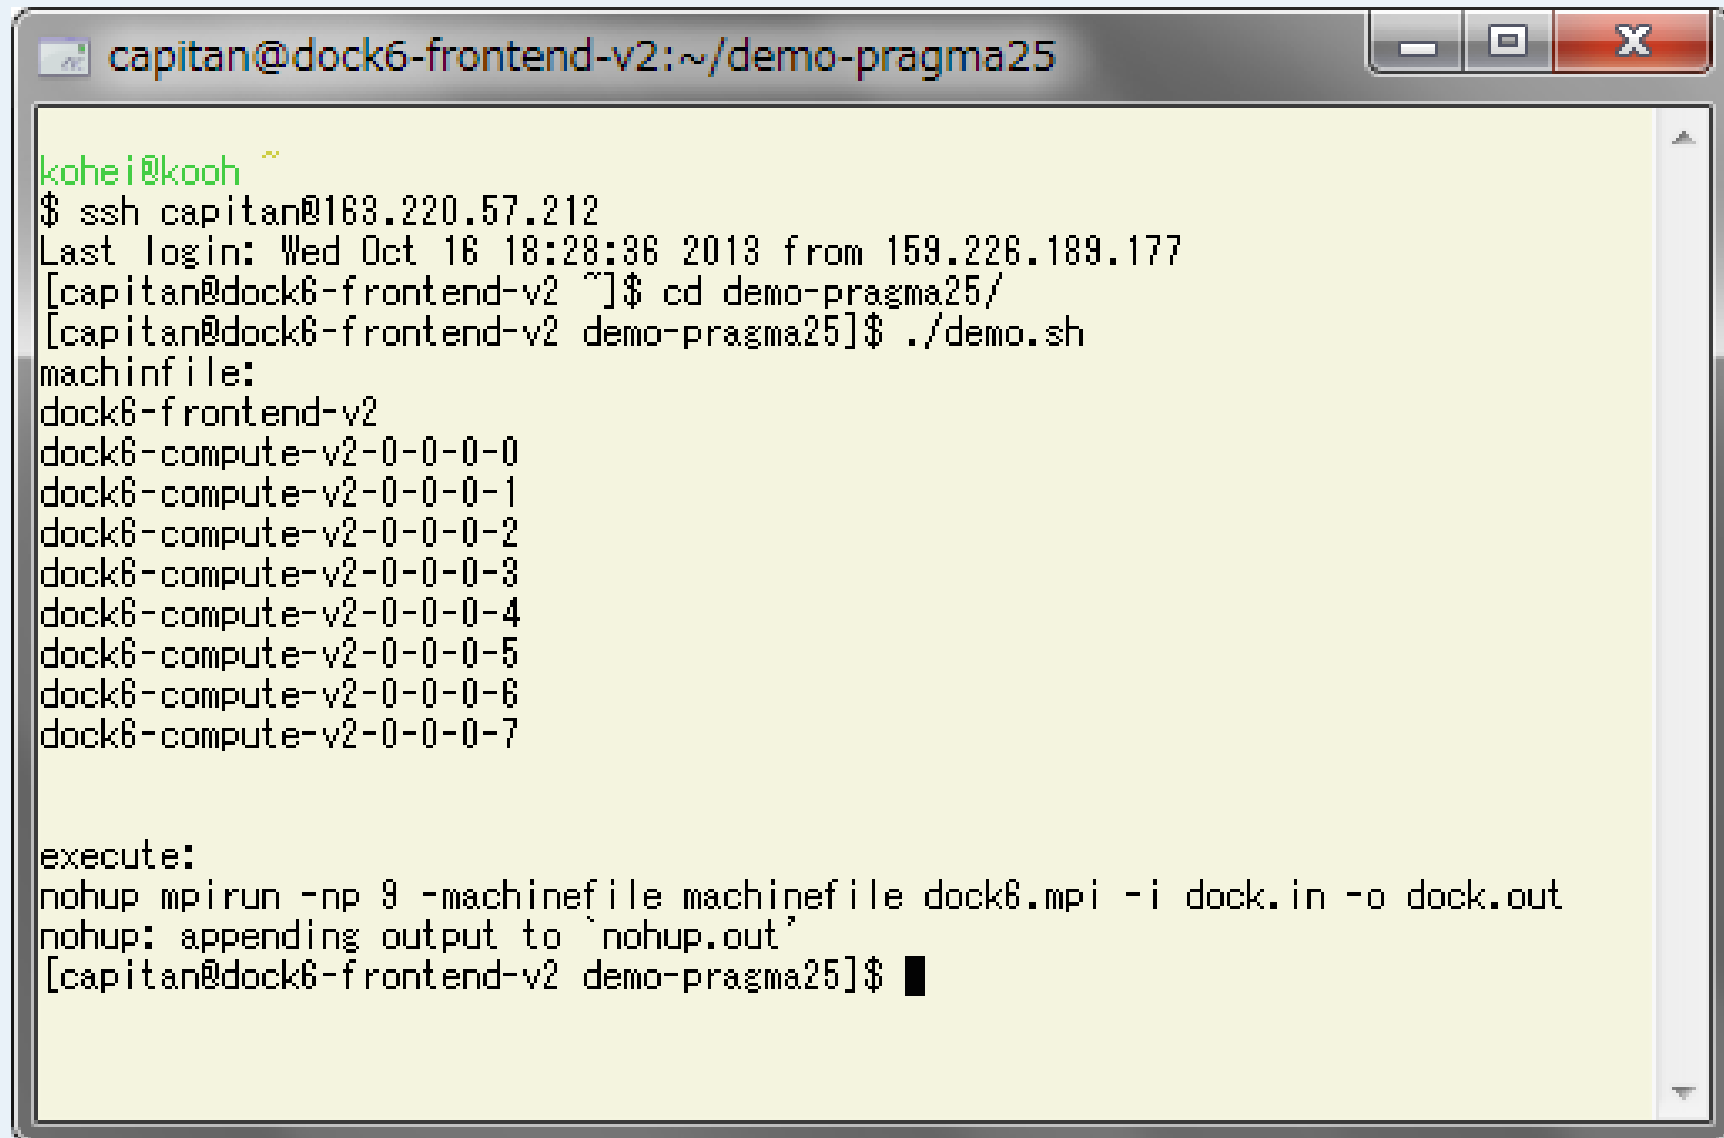
10

Showing 1 to 9 of 9 entries

<< < 1 > >>

OpenNebula 4.2.0 by C12G Labs.

# Login the VC and run applications



```
capitan@dock6-frontend-v2:~/demo-pragma25

kohei@kooh ~
$ ssh capitan@163.220.57.212
Last login: Wed Oct 16 18:28:36 2013 from 159.226.189.177
[capitan@dock6-frontend-v2 ~]$ cd demo-pragma25/
[capitan@dock6-frontend-v2 demo-pragma25]$ ./demo.sh
machinfile:
dock6-frontend-v2
dock6-compute-v2-0-0-0-0
dock6-compute-v2-0-0-0-1
dock6-compute-v2-0-0-0-2
dock6-compute-v2-0-0-0-3
dock6-compute-v2-0-0-0-4
dock6-compute-v2-0-0-0-5
dock6-compute-v2-0-0-0-6
dock6-compute-v2-0-0-0-7

execute:
nohup mpirun -np 8 -machinefile machinefile dock6.mpi -i dock.in -o dock.out
nohup: appending output to `nohup.out'
[capitan@dock6-frontend-v2 demo-pragma25]$
```

# Acknowledgements

- Luca Clementi, Nadya Williams, Philip Papadopoulos (UCSD)
- Pongsakorn U-chupala (NAIST)
- Yoshio Tanaka, Akihiko Ota (AIST)