

Fingerprint

Luca Clementi, Philip Papadopoulos




The problem

- ▶ Modern scientific research is dependent on reliable software environments
- ▶ Every applications relay on several shared libraries and external files

```
# mpirun  
mpirun: error while loading shared libraries: libxml2.so.2:  
cannot open shared object file: No such file or directory
```


ls -l

- ▶ 9 files
- ▶ 12 shared libraries
 - ▶ 4 dynamically loaded
- ▶ 4 OS packages




Paraview

- ▶ 139 files
- ▶ 382 shared libraries
 - ▶ 68 dynamically loaded
- ▶ 122 OS packages



Our approach

Fingerprinting application
dependencies



Command line tool

What is it...

Fingerprint is a software tool which can do:

- ▶ Dependencies **discovery**
- ▶ Dependencies **display**
- ▶ **Verify** dependencies
- ▶ **Compose** a new system based on a list of dependencies

Static Fingerprinting

Collect requirements from a list of binary files.

- ▶ Parse ELF (Executable Linkable Format) headers.

Static Fingerprinting

```
# fingerprint -c /bin/ls
```

Output written to a file: Swirl

But...

Linux supports dynamic loading of shared libraries...

- ▶ Explicit dependency: a dependency which is declared in the binary header
- ▶ Implicit dependency: a dependency which is loaded dynamically at runtime (`dlopen()`)

Dynamic Fingerprinting

Dynamic Fingerprinting

```
# fingerprint -c -x "ls -l"
```

- ▶ run the process and trace it
- ▶ detect each invocation of `open()`
- ▶ read `/proc/ID/mmaps`

Swirl File

The Swirl file stores info regarding all the traced files.

For each file it saves its:

- ▶ Full path and symbolic links
- ▶ Explicit dependencies (soname)
- ▶ Implicit dependencies (pointers to SwirlFile)
- ▶ Major and minor version of symbols provided
- ▶ Opened files (pointers to SwirlFile)
- ▶ Environment variables
- ▶ Hash, architecture, package, rpath, executable

Dependencies Display

A Swirl can be displayed with text


```
# fingerprint -d
/bin/ls
/lib/x86_64-linux-gnu/ld-2.15.so
/lib/x86_64-linux-gnu/libacl.so.1.1.0
/lib/x86_64-linux-gnu/libc-2.15.so
/lib/x86_64-linux-gnu/librt-2.15.so
/lib/x86_64-linux-gnu/libselinux.so.1
/lib/x86_64-linux-gnu/libattr.so.1.1.0
/lib/x86_64-linux-gnu/libpthread-2.15.so
/lib/x86_64-linux-gnu/libdl-2.15.so
/lib/x86_64-linux-gnu/libnss_compat-2.15.so --(Dyn)--
/lib/x86_64-linux-gnu/libnsl-2.15.so --(Dyn)--
/lib/x86_64-linux-gnu/libnss_nis-2.15.so --(Dyn)--
/lib/x86_64-linux-gnu/libnss_files-2.15.so --(Dyn)--
Opened files:
/etc/localtime
/etc/group
/etc/passwd
/usr/lib/locale/locale-archive
/usr/lib/x86_64-linux-gnu/gconv/gconv-modules.cache
/usr/share/locale-langpack/en/LC_MESSAGES/coreutils.mo
/etc/locale.alias
/etc/nsswitch.conf
/proc/filesystems
```

Dependencies Display

A Swirl can be displayed with Graphviz.

Non-interactive bash `bash -c "ls"`

Swirl 2013-03-28 14:24 `"/bin/bash -c ls"`



Interactive bash ls

The diagram illustrates the mapping from the Execution Domain to the Linker Domain and then to the Package Domain. The Execution Domain (left) lists various executables and scripts. The Linker Domain (middle) shows the corresponding linker scripts and object files. The Package Domain (right) shows the specific package names and versions for each linker script and object file.

Execution Domain	Linker Domain	Package Domain
/usr/.groups	/lib/.libnss_nis-2.15.so	libc6 2.15-0ubuntu10.3 amd64
/bin/bash	/lib/.libnss_compat-2.15.so	libc6 2.15-0ubuntu10.3 amd64
/usr/.sort	/lib/.libnsl-2.15.so	libc6 2.15-0ubuntu10.3 amd64
/bin/dash	/lib/.libnss_files-2.15.so	libc6 2.15-0ubuntu10.3 amd64
/usr/.ld.so.colors	ld.so.2	libc6 2.15-0ubuntu10.3 amd64
/usr/.ldname	ld.so.2	libc6 2.15-0ubuntu10.3 amd64
/bin/uname	ld.so.2	libc6 2.15-0ubuntu10.3 amd64
/usr/.basename	ld.so.2	libc6 2.15-0ubuntu10.3 amd64
/bin/sed	ld.so.2	libc6 2.15-0ubuntu10.3 amd64
/bin/ls	ld.so.2	libc6 2.15-0ubuntu10.3 amd64
	libc.so.5	libc6 2.15-0ubuntu10.3 amd64
	libpthread.so.0	libc6 2.15-0ubuntu10.3 amd64
	ld-linux-x86-64.so.2	libc6 2.15-0ubuntu10.3 amd64
	libselinux.so.1	libselinux 2.10-4.1ubuntu1 amd64
	librt.so.1	librt 1.2.2-51.5ubuntu1 amd64
	libacl.so.1	libacl 2.2.51-5ubuntu1 amd64

System Validation

Does this system satisfy this set of requirements?

```
# fingerprint -yv  
The file output.swirl failed.
```

Missing Dependencies:

```
libunwind-x86_64.so.8()(64bit)  
libunwind-ptrace.so.0()(64bit)
```

Software Stack Integrity

Did any of the requirements change?

```
# fingerprint -yvi  
The file output.swirl failed.
```

Modified Dependencies:

```
libunwind-ptrace.so.0()(x86_64)  
wrong hash
```

Compose a Cluster from a Swirl

- ▶ Given a Swirl file we can compose a software environments which provides all the requirement specified in a given Swirl
- ▶ It supports only Rocks Clusters

Definition

Rocks is an open source Linux distribution based on Centos that enables users to easily build computational clusters

Compose a Roll

- ▶ Create a Swirl (with dynamic tracing)

```
# fingerprint -cx "ls -l"
# ls -lh output.swirl
-rw-rw-r-- 1 root root 26K Oct 16 09:10 output.swirl
```

- ▶ Create a Swirl archive

```
# fingerprint -cr
# ls -lh output.tar.gz
-rw-rw-r-- 1 root root 13M Oct 16 09:16 output.tar.gz
```

Compose a Roll

- ▶ Copy output.tar.gz on a Rocks Cluster and create a Roll.

To way to port an application:

- ▶ Native Mode
fingerprint -cm
- ▶ Sandbox Mode
fingerprint -cmz

Definition

Roll is the name for a software package in Rocks

Port Dock6 on RedHat 6

Dock6 is a virtual machine used in Pragma 25 and 26 for docking (NAIST).

- ▶ It is based on Centos 5.9
- ▶ On a 32 bit system
- ▶ Uses mpich 3.0.4

Port binaries on on a Rocks 6.1 (Based on Centos 6.3) 64 bit.

Port Dock6 on RedHat 6

- ▶ Origin: Dock6 (Centos 5.9 32 bit)
@rocks-213
- ▶ Destination: Rocks 6.1 (Centos 6.3 64 bit)
@cluster

Questions?

Software available at:

<https://github.com/rocksclusters/FingerPrint>

For info:

- ▶ lclementi at ucsd.edu