Practicality and Feasibility of Improving Linux Container Utilization with Task Rebalancing Strategy

Pongsakorn U-chupala
Nara Institute of Science and Technology

PRAGMA30 | Manila Philippines

HPC is everywhere

HPC workloads are diverse

CPU-bound

NGINX

MPI

Less fault-tolerant

More fault-tolerant

redis

Memory-bound

Resource sharing with VM is cost effective

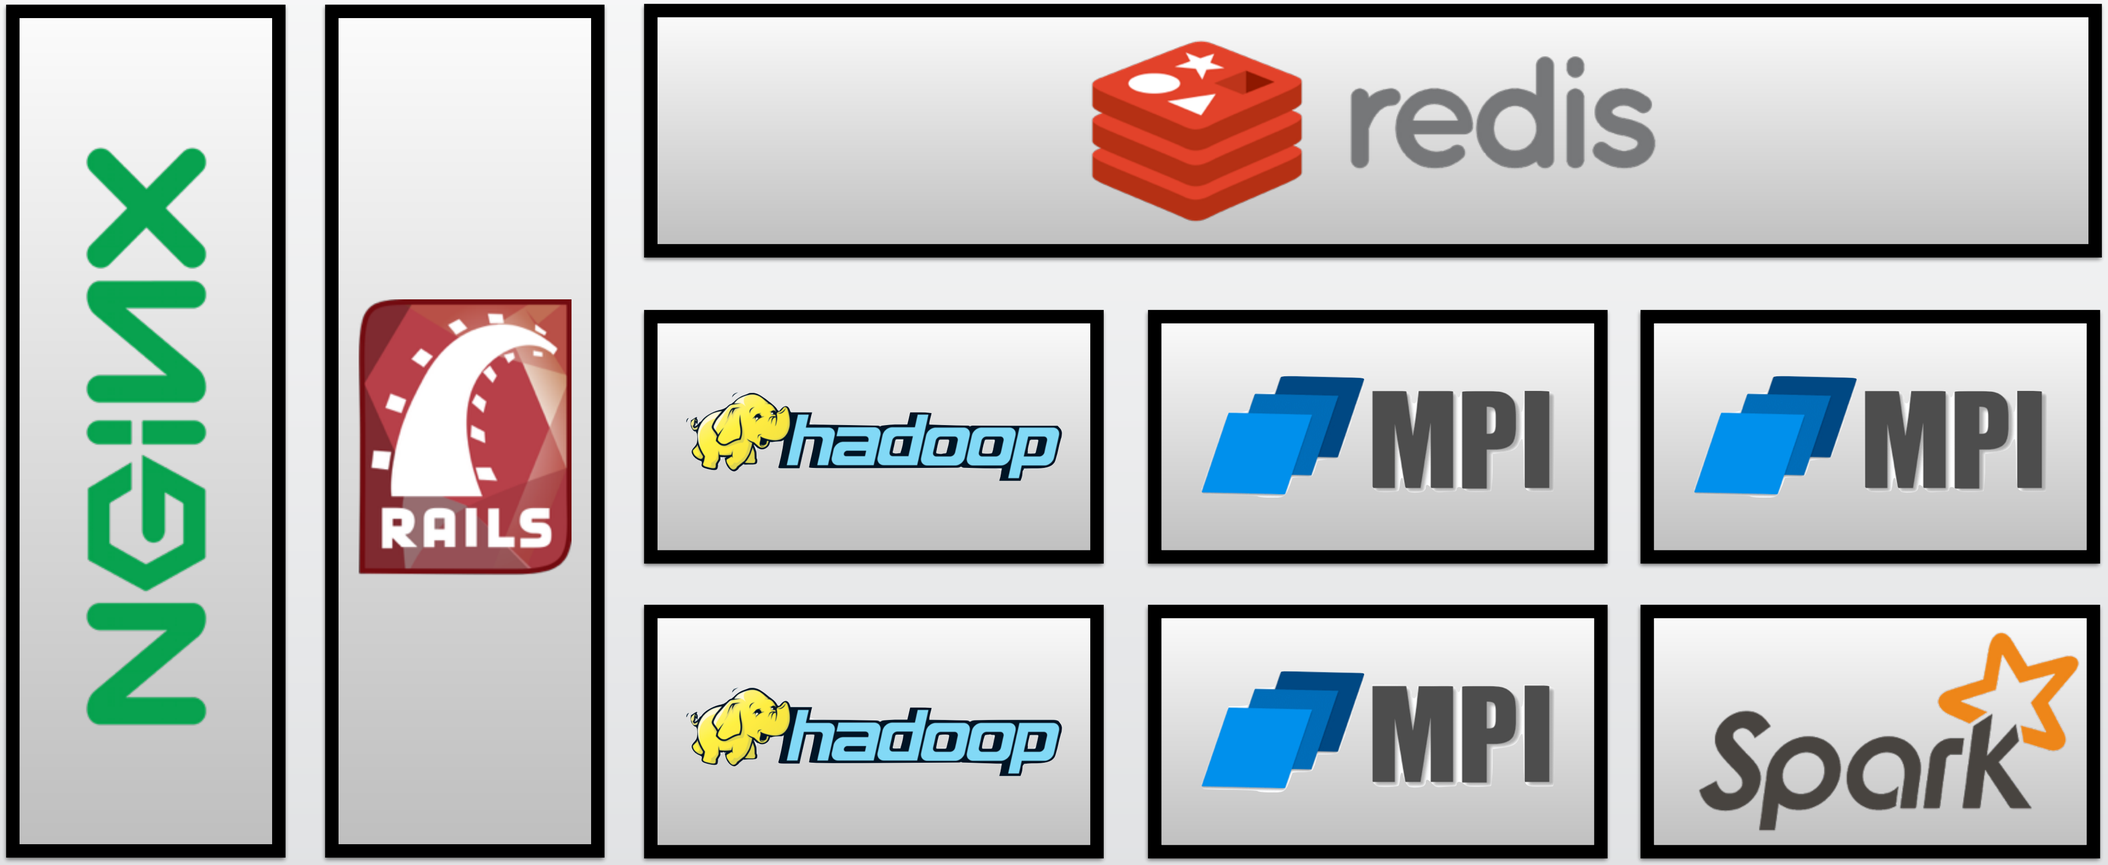

PROBLEM: Scheduling is inefficient

Moving Hourly Average
from Google Cluster Data
(color indicates priority)

CPU

Memory

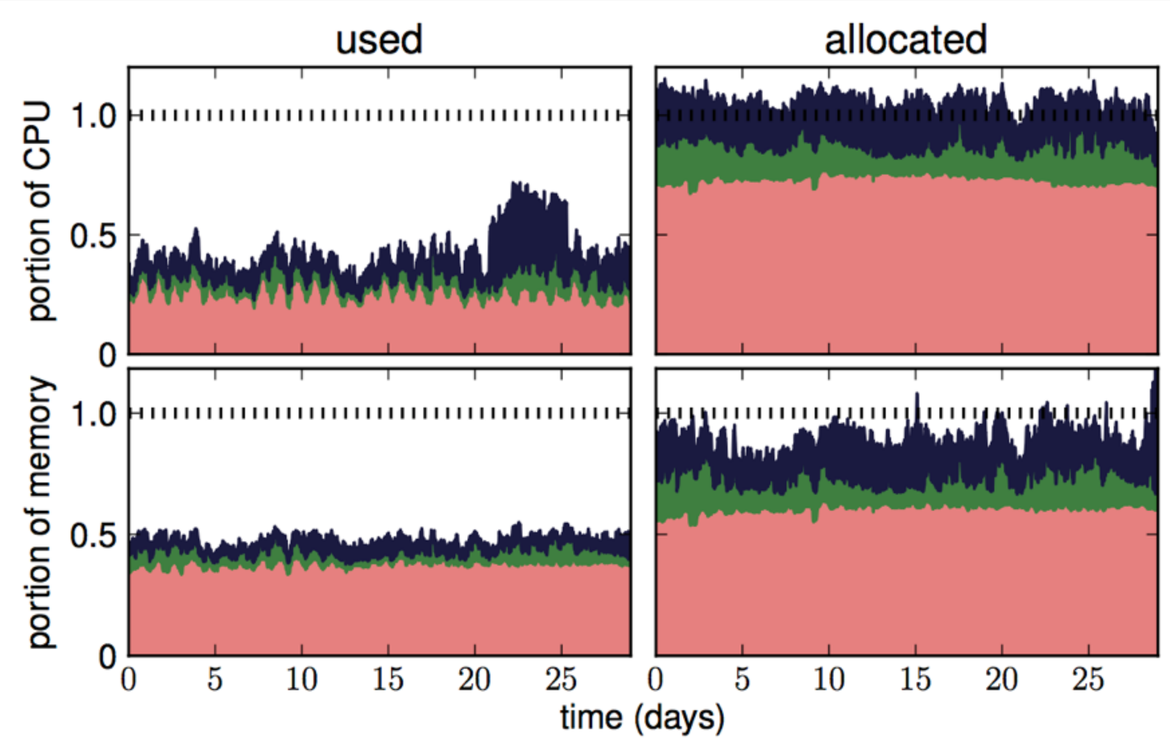

C. Reiss, A. Tumanov, G. R. Ganger, R. H. Katz, and M. a. Kozuch, “Heterogeneity and dynamicity of clouds at scale,” Proc. Third ACM Symp. Cloud Comput. - SoCC ’12, pp. 1–13, 2012.

Overcommitting

- Allocating more resource than available => Increase chance of task failure
- **Overcommit factor:** How much over-commit is allowed? (example: 1.2x of available resources)
  - Too high => Increase task failure rate
  - Too low => Resources are underutilized

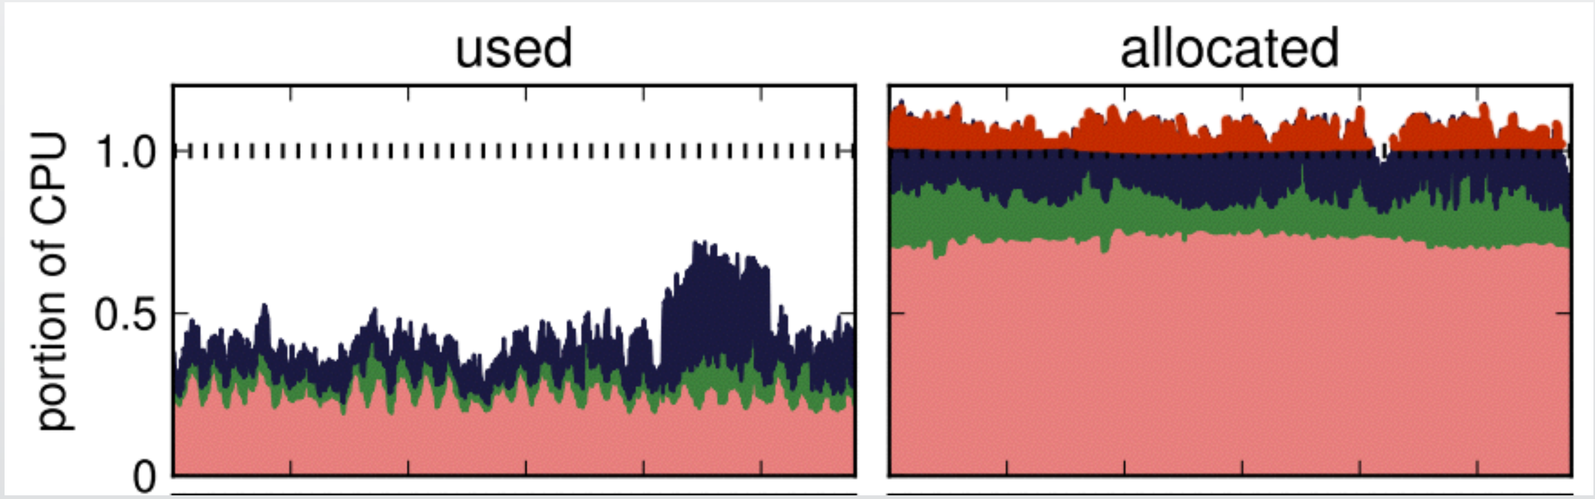

Advent of Linux Container

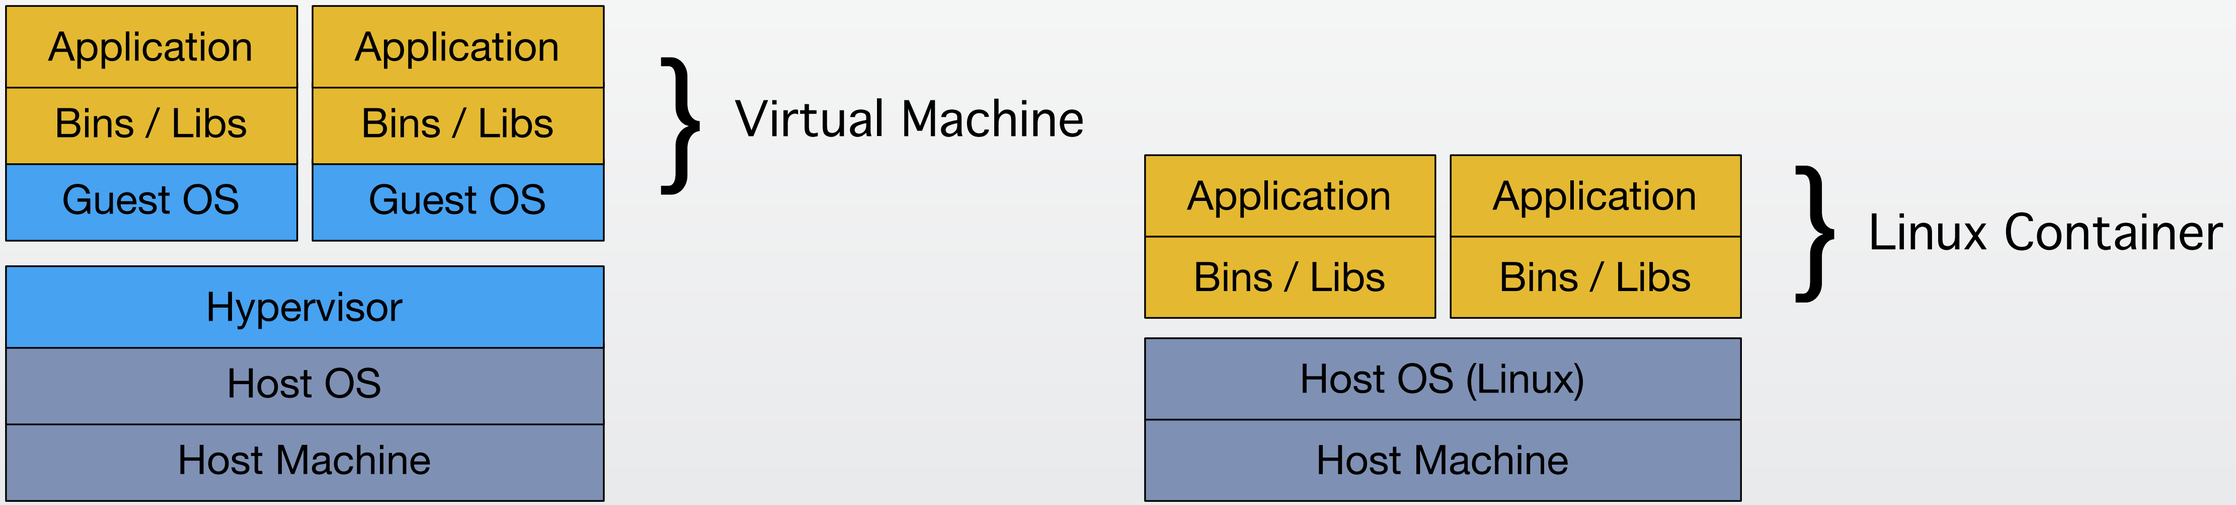

Less overhead → Lower spin-up time

Migration Time

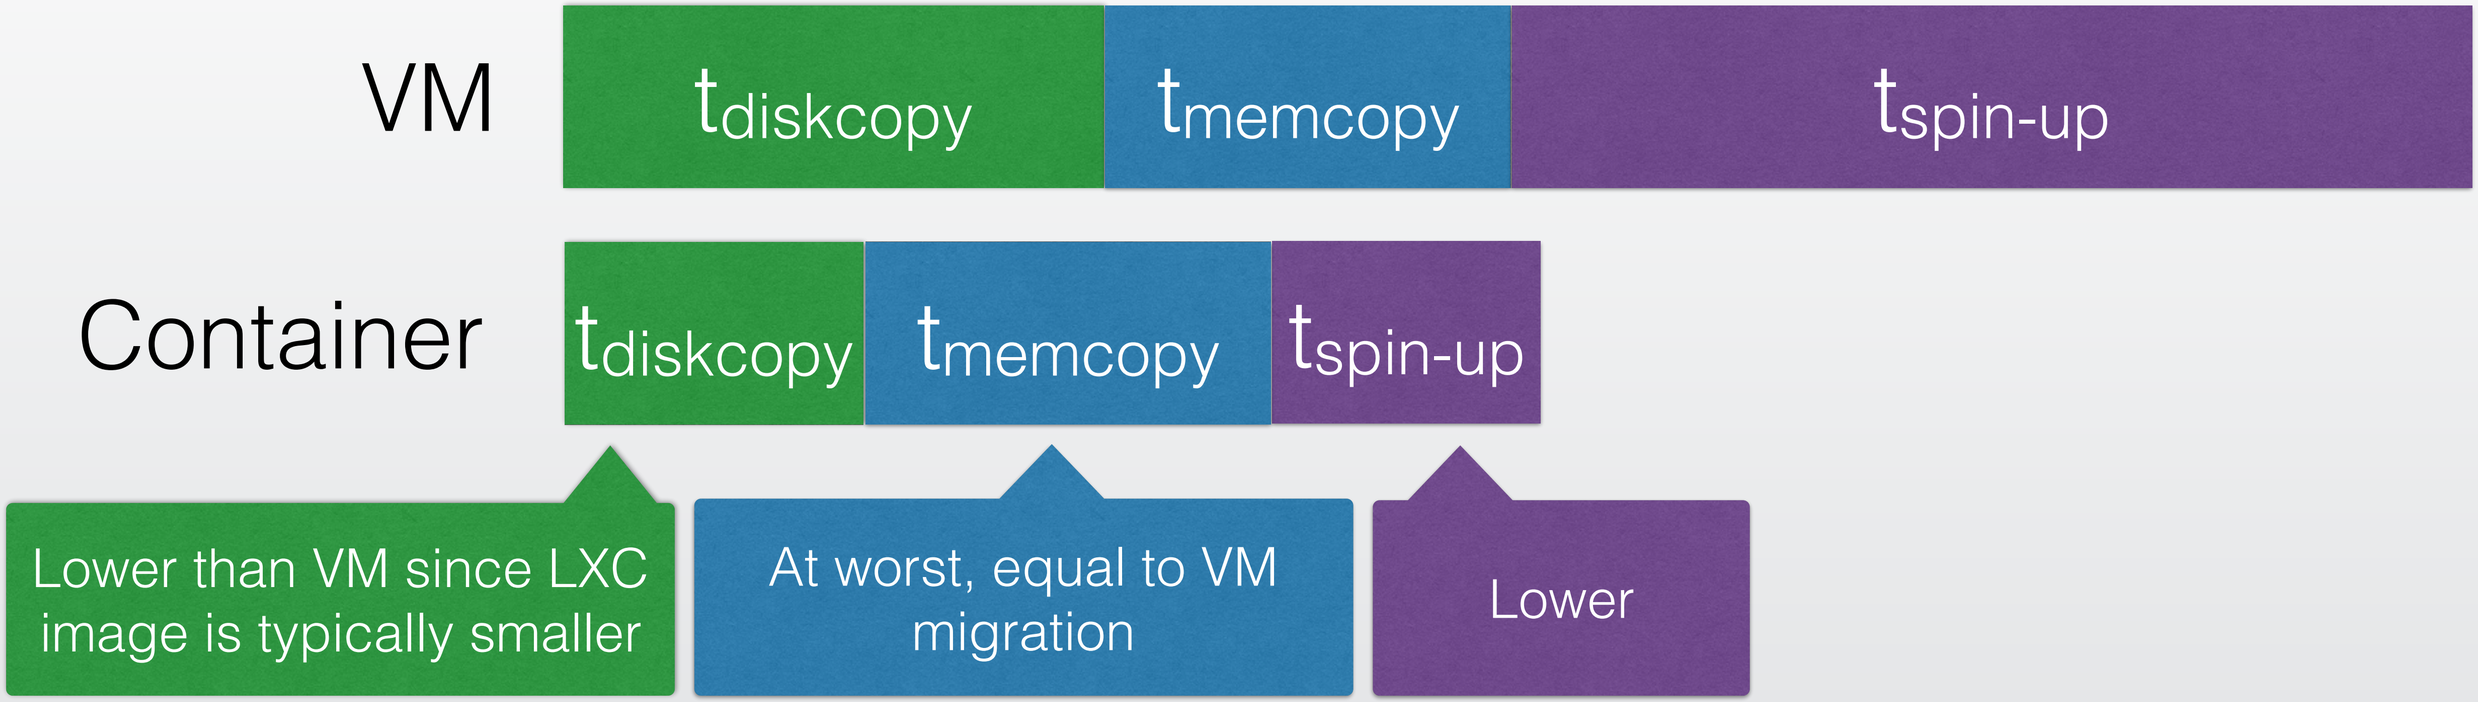

Faster migration → Enabling task micromanagement

PROPOSAL: Task Rebalancing

- Real-time host load-balancing
  - Increase optimal overcommit factor => Increase utilization
- Minimal interference to the scheduler and scheduled tasks
- Easier to explain with example

Example

Task


AVG Resource Utilization
(Abstract Resource Unit)
Task
Resource Allocation
(Abstract Resource Unit)

Scheduling

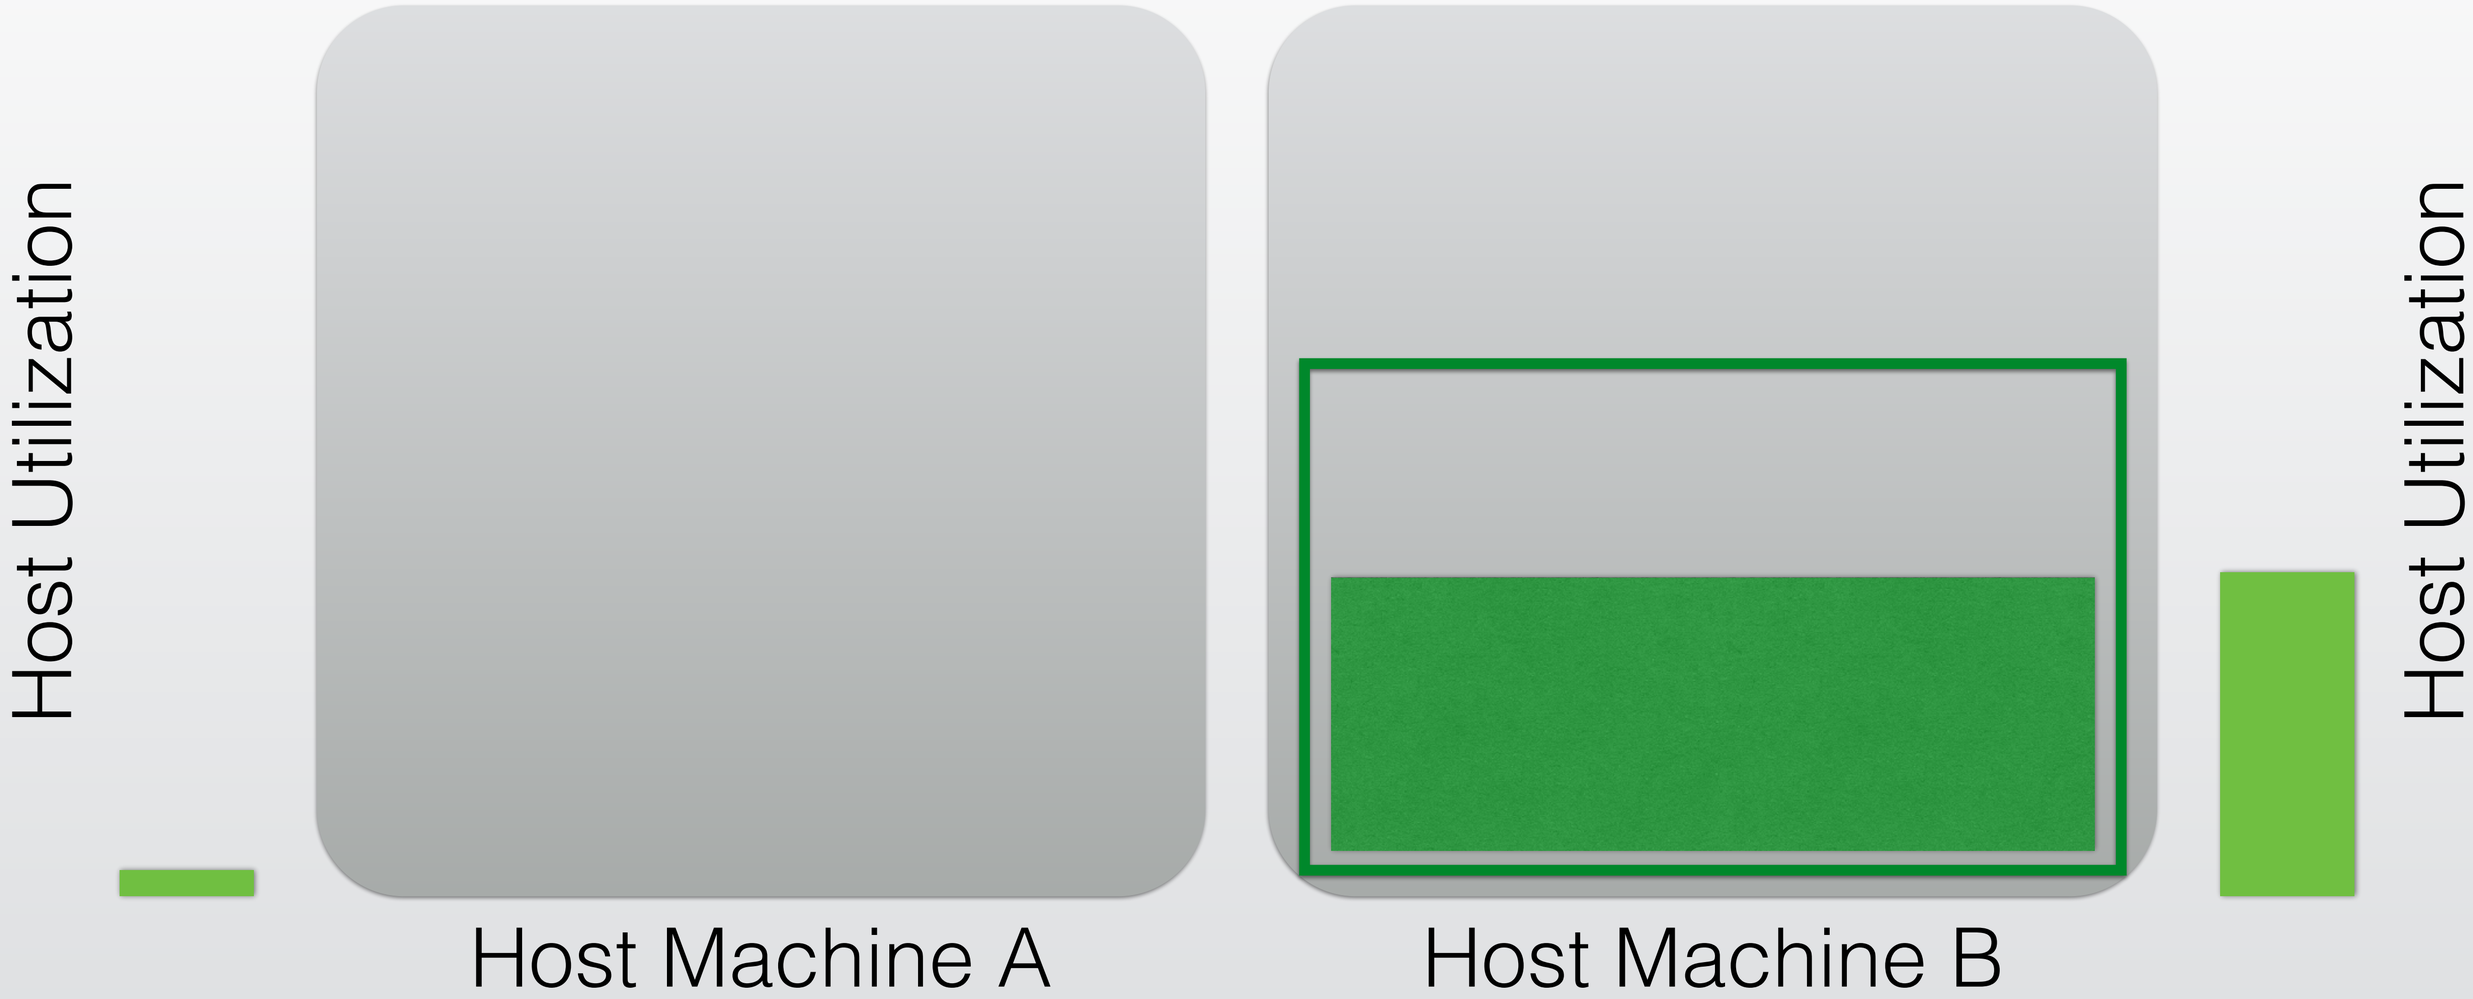
Host Utilization
Host Machine A
Host Machine B
Host Utilization

Scheduling

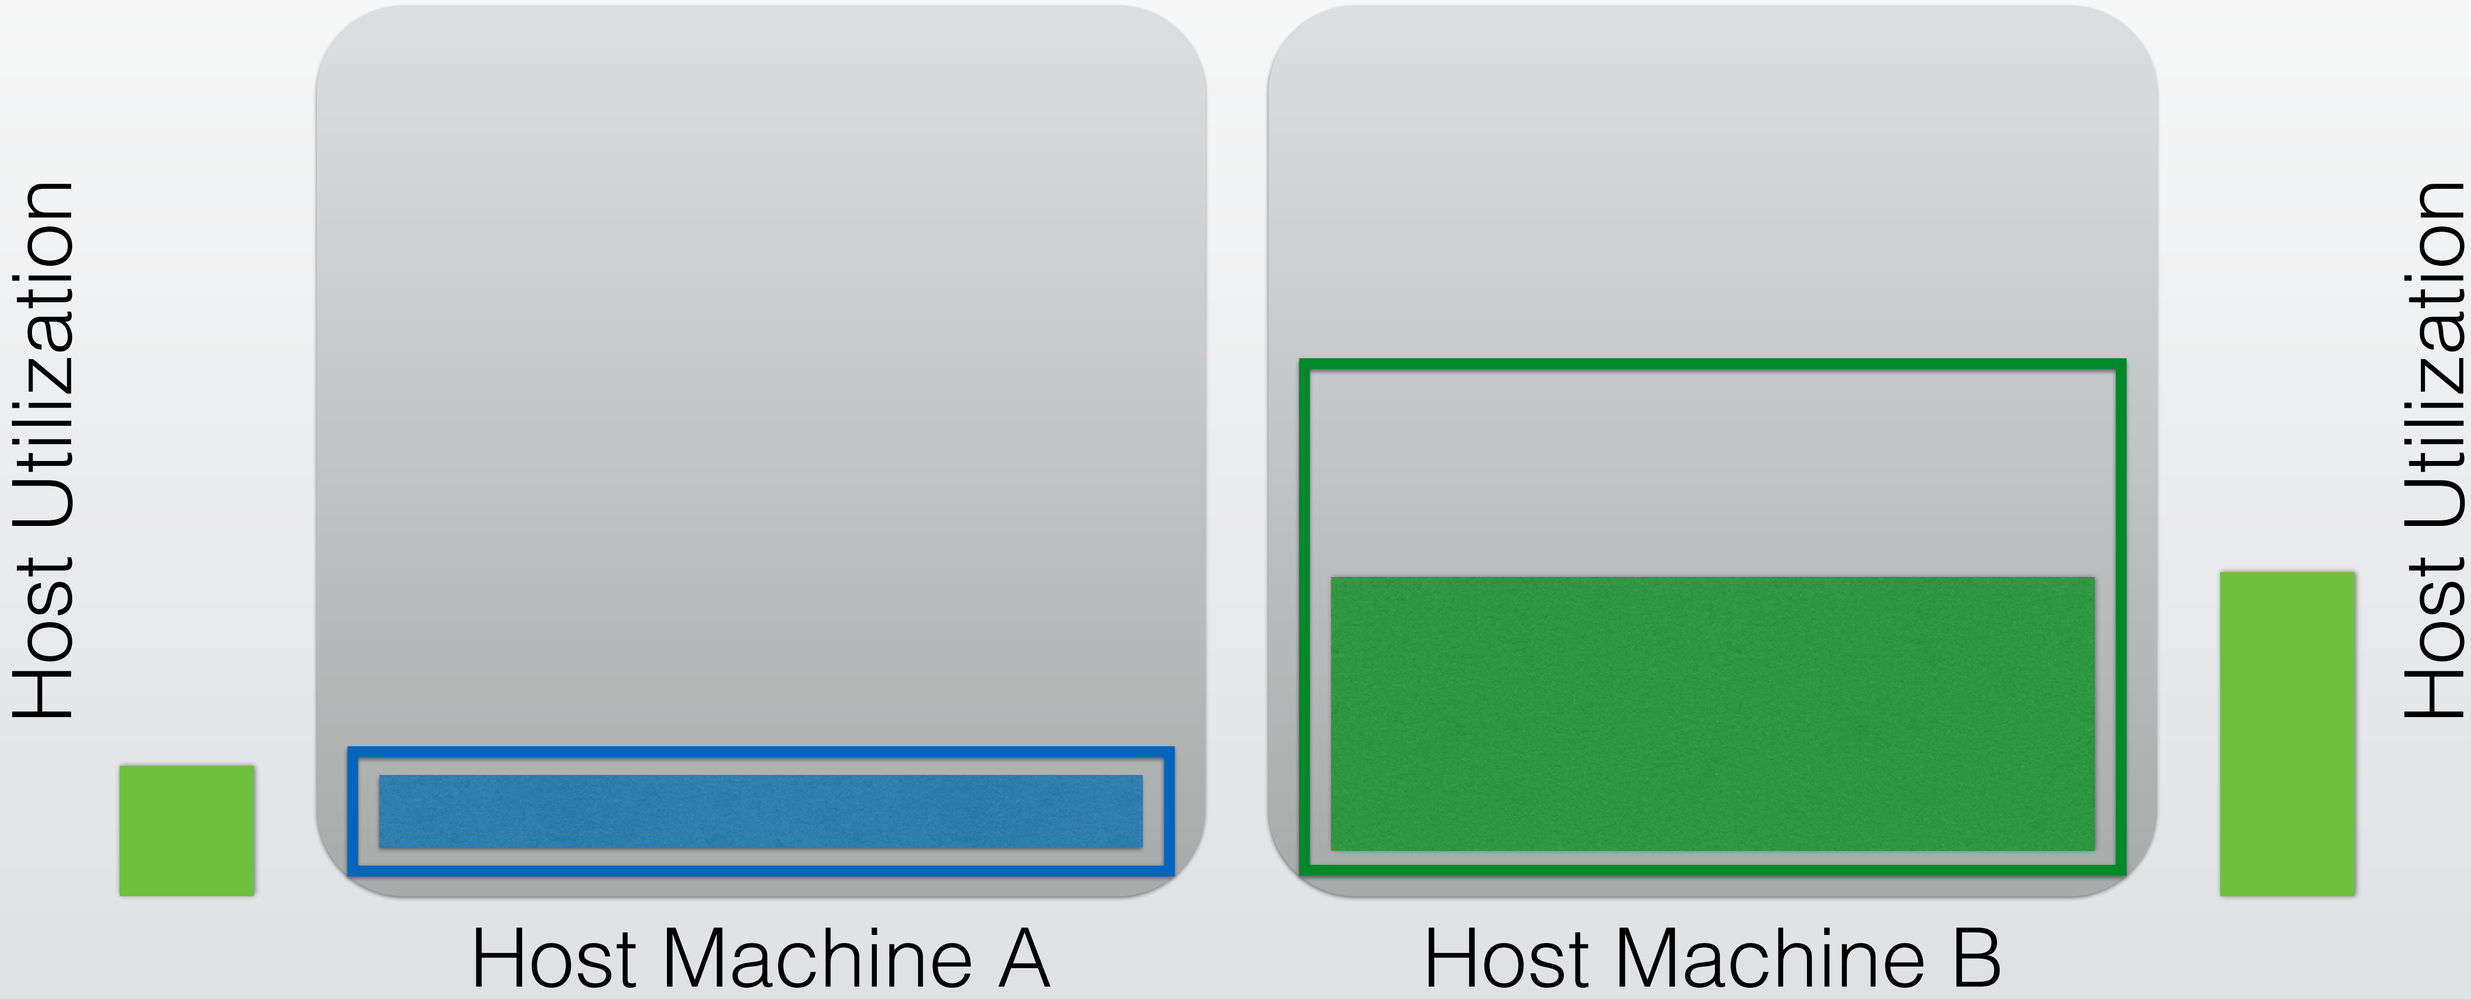
Host Utilization
Host Machine A
Host Machine B
Host Utilization

Scheduling

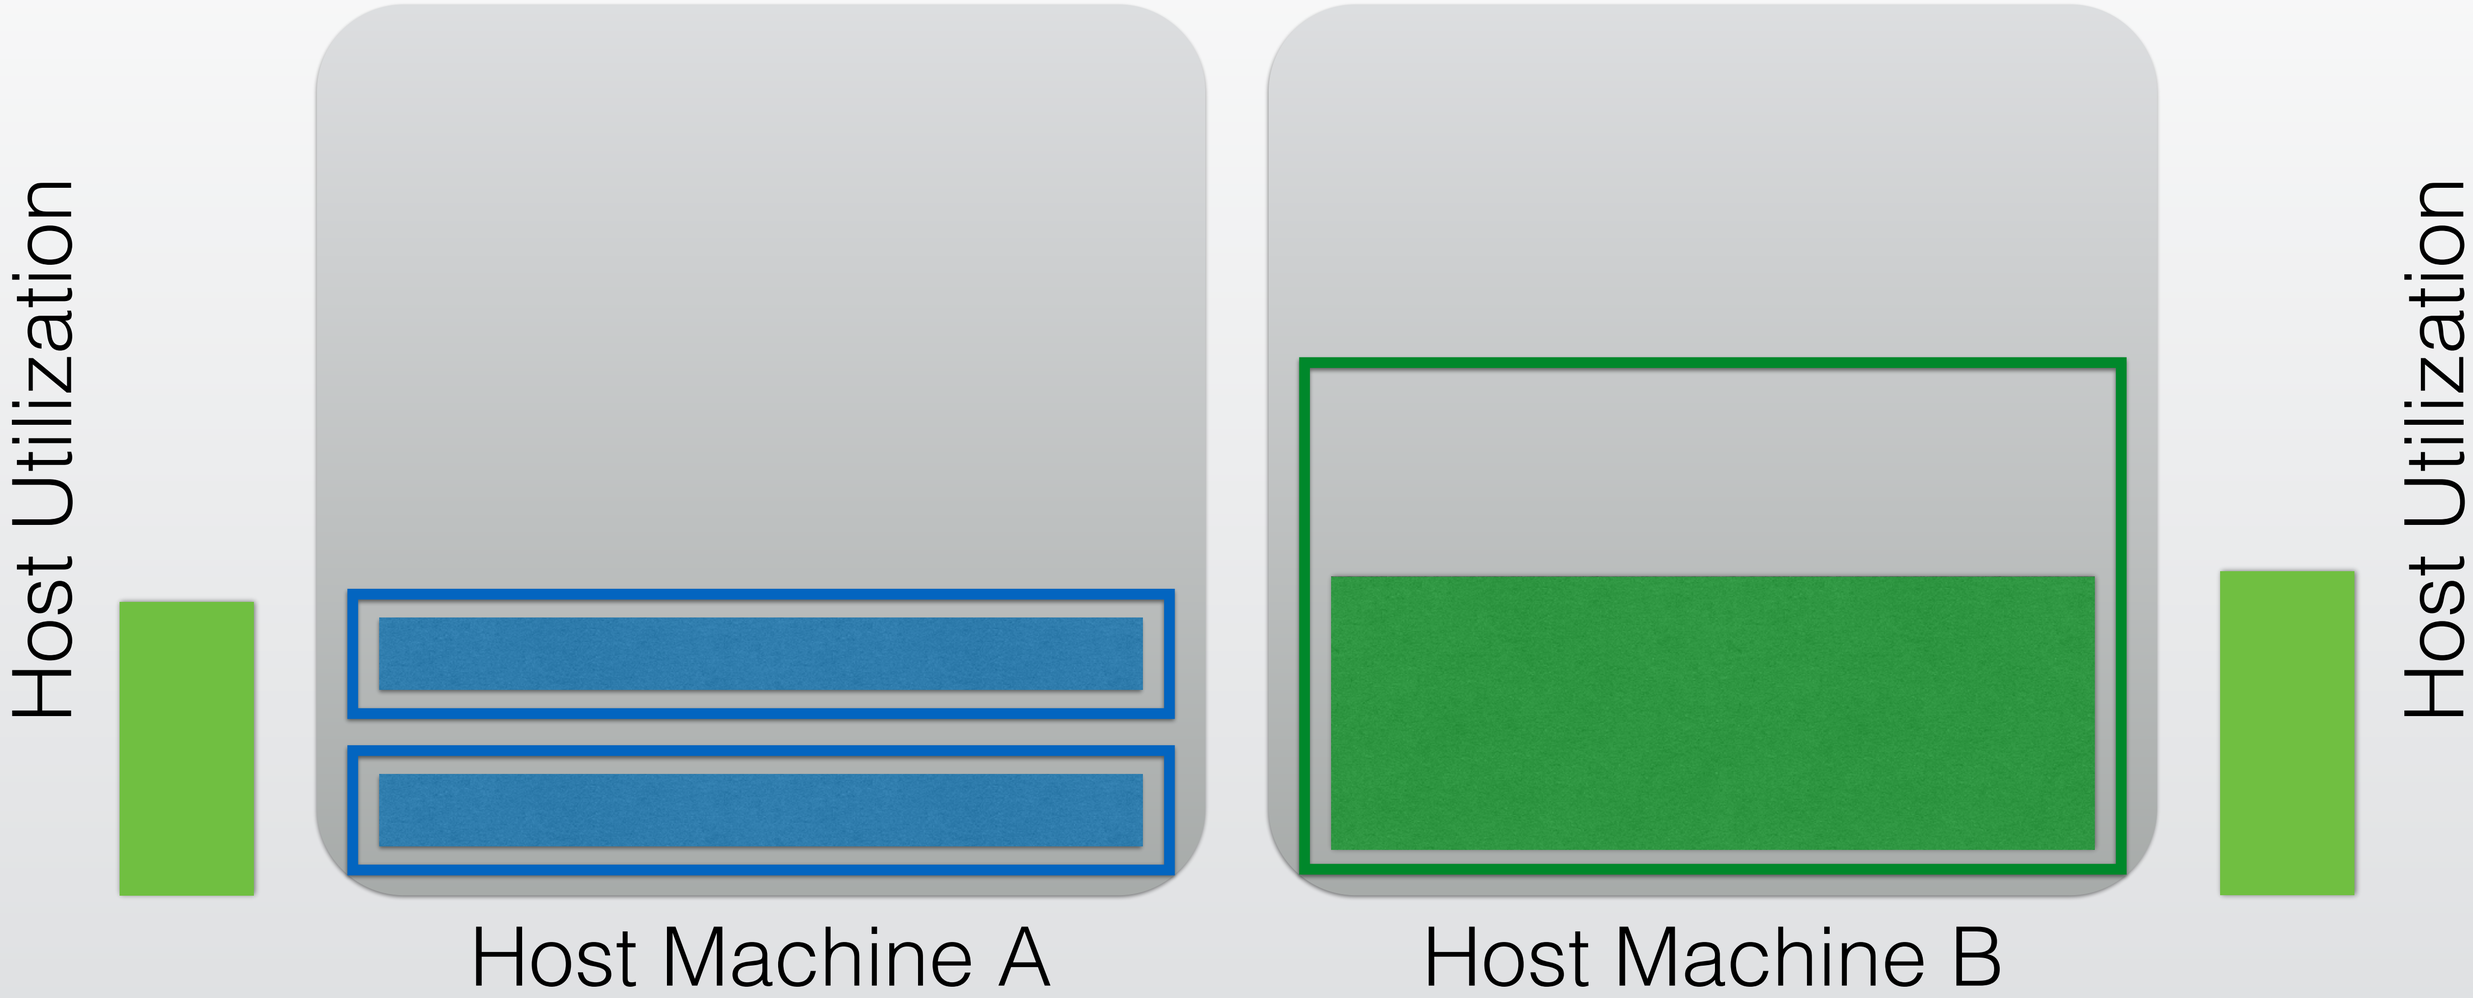
Host Utilization
Host Machine A
Host Machine B
Host Utilization

Scheduling

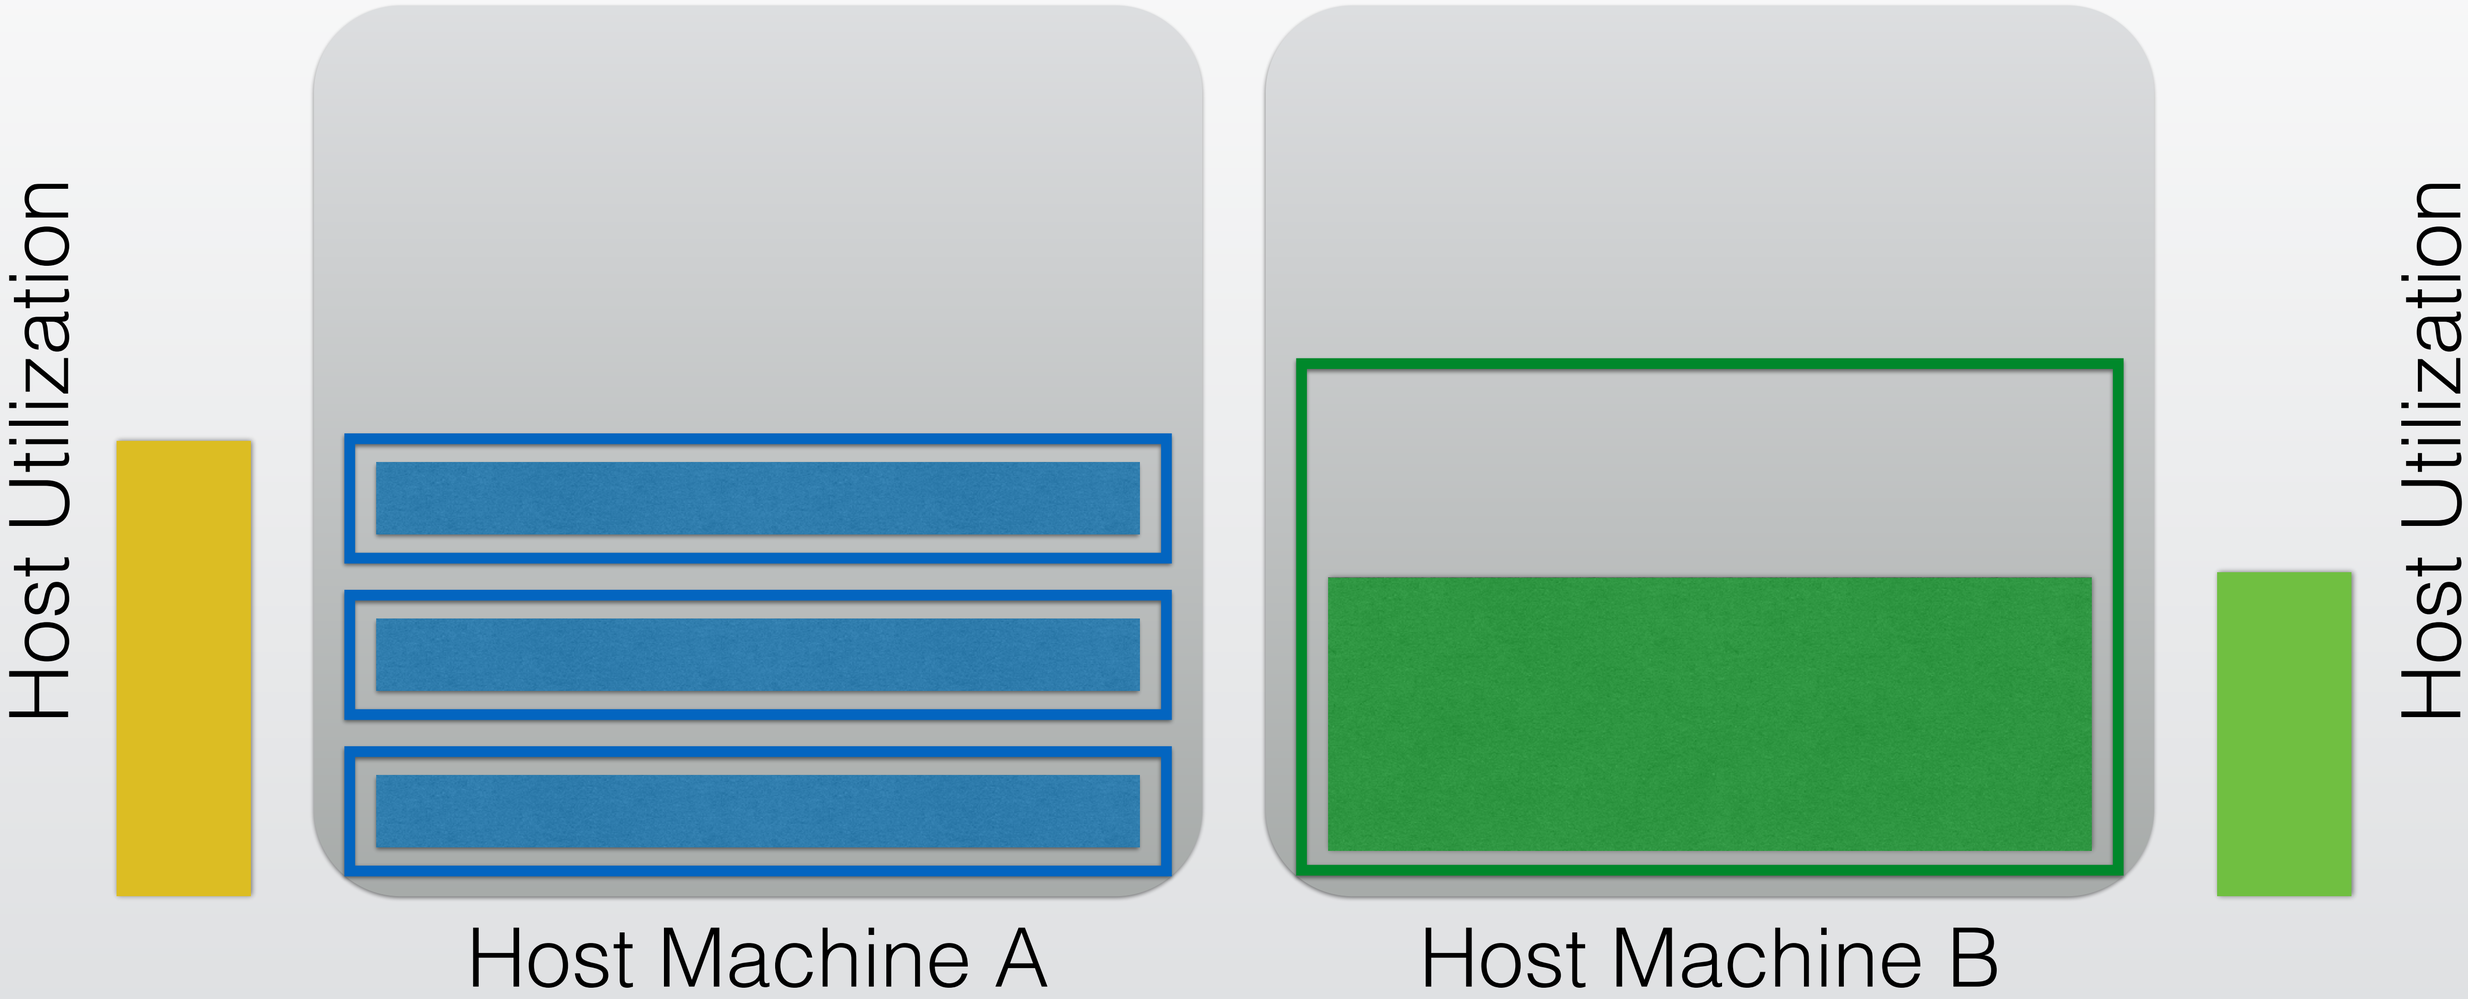
Host Utilization
Host Machine A
Host Machine B
Host Utilization

Scheduling

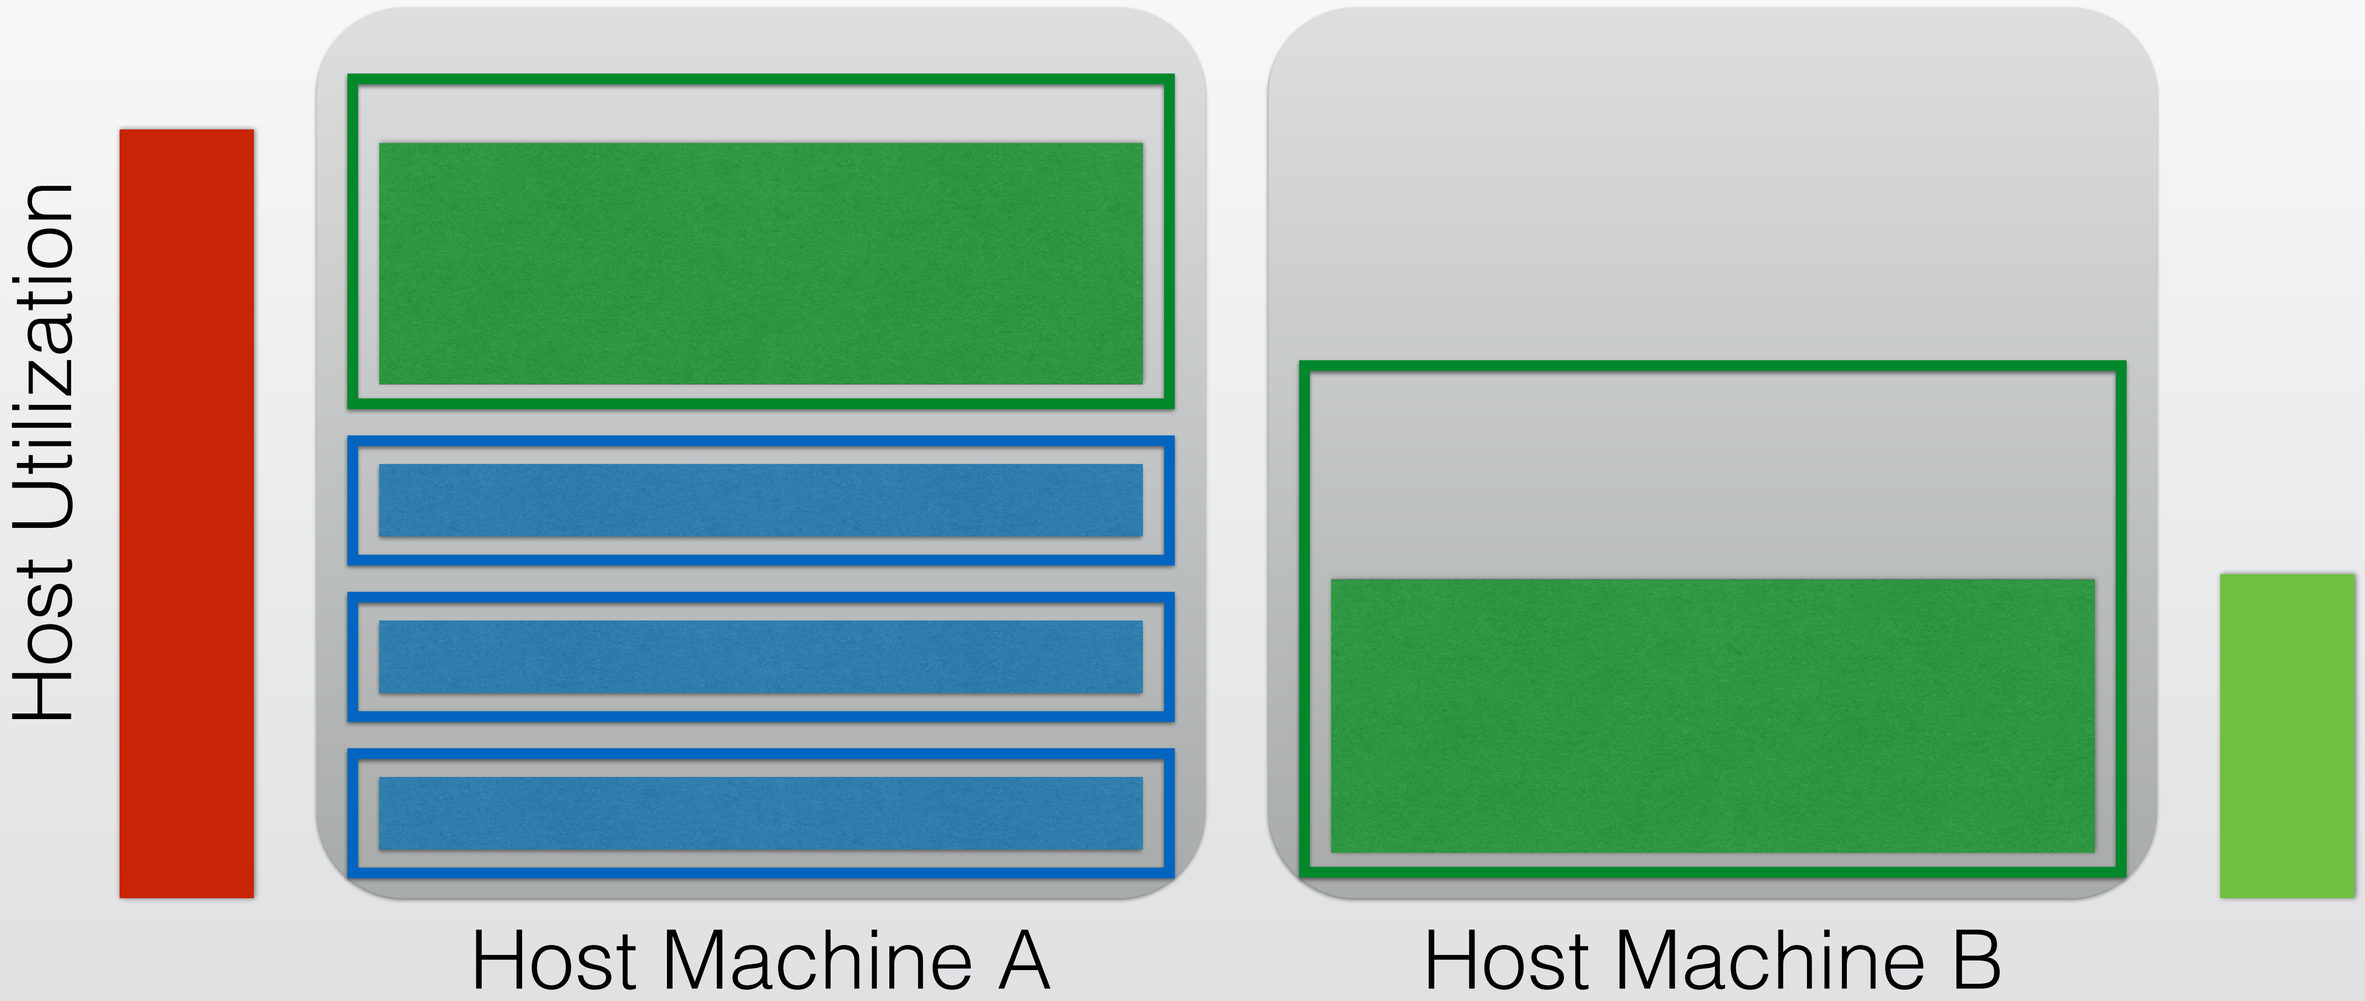
Host Utilization
Host Machine A
Host Machine B
Host Utilization

Scheduling

Unable to overcommit

Open for overcommit

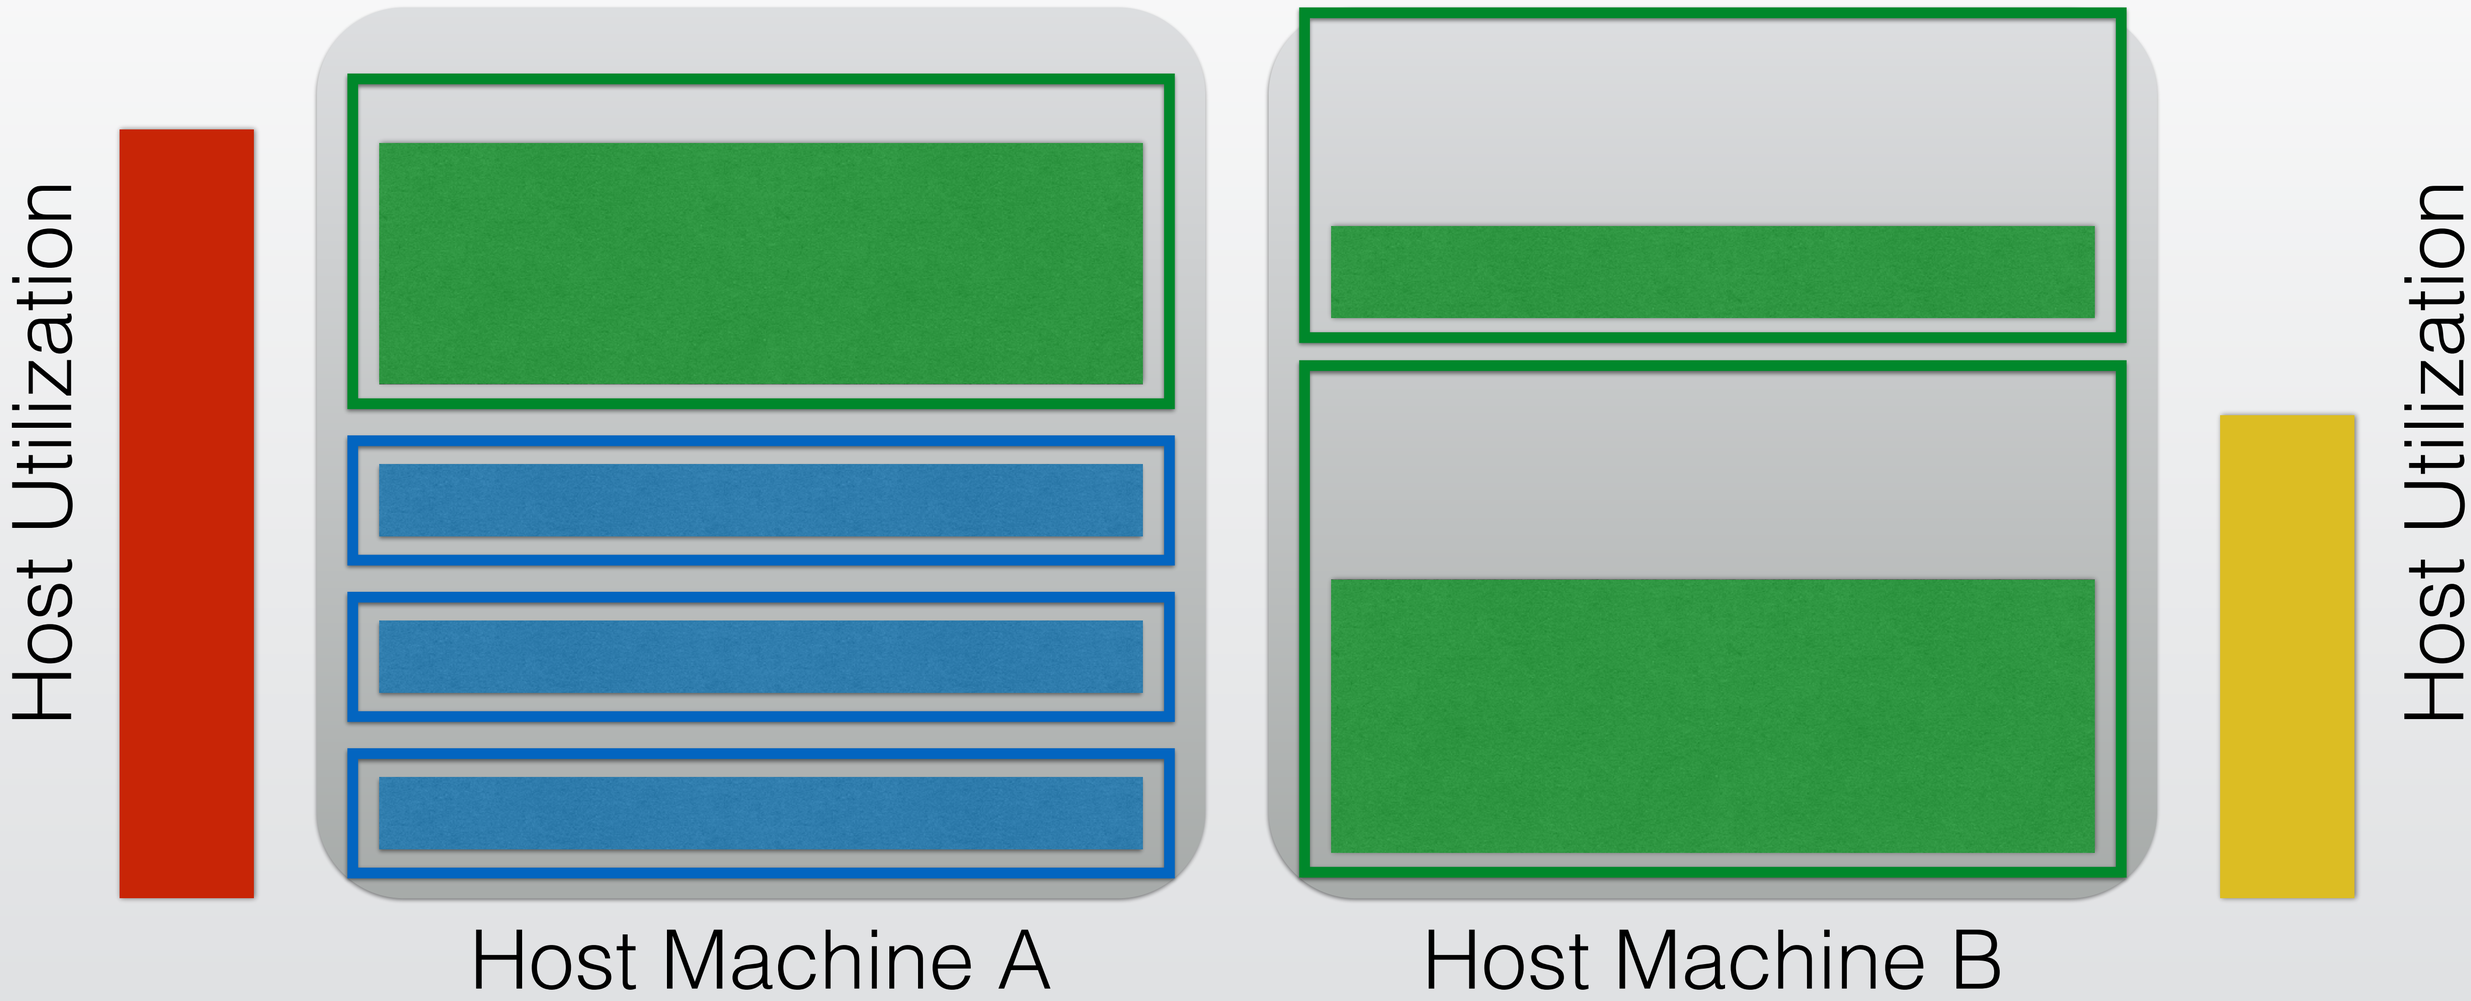

Scheduling + Rebalancing

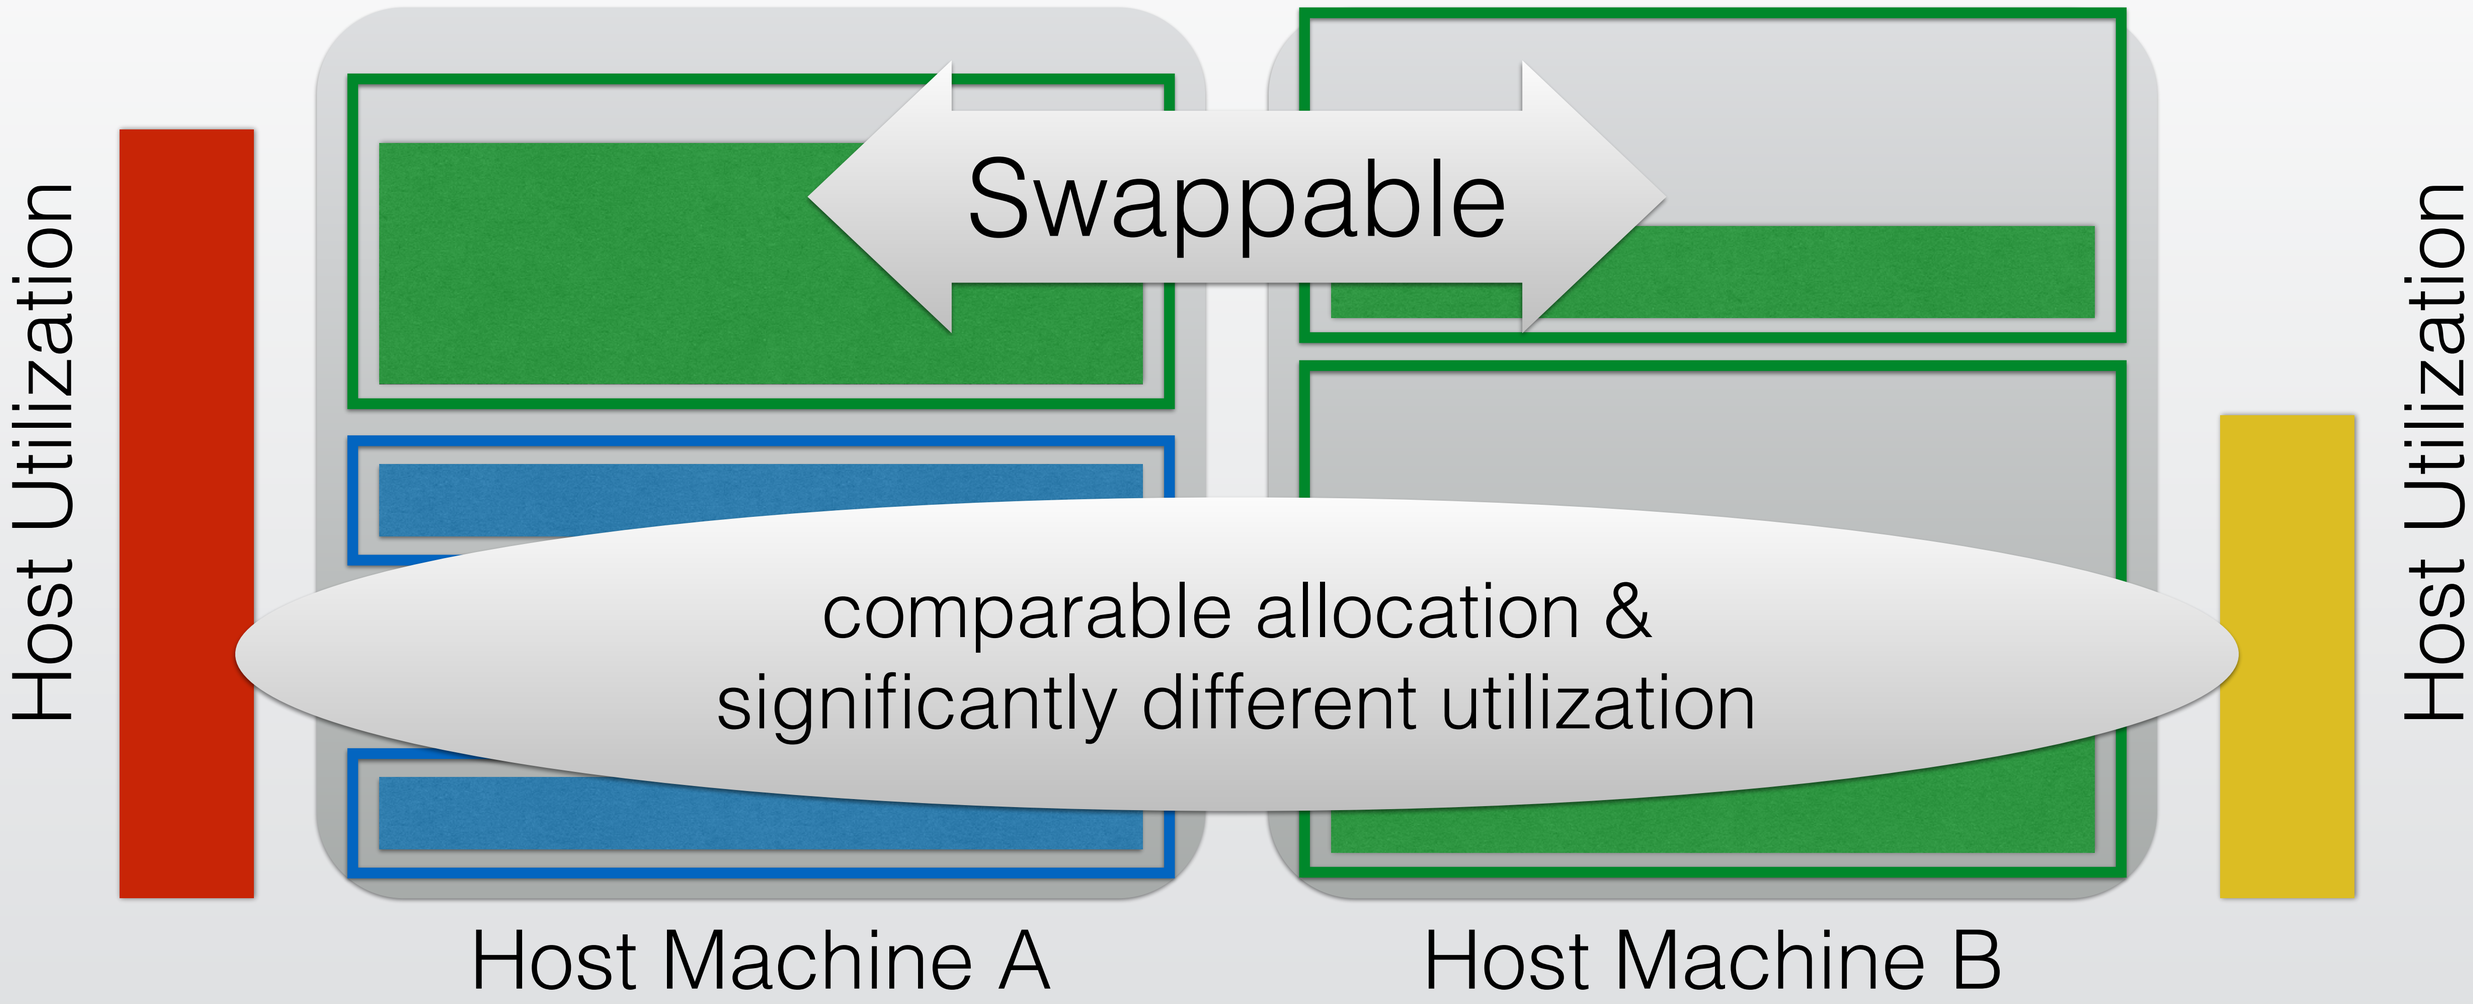

Scheduling + Rebalancing

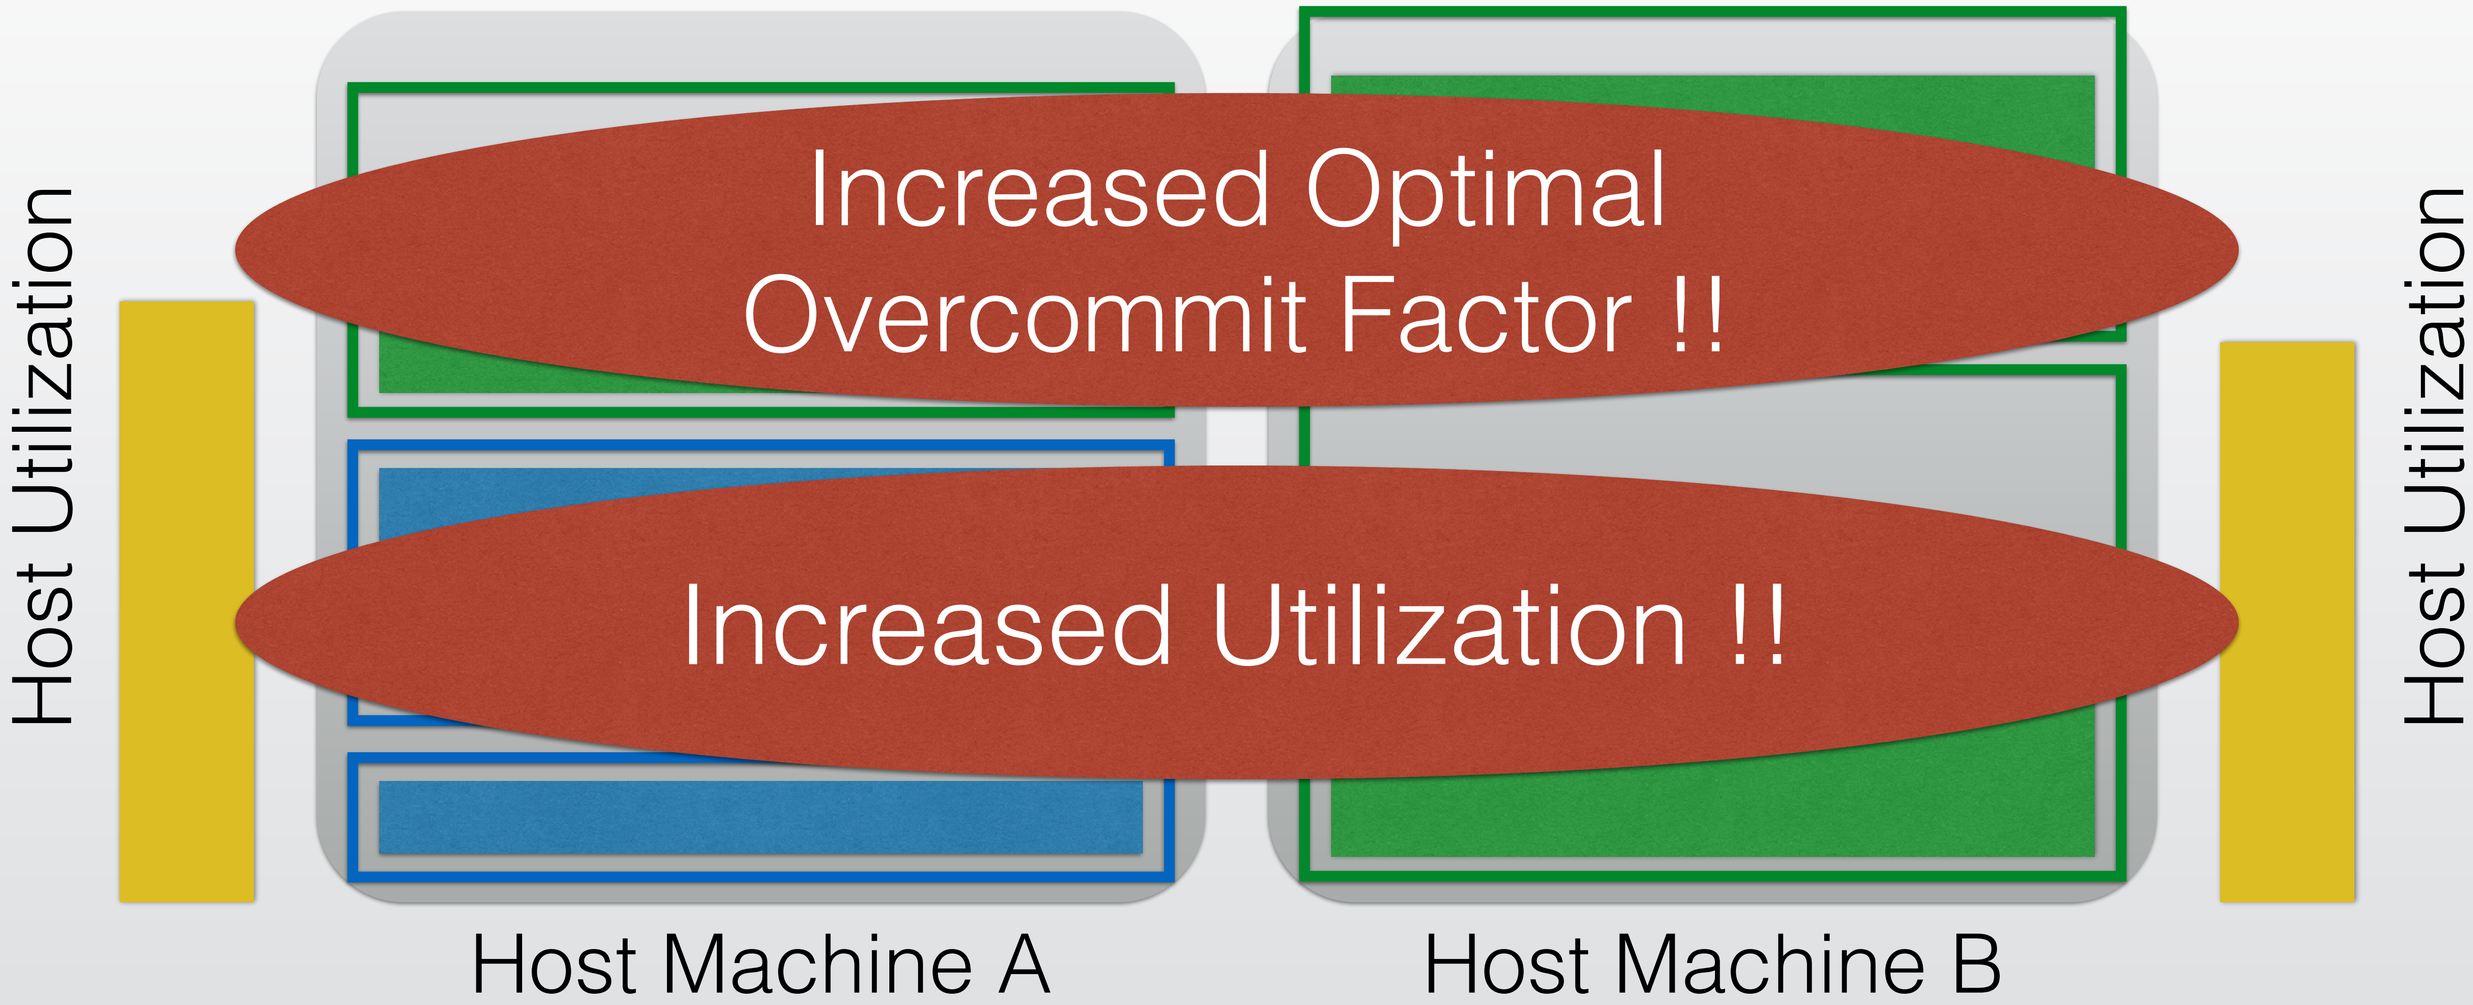

Evaluation

- Simulate cluster with/without task rebalancing
- Driven by Google's cluster data traces *,**
- Measure performances
  - Metrics: cpu utilization, # of completed tasks, etc.

* J. Wilkes, *"More Google cluster data."* Nov-2011.

** C. Reiss, J. Wilkes, and J. L. Hellerstein, *"Google cluster-usage traces: format + schema,"* Mountain View, CA, USA, Nov. 2011.

Preliminary Results

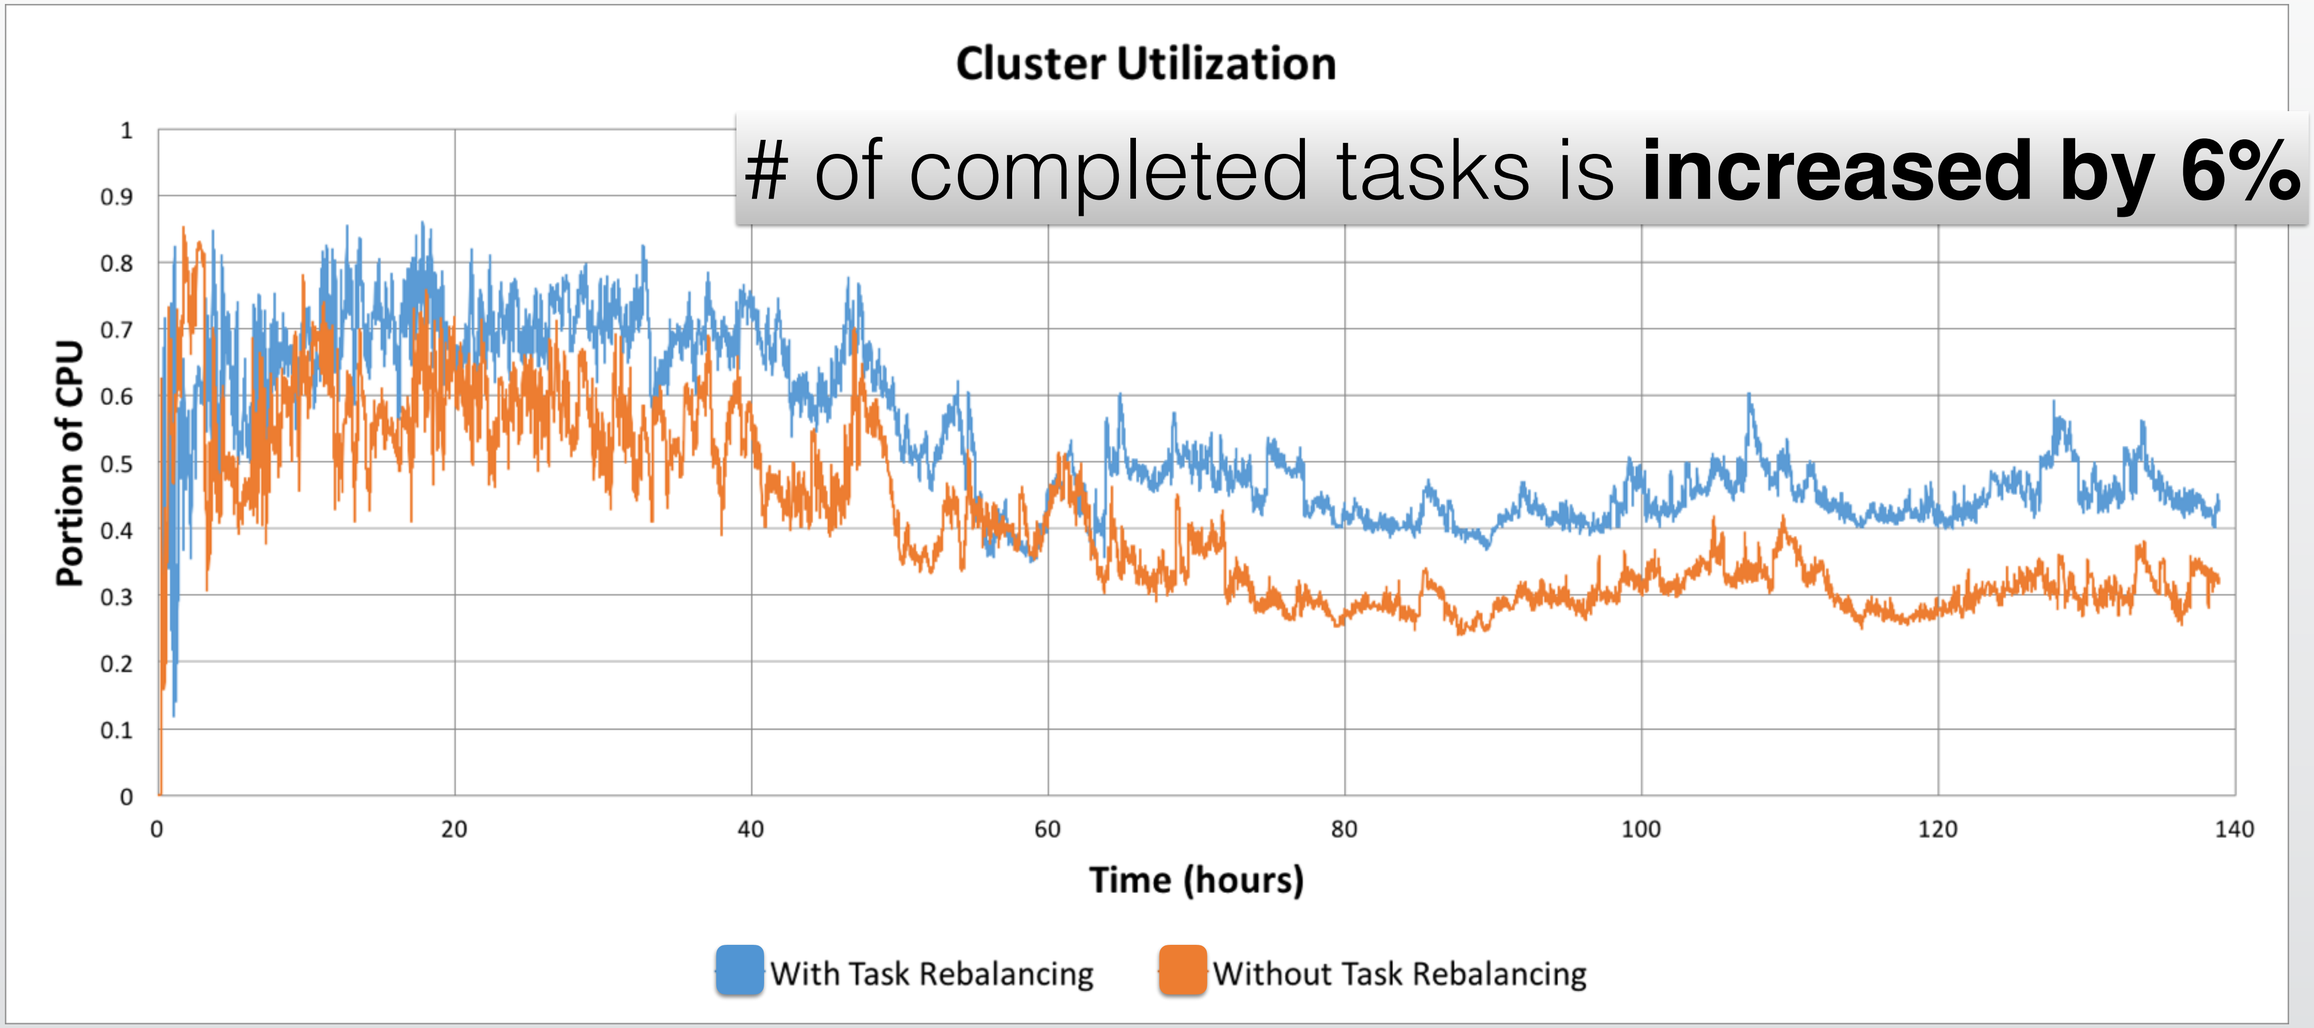

Conclusion

- PROBLEM: Scheduling is inefficient
- INSIGHT: Linux container enables task micromanagement
- PROPOSAL: Task rebalancing
  - Find and swap *swappable task pair* to load-balance hosts
  - Increase optimal overcommit factor => Increase cluster utilization
- EVALUATION: Simulation driven by Google's cluster data