

Resources Working Group Update

PRAGMA 30

Manila, Philippines

istributed Resources

How do we make a worldwide collection of resources usable for sharing of data/computation?

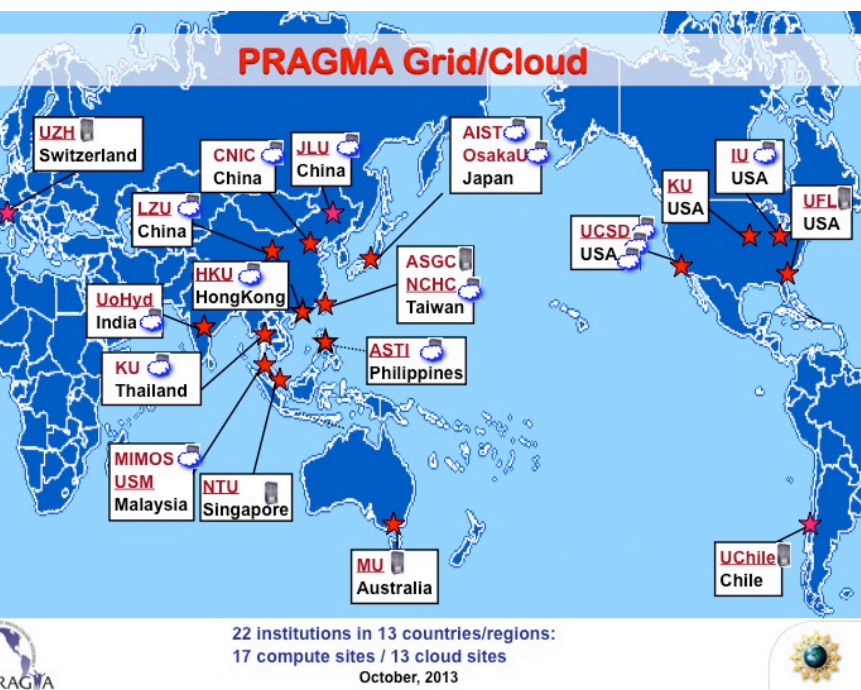
Drive infrastructure by needs of specific use cases

- Biodiversity expedition
- Lake Ecology Expedition
- Experimental Networking Testbed
- PRAGMA RDA (Research Data Alliance) Sprint

Key update on Cloud Scheduler (Demo)

Updated disk image handling dramatically improves boot times

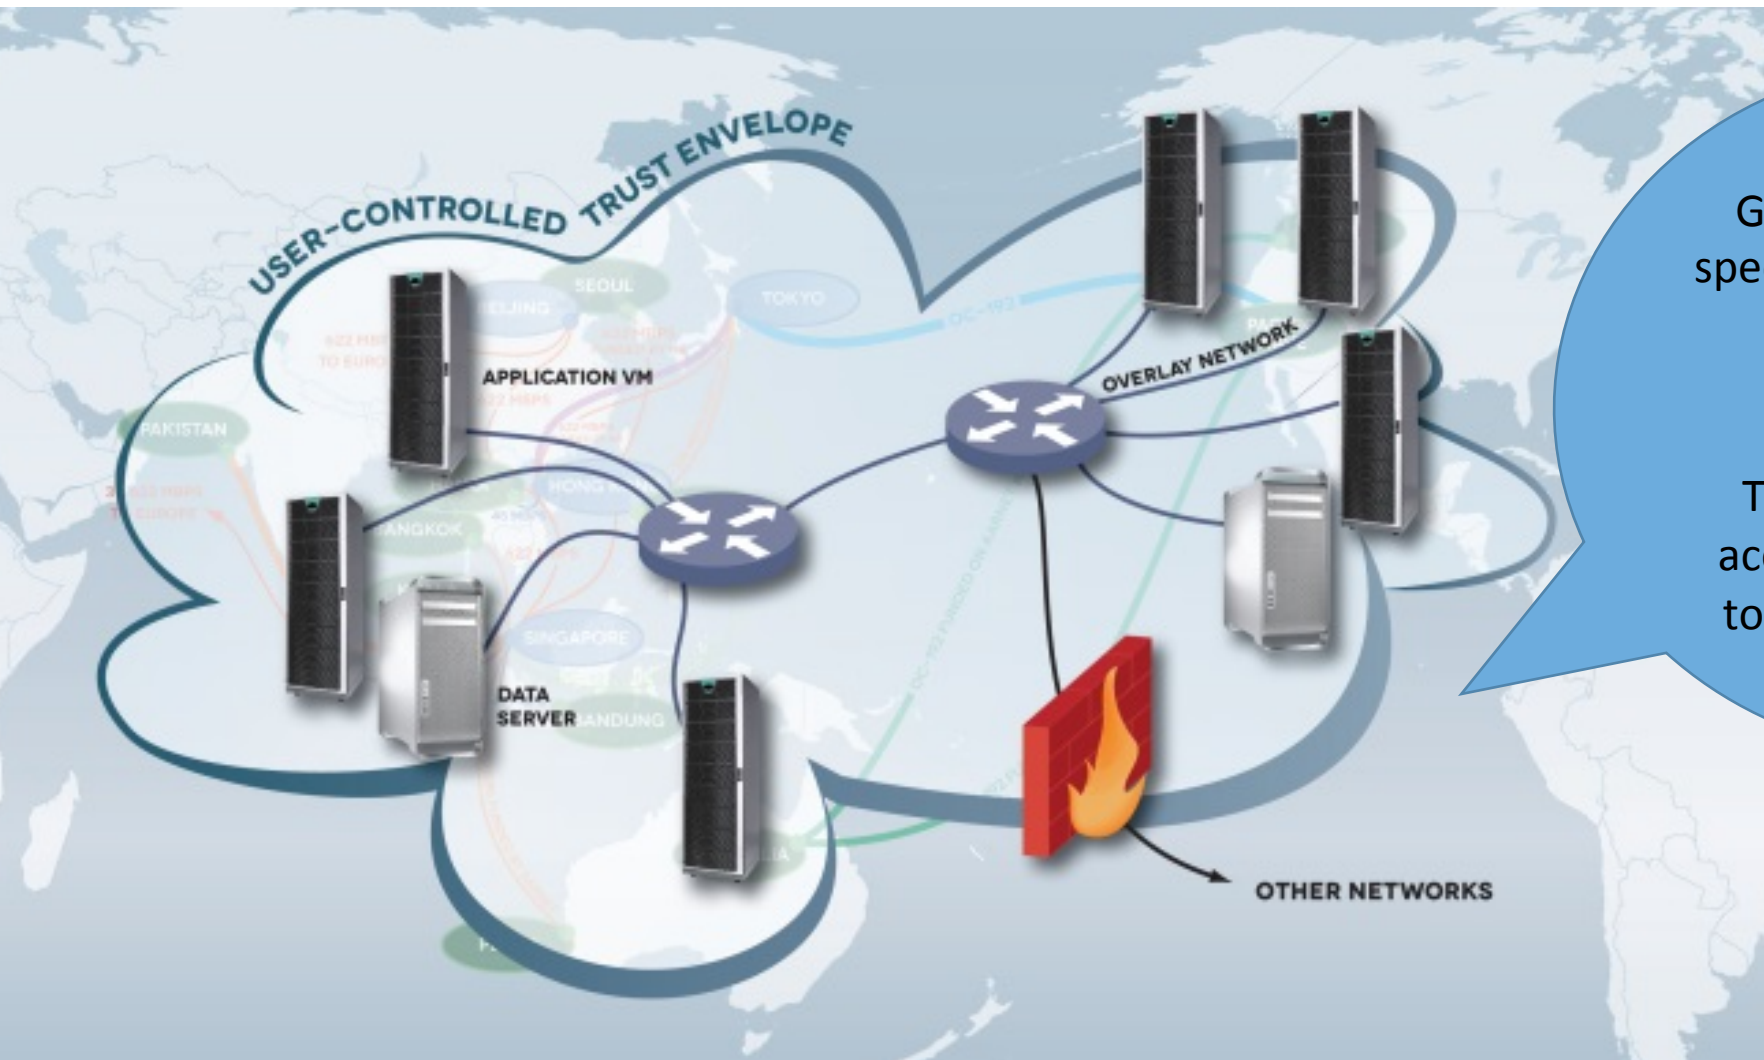
Progress goal: Rebuilding the PRAGMA Globalwide Testing Infrastructure



Next
Generation

Trying to rebuild infrastructure AND make scientific progress

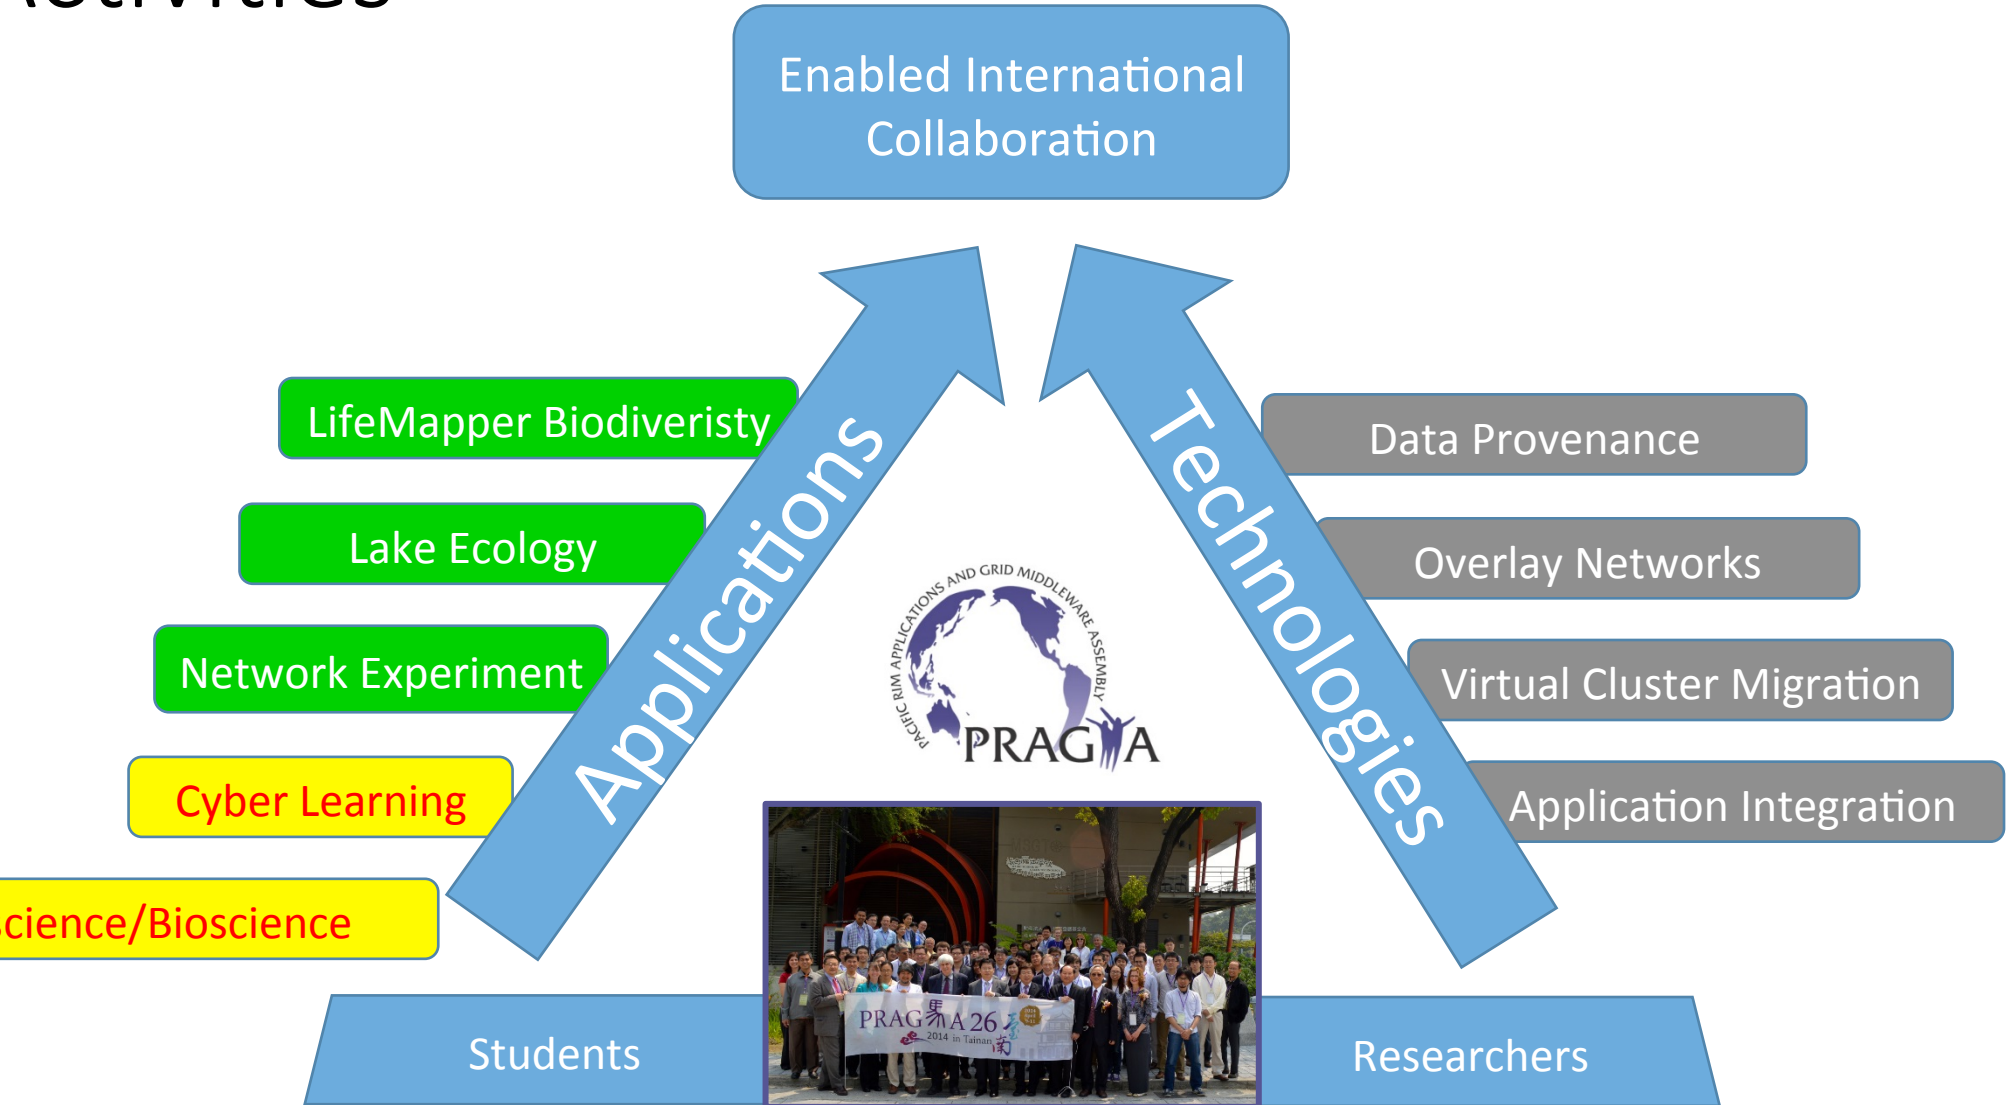
We need controlled sharing for resources



General Idea: identify specific resources needed by distributed researchers.

Then build a "limited access" network to keep together the resources.

Where Resources fits in for PRAGMA Activities



Resources Working Group Update

Lifemapper PRAGMA 30 Update

KU infrastructure

- Moving towards enabling virtual clusters for beta and production systems

SDSC infrastructure

- Updated fiji PRAGMA-dedicated cluster to Rocks 6.2
- PRAGMA Boot/Image Storage capable

Lifemapper code

- Start enabling Lifemapper configuration and all setup/test invocations as a command line infrastructure. Simplify commands so that install, configuration, and testing are unified and extensible. Example:
 - `lm init db`
 - `lm star/stop pipeline`
 - `lm list users`
 - `lm update ip`
- Bug fixes
- Work on adding solr (Apache search) (for indexing)
- Continue work on formalizing requirements and code for fully described data allowing easy use of different input datasets

Lifemapper as part of the RDA Sprint

RDA

- Discussions about RDA framework to determine the trustworthiness of Lifemapper projection set through the mutation of occurrence set and Lifemapper VMs
- Identify metadata for LM objects (projections sets, occurrences sets, VMs)
- Set up 2 VMs for the demo:
 - **rocks-204** with SEA dataset and LM v.1.0.3
 - **pc-170** – use pragma_boot to instantiate rocks-204 as a new VM in a different network

Both VMs have metadata and PIDs that can be accessed via

<http://rocks-204.sdsc.edu/lm.json>

<http://pc-170.calit2.optiouter.net/lm.json>

Occurrence sets and projections are available on both VMs at <http://<ip>/services>

IPOP/Lake expedition Updates – PRAGMA-30

Renato Figueiredo, Ken Subratie, Kyuho Jeong, Paul Hanson,
Cayelan Carey, Kohei Ichikawa

U. Florida, U. Wisconsin, Virginia Tech, NAIST/UCSD

Technology development

IPOP (Easy to Use VPN) Switch mode

- Handle L2 broadcasts, ARP
- SDN-based overlay bypass for fast communication within network

Revamped IPOP controller framework

- Structured P2P GroupVPN
 - Chord-based P2P self-organizing topology and routing
 - On-demand IPOP links based on traffic inspection
 - Bootstrap nodes from XMPP

PRAGMA deployments/demos

More details: demo/presentations, Friday

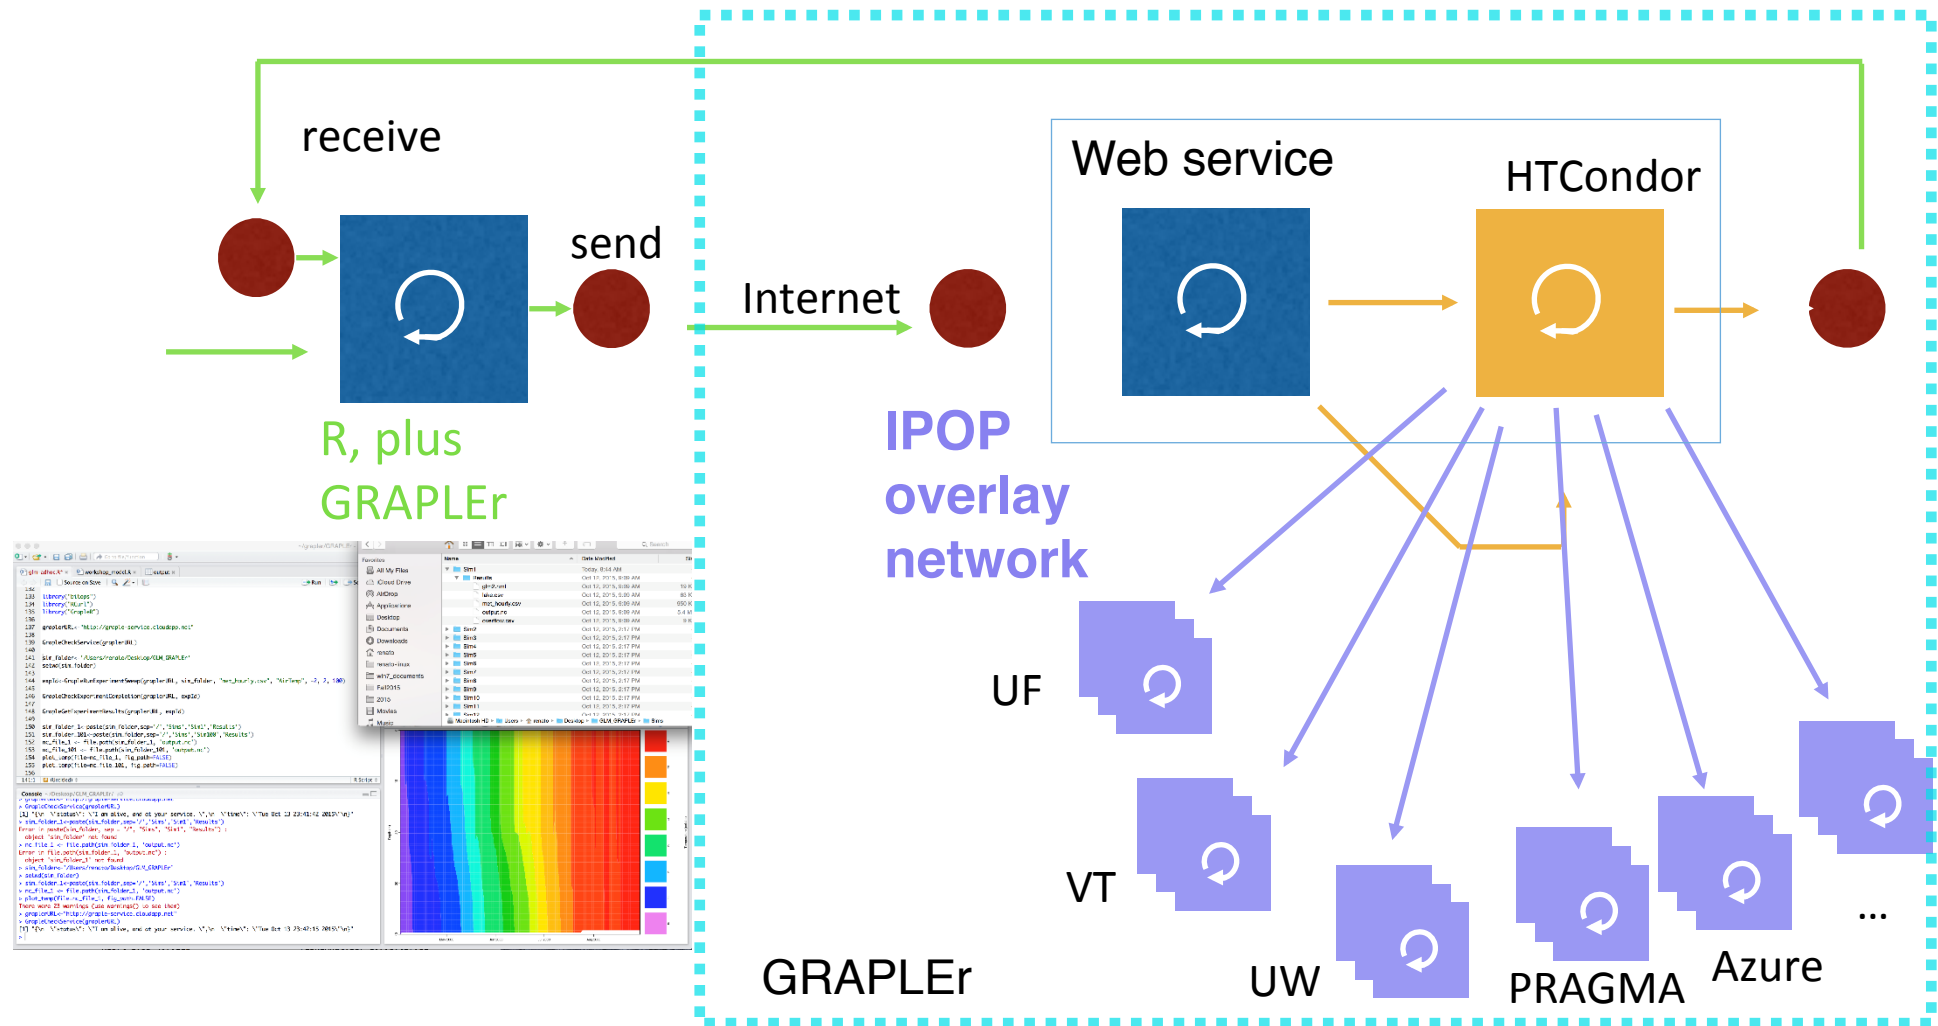
Extending PRAGMA-ENT with IPOP

- Collaboration with NAIST/UCSD (Ichikawa)
- Provide access to testbed through IPOP L2 mode

GRAPLEr

- Collaboration with UWisc (Hanson), VT (Carey)
- Web service + IPOP + HTCondor pool for lake modeling (GLM)
 - Development and use of R interface and Web service
 - Azure+UF; startup allocation in Comet/SDSC – will evaluate Virtual Cluster capability

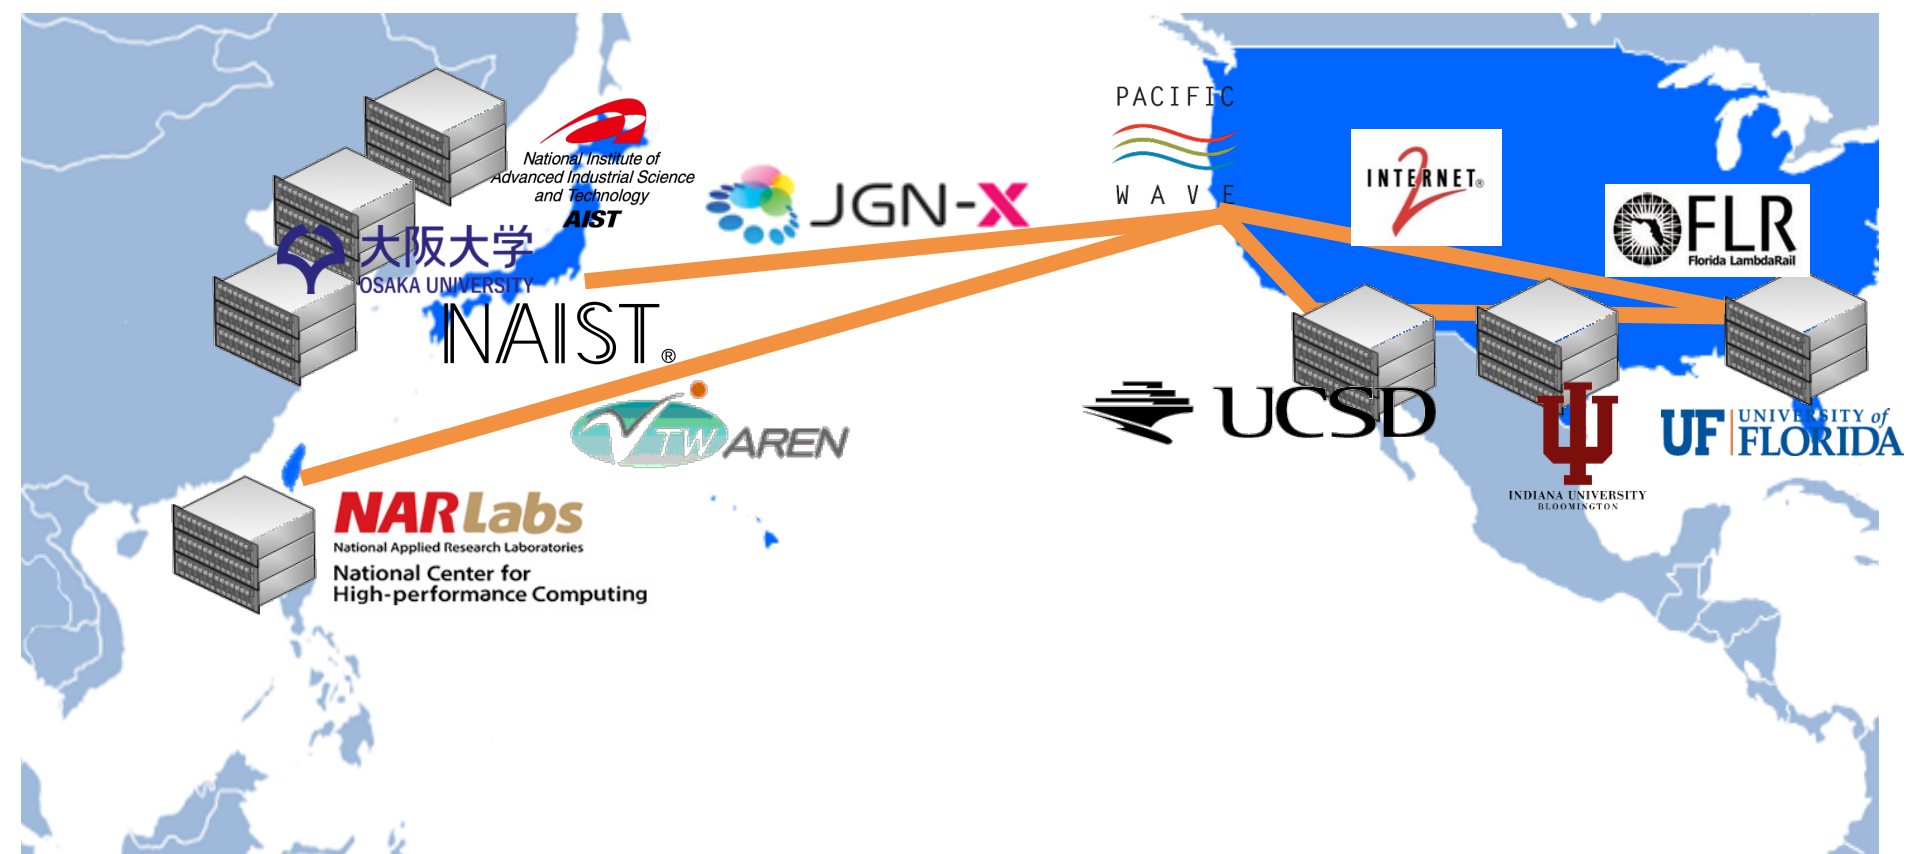
GLM through GRAPLER



PRAGMA-ENT (Experimental Network Testbed)

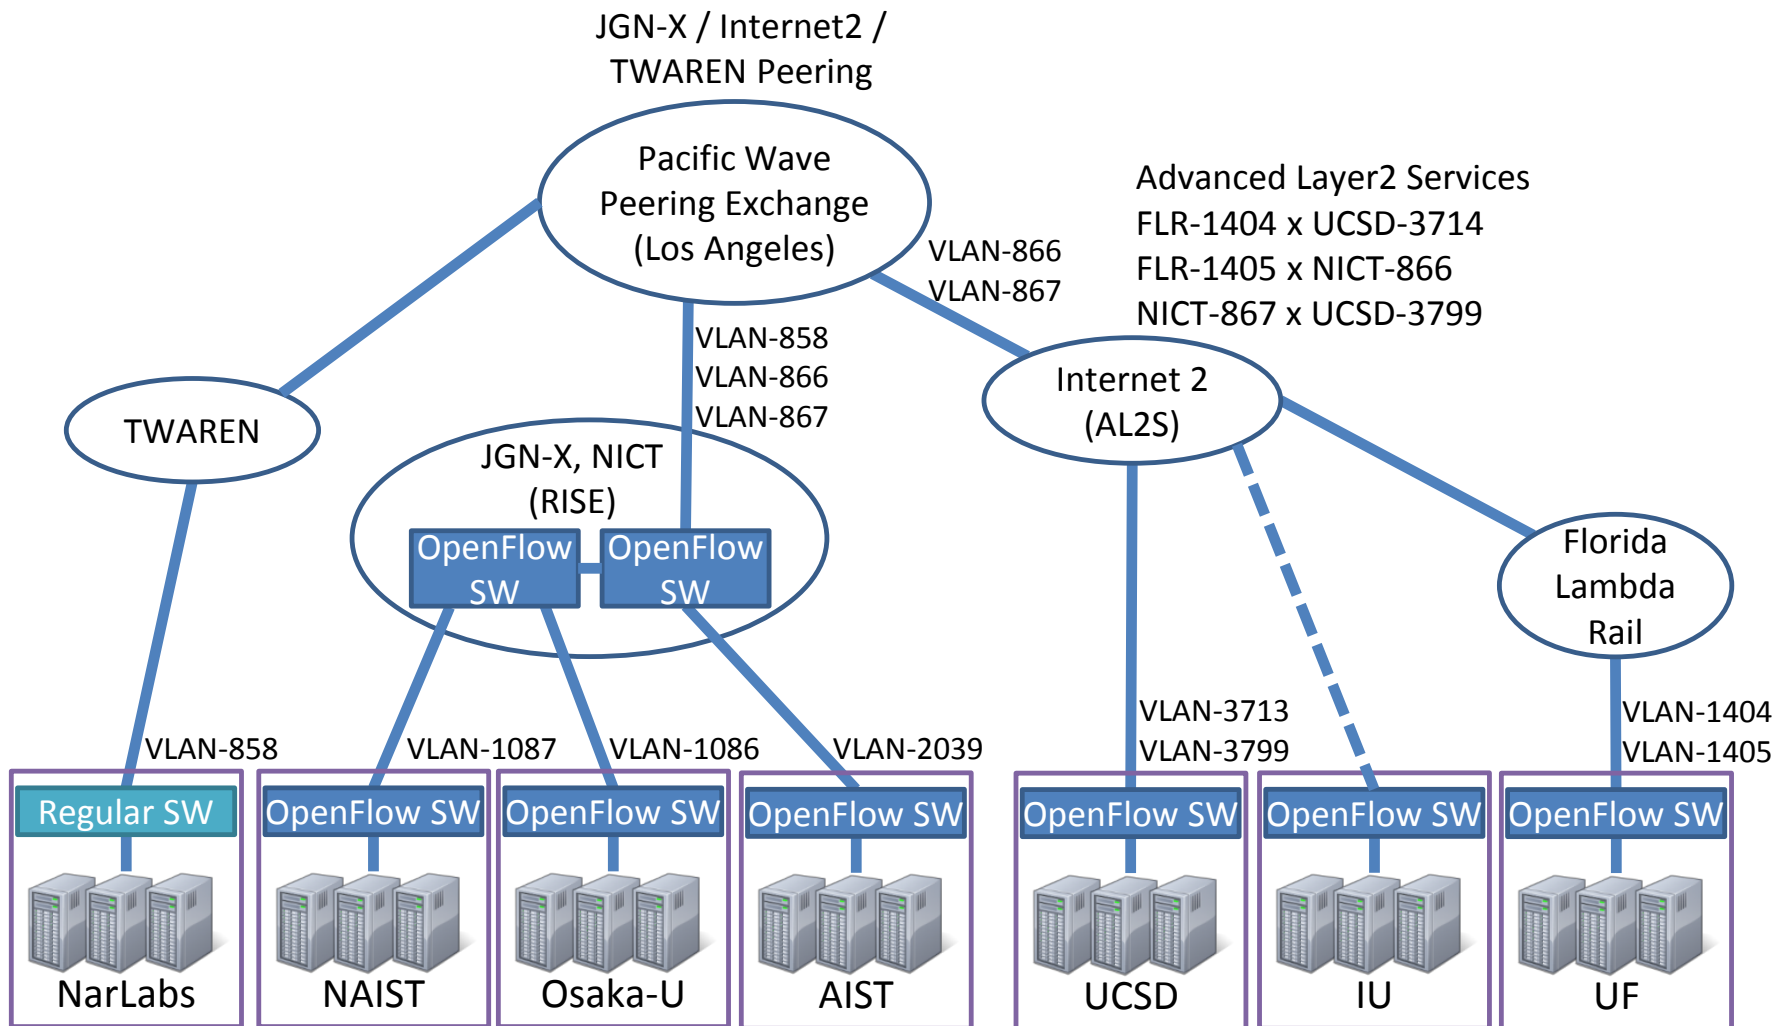
An international SDN/OpenFlow testbed for use by PRAGMA researchers and collaborators

- provides complete freedom to access and configure network resources

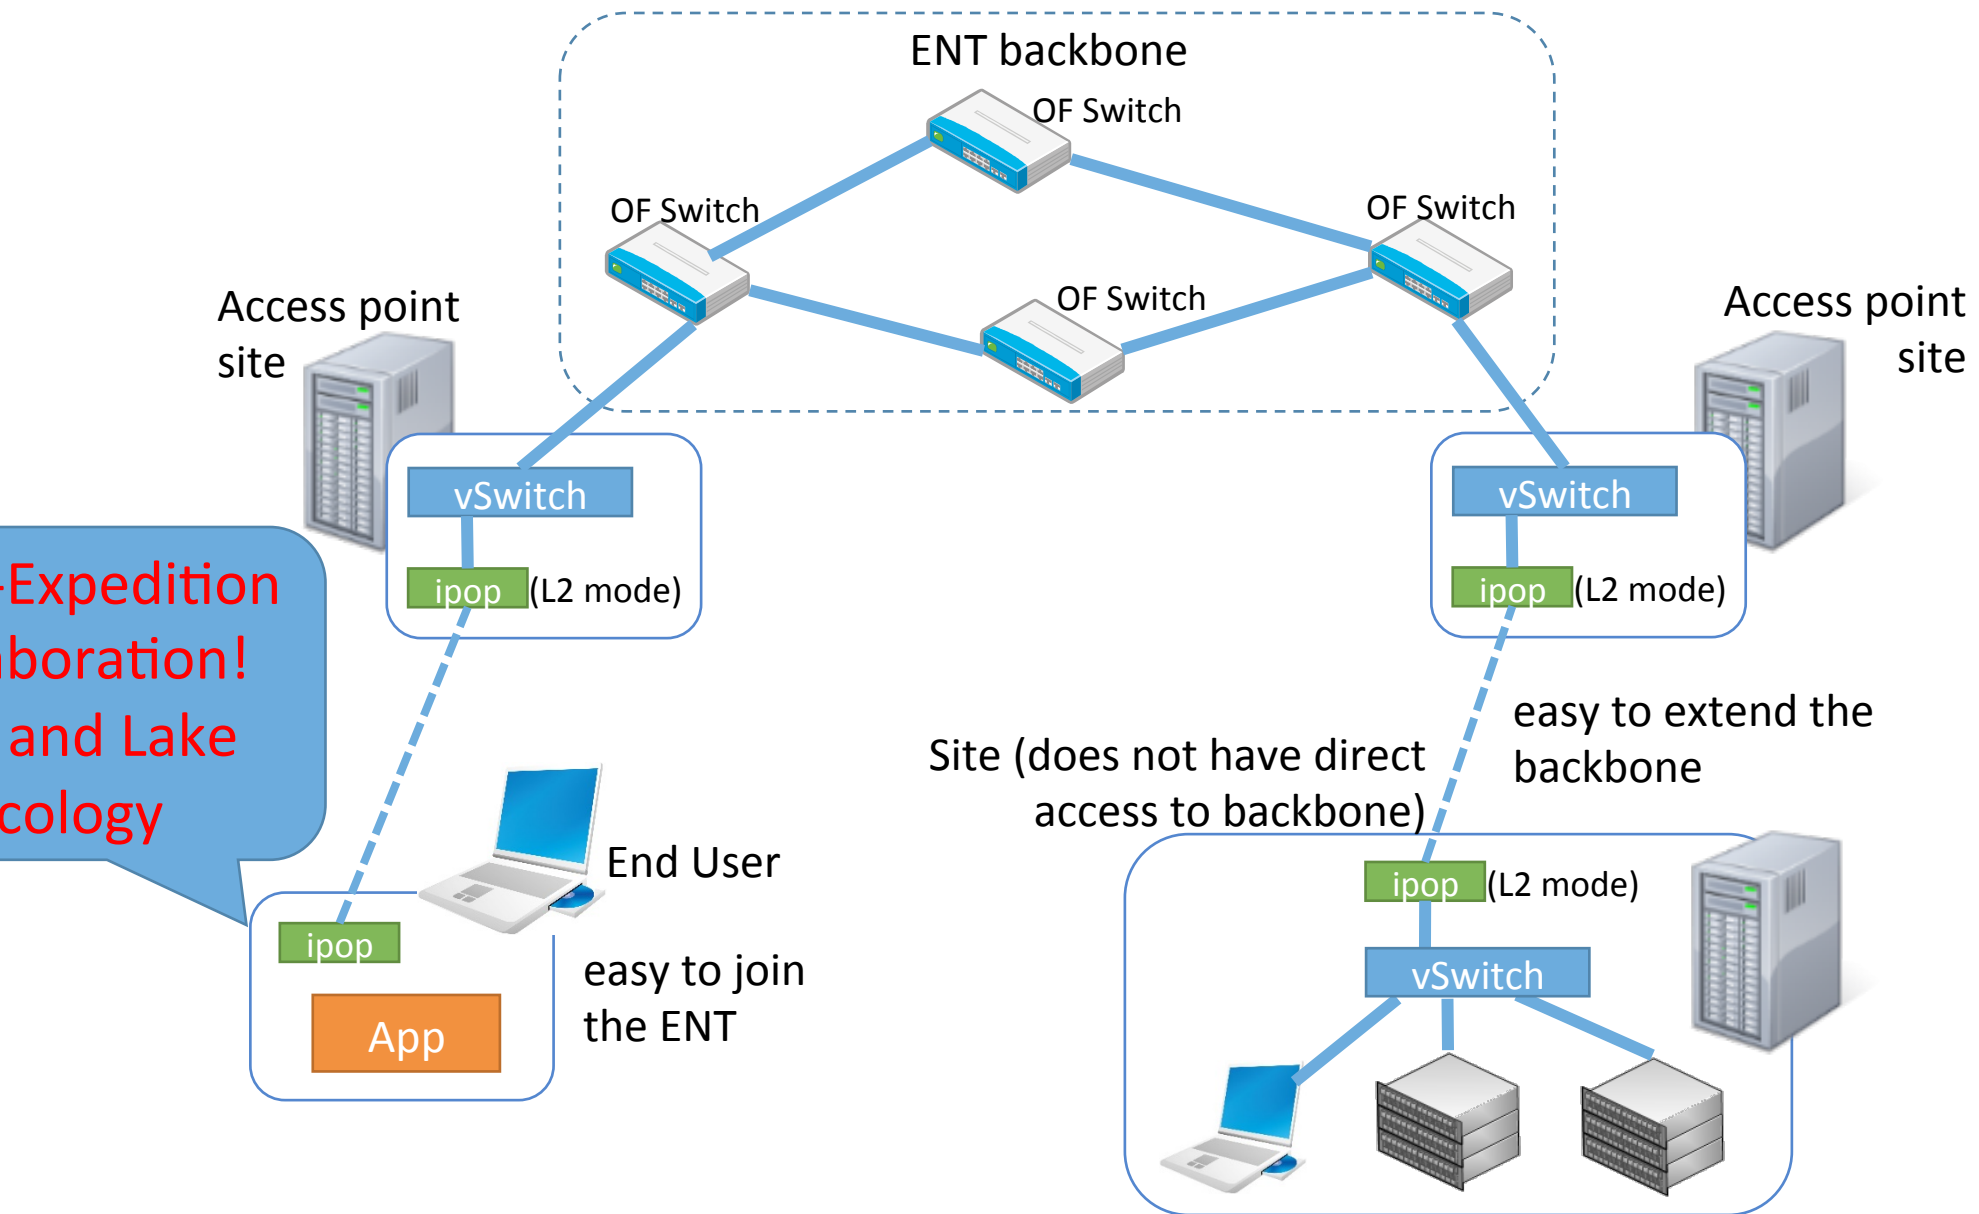


Backbone of PRAGMA-ENT

Currently, seven institutions are involved in the project. Their resources are connected each other through academic network backbones



Activities: Extending backbone through IPOP (demo)



Activities: Usability study of ENT

Based on usability study, we have been re-organizing the PRAGMA-ENT documents on github wiki

- https://github.com/pragmagrid/pragma_ent/wiki

Documentation Then

Official PicOS documents – 34 pages

vSwitch documents – 257 pages

Single user thread

Documentation Now

PicOS documents for ENT – 4 pages

vSwitch documents for ENT – 5 pages

Multiple user thread

- Application user
- Infrastructure user

RDA-PRAGMA Sprint -- PRAGMA 30 @ Philippines

Quan (Gabriel) Zhou, Nadya Williams, Aimee Stewart
Jason Haga, Beth Plale

2/4/16

Objectives

- Assess recently released tools and best practices from RDA for contribution to PRAGMA services. Carry out assessment through 2 phase demo.
- Demo: verify lineage of projection data objects, and enable rerun when new data exists
- Enhancements to PRAGMA testbed: Provide common persistent identifiers and landing pages to VMs and datasets of Lifemapper
- Feed results back to RDA (Research Data Alliance)

Demo Phases

Phase I: Use static GBIF subset for Southeast Asia as input to Lifemapper,

- Input datasets bundled into VM.
- User has ID of two projection result datasets (both result sets have same internal ID (e.g., 317)), and uses RDA services to determine whether they came from the same VM or from the primary VM and its clone

Phase II:

- Input datasets ingested dynamically into VM (workflow dynamically accesses iDigBio.)
- After seeing change to iDigBio input dataset, use new PRAGMA data infrastructure to identify, download, and faithfully replay run with new iDigBio input dataset to visually compare before and after.

**Adding a Significant Reproducibility
Dimension to Lifemapper
Experiments**

New architectural components

PRAGMA-RDA Data Service

- Stores metadata, objects, and landing pages
- Maintains metadata about both data sets and VMs
- Assigns unique handle to incoming objects
- Displays landing page for each object
- Interacts with RDA PIT/DTR service
 - PID – Persistent Identifier, PIT = PID Information Types, DTR – Data Type Registry

RDA PIT/DTR service: stores type information about minimal metadata that allows interpretation of the metadata.

Handle service : obtain handle PID

Landing Page

URL:
<https://hdl.handle.net/11723/5e9-54f8-4c5c-9e95-054a26c>

base example:
<https://github.com/Gabriel-Zhou/landingpage>




RDA-PRAGMA Landing Page for Lifemapper

| | |
|-------------------|--|
| ProjectionSet ID | 7575077 |
| Display Name | Carex pensylvanica |
| Scenario Code | WC-10min |
| Bounding Box | -180.0, -60.0, 180.0, 90.0 |
| Resolution | 0.16667 |
| Last Modified | 2015-12-16 20:32:54 |
| OccurrenceSet PID | 11723/374e7cff-2195-42d5-b48d-311d0e65cc66 |
| Experiment ID | 156119 |
| Checksum | f540c8cc528596967fde3c9925e140c9 |

[Download](#)[Go to Occurrence Set](#)

Projection Set Provenance



Checksum Information Type

PID: 11723/377739b4-14df-441a-b219-15881cf6ae52

Checksum

Type: dataType



Digital Object View JSON View **Versions View** Show Relationships

Identifier

20.5000.239/d8fcd1cd020581d6d23f

Type Name *

Checksum

Description *

A property that holds a checksum String for a digital object.

Provenance

Contributors of this Record

| Identified Using * | Name * | Details |
|--------------------|-----------|--------------------------------|
| Handle | Quan Zhou | Indiana University Bloomington |

Creation Date

2015-11-15T02:56:31.628Z

Last Modification Date

2015-11-15T02:56:31.643Z

Expected Uses

Use *

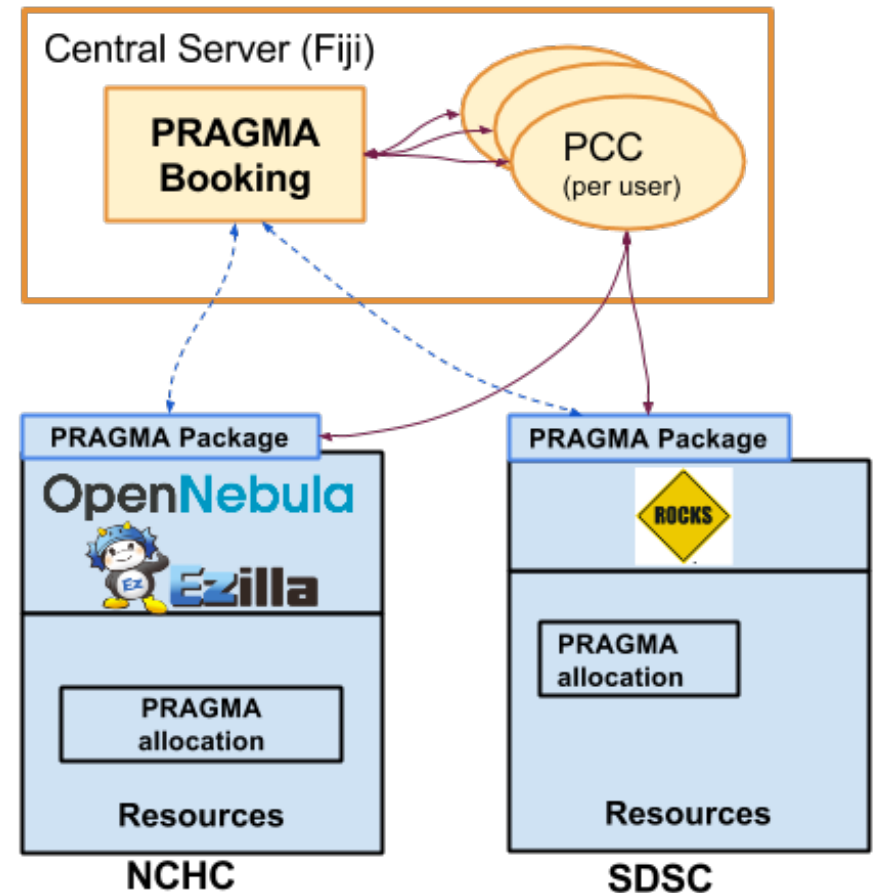
A property that holds a checksum String for a digital object.

U Data Node Resources – Pathfinder for Data-centric nodes in PRAGMA

- RDA PIT/DataType Registry Service
- PRAGMA-RDA Data Service
- Handle V8 service (generates PIDs, runs at CNRI)
- PRAGMA-ENT Mesh
- Open HathiTrust Corpus

Cloud Scheduler Updates

- **Goal: Provide a lightweight cloud scheduler for sharing of PRAGMA resources**
- **New:** Migration from prototype to (very) early users
 - Setup Fiji (UCSD/SDSC) as our production cloud scheduler server
 - Initial integration of NCHC Ezilla deployment using OpenNebula pragma_boot driver
 - Added automated shutdown of virtual clusters



Pragma_boot updates

- Goal: Boot virtual clusters for users across PRAGMA institutions using local VM provisioner
 - **New:** Rewrote pragma_boot framework Python scripts and kvm_rocks driver (previously bash scripts)
 - **New:** Enhanced vc-out-parser to work for Ubuntu
 - **New:** Made virtual cluster images of the hku_biolinux.img from Biosciences WG (frontend and compute images) and leveraged Clonezilla to create sparsified ZFS volumes

Cross – working
Group interaction

Using Tech from
NCHC Taiwan

Practically dealing with Large disk images

Previous versions of PRAGMA boot had to fully copy Frontend and Compute Node images to remote hosts BEFORE booting

→ Could take 30 – 60 minutes to boot a virtual cluster

OSC Developed tech –

• Boot a virtual machine with its system disk homed on a central disk server
• Migrate image to remote. Then over iSCSI link → VM disk access become local

Adapted for PRAGMA to support cloning of compute node images

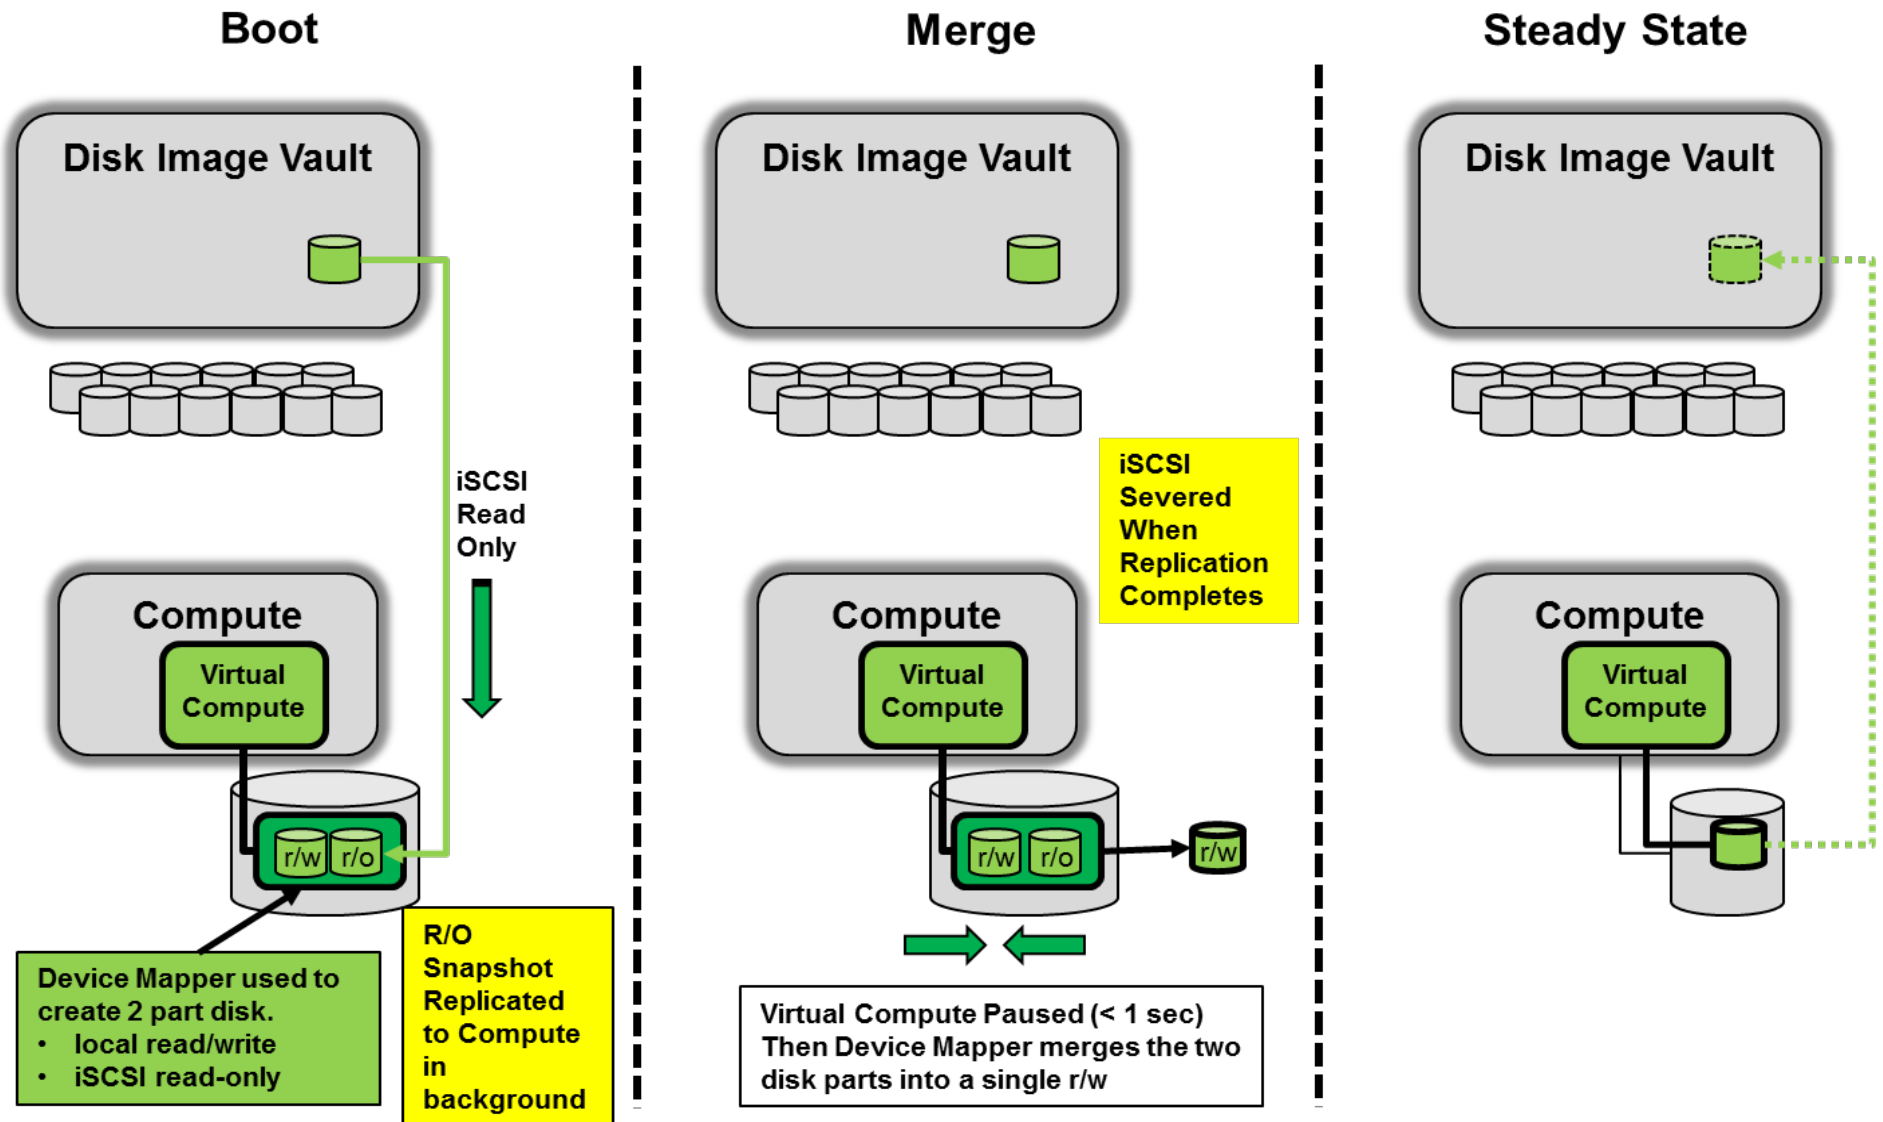
→ Frontend + 2 Compute node boot in less than 90 seconds

```
-01-27 03:08:03,026 INFO Loading driver pragma.drivers.kvm_rocks.Driver
```

```
-01-27 03:08:30,320 DEBUG Executing command: '/opt/rocks/bin/rocks start l  
ocks-52'
```

```
-01-27 03:09:25,615 DEBUG Executing command: '/opt/rocks/bin/rocks start l  
m-rocks-52-1'
```

leveraging and expanding tech from SDSC



Advantages

Virtual Machines boot very quickly

- Disk Image on Storage Vault is replicated in background

On Shutdown, final disk state is returned to vault

All Disks are lazily replicated

- If physical compute node fails, state of virtual disk is close to “up to date”

Utilizes the parallel I/O busses of all the compute nodes

- Disk vault can be built using commodity components.

→ More Details in Working Group

Planned Topics for PRAGMA 30

Beth - RDA, where are we, what are the near term plans, mid-term goals, how others can get involved

Shava - How PRAGMA boot works, thoughts on using Clonezilla within the workflow.

Renato/Paul/Cayelan -- How to get others involved in the Lake Eco expedition. Contribute resources, help find domain collaborators

Kohei - PRAGMA ENT and work done to make it easier to add users. Detailed status.

Yoshio -- concrete steps to rebuild the PRAGMA cloud

Phil - Update on image storage used on comet (reducing dependency on Rocks)