Pacific Research Platform and The Pacific Rim

Presenter: Tom DeFanti, Research Scientist, QI, Co-PI

Larry Smarr, Calit2 Director and PI

Ilkay Altintas, Camille Crittenden, Ken Kreutz-Delgado, Phil Papadopoulos, Tajana Rosing, Frank Wuerthwein, co-PIs

John Graham, Senior Development Engineer

Dima Mishin, Isaac Nealey, Joel Polizzi, Mark Yashar, Programmers

UC San Diego and UC Berkeley

PACIFIC RESEARCH PLATFORM

it²

2015 Vision: The Pacific Research Platform will Connect Science DMZs Creating a Regional End-to-End Science-Driven Community Cyberinfrastructure

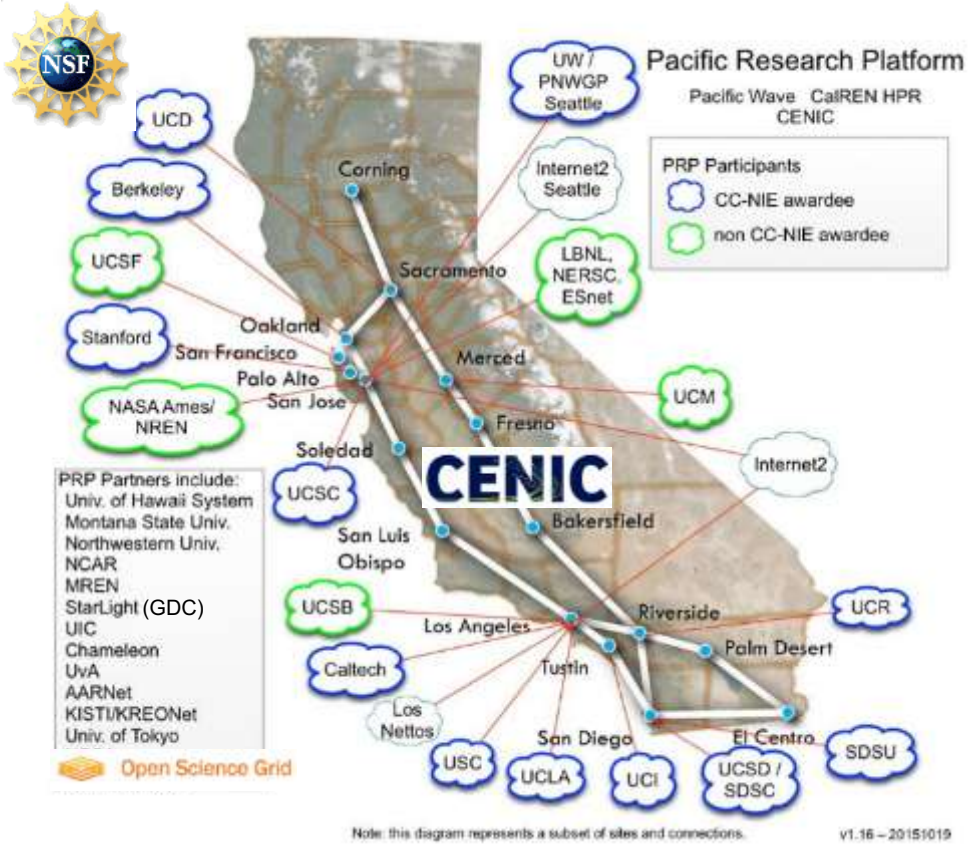

Source: John Hess, CENIC

NSF CC*DNI Grant
$6.3M 10/2015-10/2020
Year 5 Starts in 3 Weeks!

PI: Larry Smarr, UC San Diego Calit2
Co-PIs:

- **Camille Crittenden, UC Berkeley CITRIS,**
- **Tom DeFanti, UC San Diego Calit2/QI,**
- **Philip Papadopoulos, UCI**
- **Frank Wuerthwein, UCSD Physics and SDSC**

Letters of Commitment from:

- **50 Researchers from 15 Campuses**
- **32 IT/Network Organization Leaders**

ESnet: Given Fast Networks, Need DMZs and Fast/Tuned DTNs

PACIFIC RESEARCH PLATFORM

SDSC CITRIS BANATAO INSTITUTE it²

PRP Technical Deliverables
2015 - 2017

- **Phase 0: Tested Layer2 CENIC Networks and FIONAs—early 2015**
- **Phase 1: A Scalable Network for Optimizing Data Transfer—2015-2017**
  - **Layer 3 Data Transfer & Measurement Network**
  - **Tested, Debugged, Measured, Optimized, and MaDDash'ed Layer 3**
  - **Supported Rates up to 9.7/10, 37/40 Gb/s in 10GB Bursts**
  - **Included UvAmsterdam and Korea (KISTI)**
  - **Showed Full Bandwidth Utilization**
  - **Essentially No TCP Backoff on Long Distance Best-Efforts Networks**
- **This is What Most Other Research Platforms are Focusing on—Big Data Transfer**

PRP PACIFIC RESEARCH PLATFORM

it²

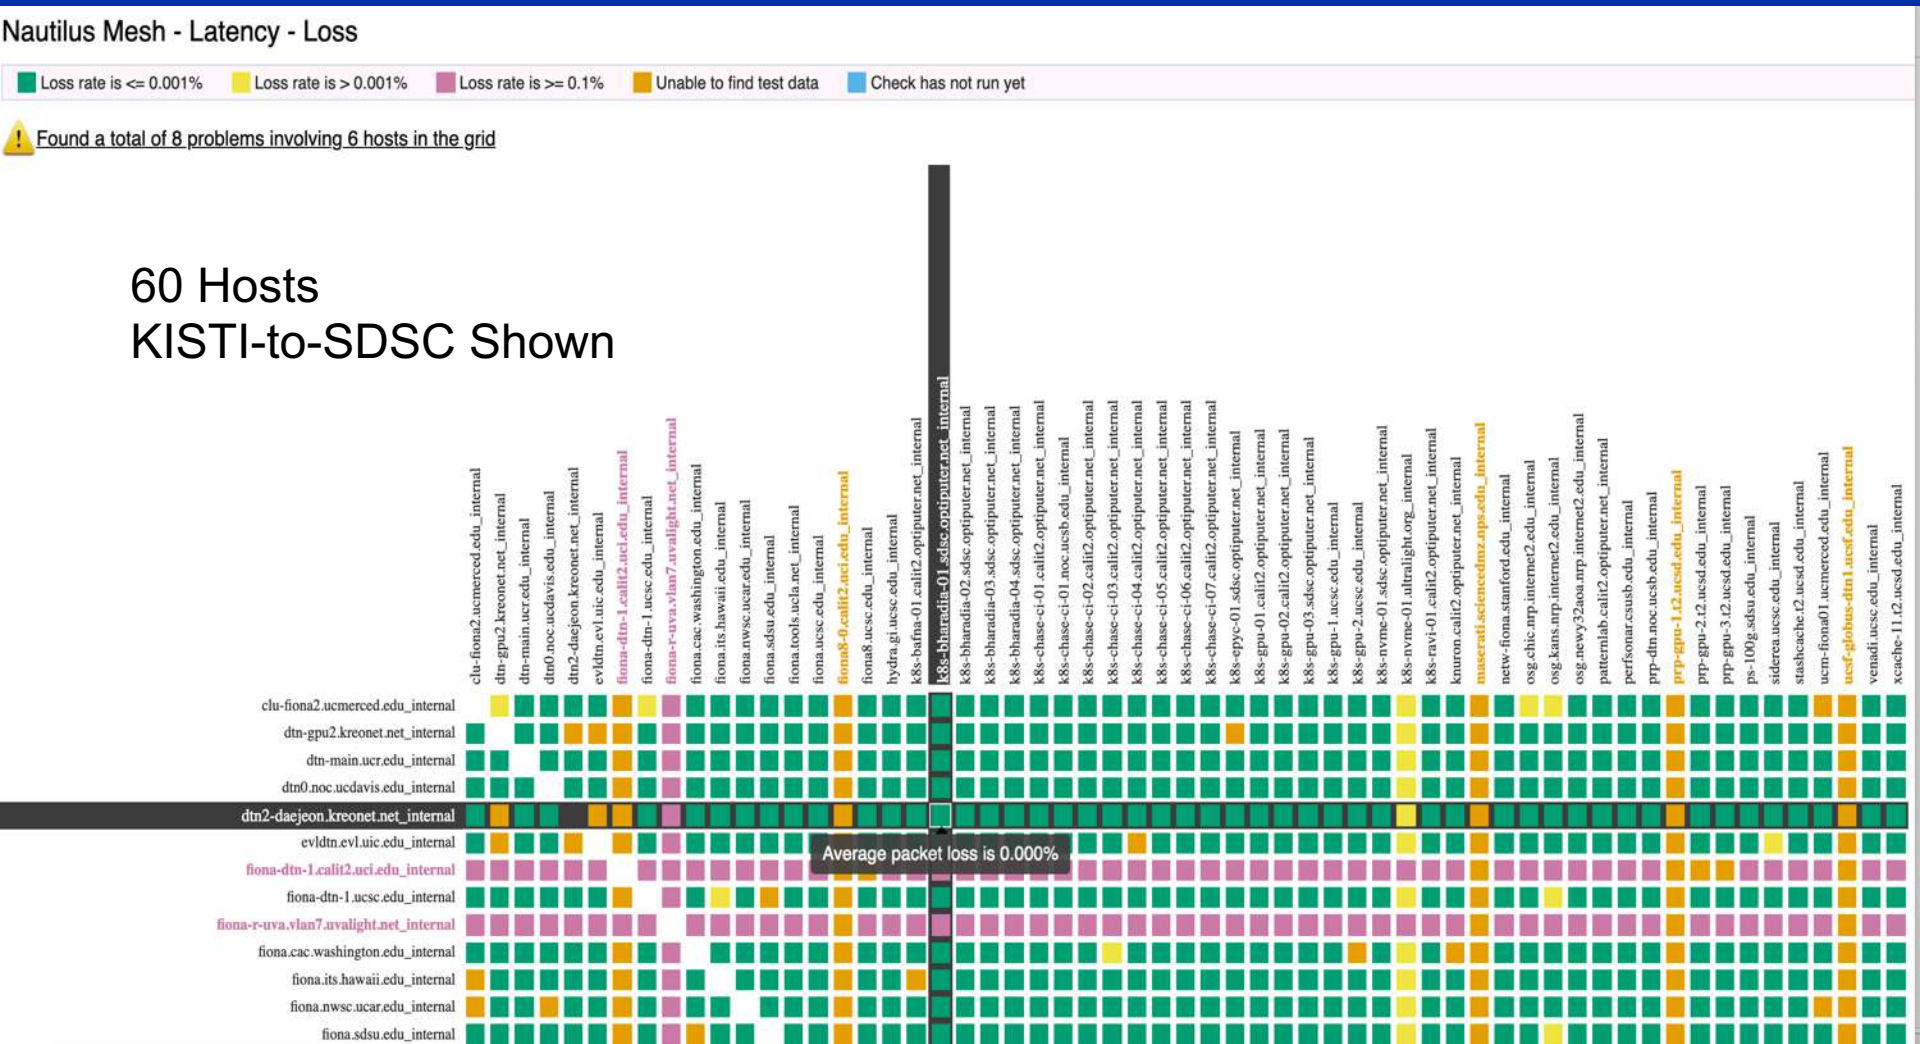
Nautilus Mesh - Latency - Loss
Loss rate is <= 0.001%
Loss rate is > 0.001%
Loss rate is >= 0.1%
Unable to find test data
Check has not run yet
Found a total of 8 problems involving 6 hosts in the grid
60 Hosts
KISTI-to-SDSC Shown
clu-fiona2.ucmerced.edu_internal
dtn-gpu2.kreonet.net_internal
dtn-main.ucr.edu_internal
dtn0.noc.ucdavis.edu_internal
dtn2-daejeon.kreonet.net_internal
evldtn.evl.uic.edu_internal
fiona-dtn-1.calit2.uci.edu_internal
fiona-dtn-1.ucsc.edu_internal
fiona-r-uva.vlan7.uvalight.net_internal
fiona.cac.washington.edu_internal
fiona.its.hawaii.edu_internal
fiona.nwsc.ucar.edu_internal
fiona.sdsu.edu_internal
fiona.tools.ucla.net_internal
fiona.ucsc.edu_internal
fiona8-0.calit2.uci.edu_internal
fiona8.ucsc.edu_internal
hydra.gi.ucsc.edu_internal
k8s-bafna-01.calit2.optiputer.net_internal
k8s-bharadia-01.sdsc.optiputer.net_internal
k8s-bharadia-02.sdsc.optiputer.net_internal
k8s-bharadia-03.sdsc.optiputer.net_internal
k8s-bharadia-04.sdsc.optiputer.net_internal
k8s-chase-ci-01.calit2.optiputer.net_internal
k8s-chase-ci-01.noc.ucsb.edu_internal
k8s-chase-ci-02.calit2.optiputer.net_internal
k8s-chase-ci-03.calit2.optiputer.net_internal
k8s-chase-ci-04.calit2.optiputer.net_internal
k8s-chase-ci-05.calit2.optiputer.net_internal
k8s-chase-ci-06.calit2.optiputer.net_internal
k8s-chase-ci-07.calit2.optiputer.net_internal
k8s-epyc-01.sdsc.optiputer.net_internal
k8s-gpu-01.calit2.optiputer.net_internal
k8s-gpu-02.calit2.optiputer.net_internal
k8s-gpu-03.sdsc.optiputer.net_internal
k8s-gpu-1.ucsc.edu_internal
k8s-gpu-2.ucsc.edu_internal
k8s-nvme-01.sdsc.optiputer.net_internal
k8s-nvme-01.ultralight.org_internal
k8s-ravi-01.calit2.optiputer.net_internal
knuron.calit2.optiputer.net_internal
maserati.sciencedmz.nps.edu_internal
netw-fiona.stanford.edu_internal
osg.chic.nrp.internet2.edu_internal
osg.kans.nrp.internet2.edu_internal
osg.newy32aoa.nrp.internet2.edu_internal
patternlab.calit2.optiputer.net_internal
perfsonar.csusb.edu_internal
prp-dtn.noc.ucsb.edu_internal
prp-gpu-1.t2.ucsd.edu_internal
prp-gpu-2.t2.ucsd.edu_internal
prp-gpu-3.t2.ucsd.edu_internal
ps-100g.sdsu.edu_internal
siderea.ucsc.edu_internal
stashcache.t2.ucsd.edu_internal
ucm-fiona01.ucmerced.edu_internal
ucsf-globus-dtn1.ucsf.edu_internal
venadi.ucsc.edu_internal
xcache-11.t2.ucsd.edu_internal
Average packet loss is 0.000%

PRPv1 to PRPv2:
The Transition from Network Diagnosing to Application Support

- **PRPv1 Designed, Built, and Installed ~40 Purpose-built FIONAs, Tuned to Measure and Diagnose End-to-End 10G, 40G and 100G**

- **But, Our PRP NSF Funding Requires Showing Use of the PRP by Scientists and Engineers—It's a Data Grant, not a Networking Grant**

- **Note: Our Scientists Clearly Need More than Bandwidth Tests**
  - **Teams of Scientists Want to Share Their Data at High Speed and Compute on It**
  - **They Need to Interoperate with Commercial and University Clouds**

- **So PRPv2 Added DMZ-Distributed Temporary Storage**
  - **1.7PB total in 14 ~200GB previous PRPv1 FIONAs in Campus DMZs**

PRP
PACIFIC RESEARCH PLATFORM

it²

Detailed Real-Time Monitoring of PRP Nautilus: UCD, UCSD, UCI, UCSB, UCLA, UCR, Stanford, UCAR, UCM, UCSC, UHM Ceph

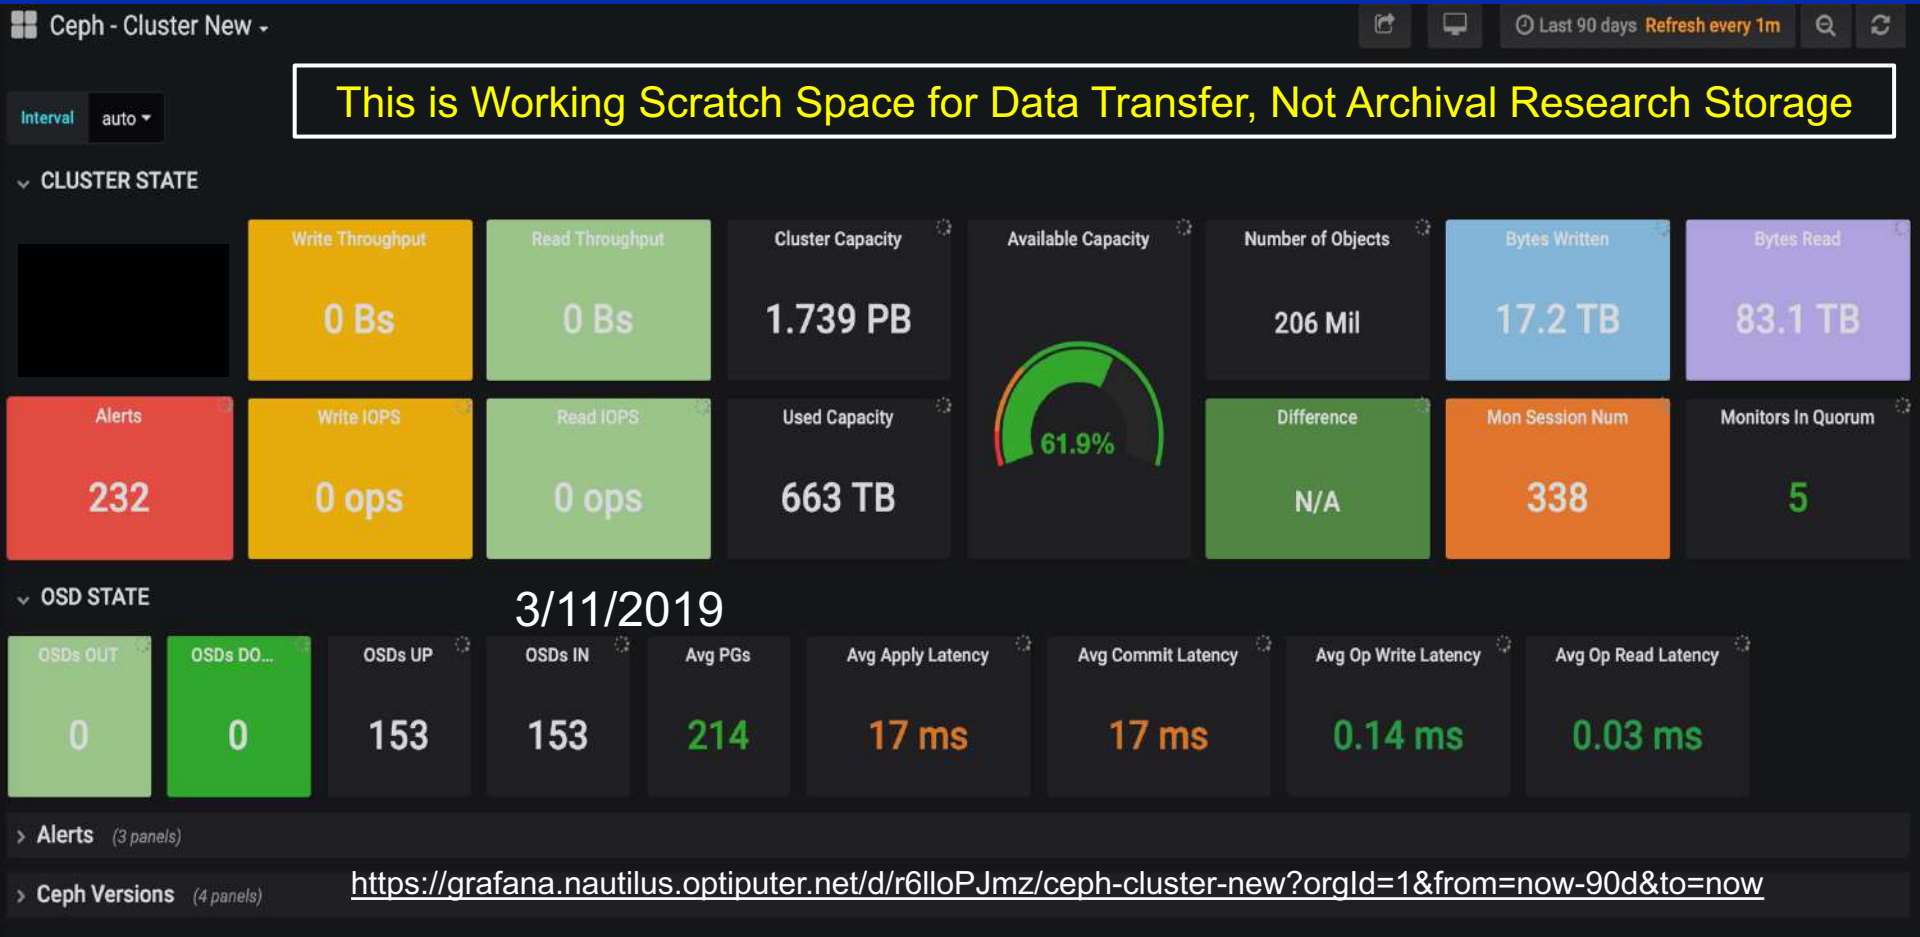

https://grafana.nautilus.optiputer.net/d/r6lloPJmz/ceph-cluster-new?orgId=1&from=now-90d&to=now

Regional Scale Cluster: Connected by PRP's Use of CENIC 100G Network
PRP's Nautilus Hypercluster

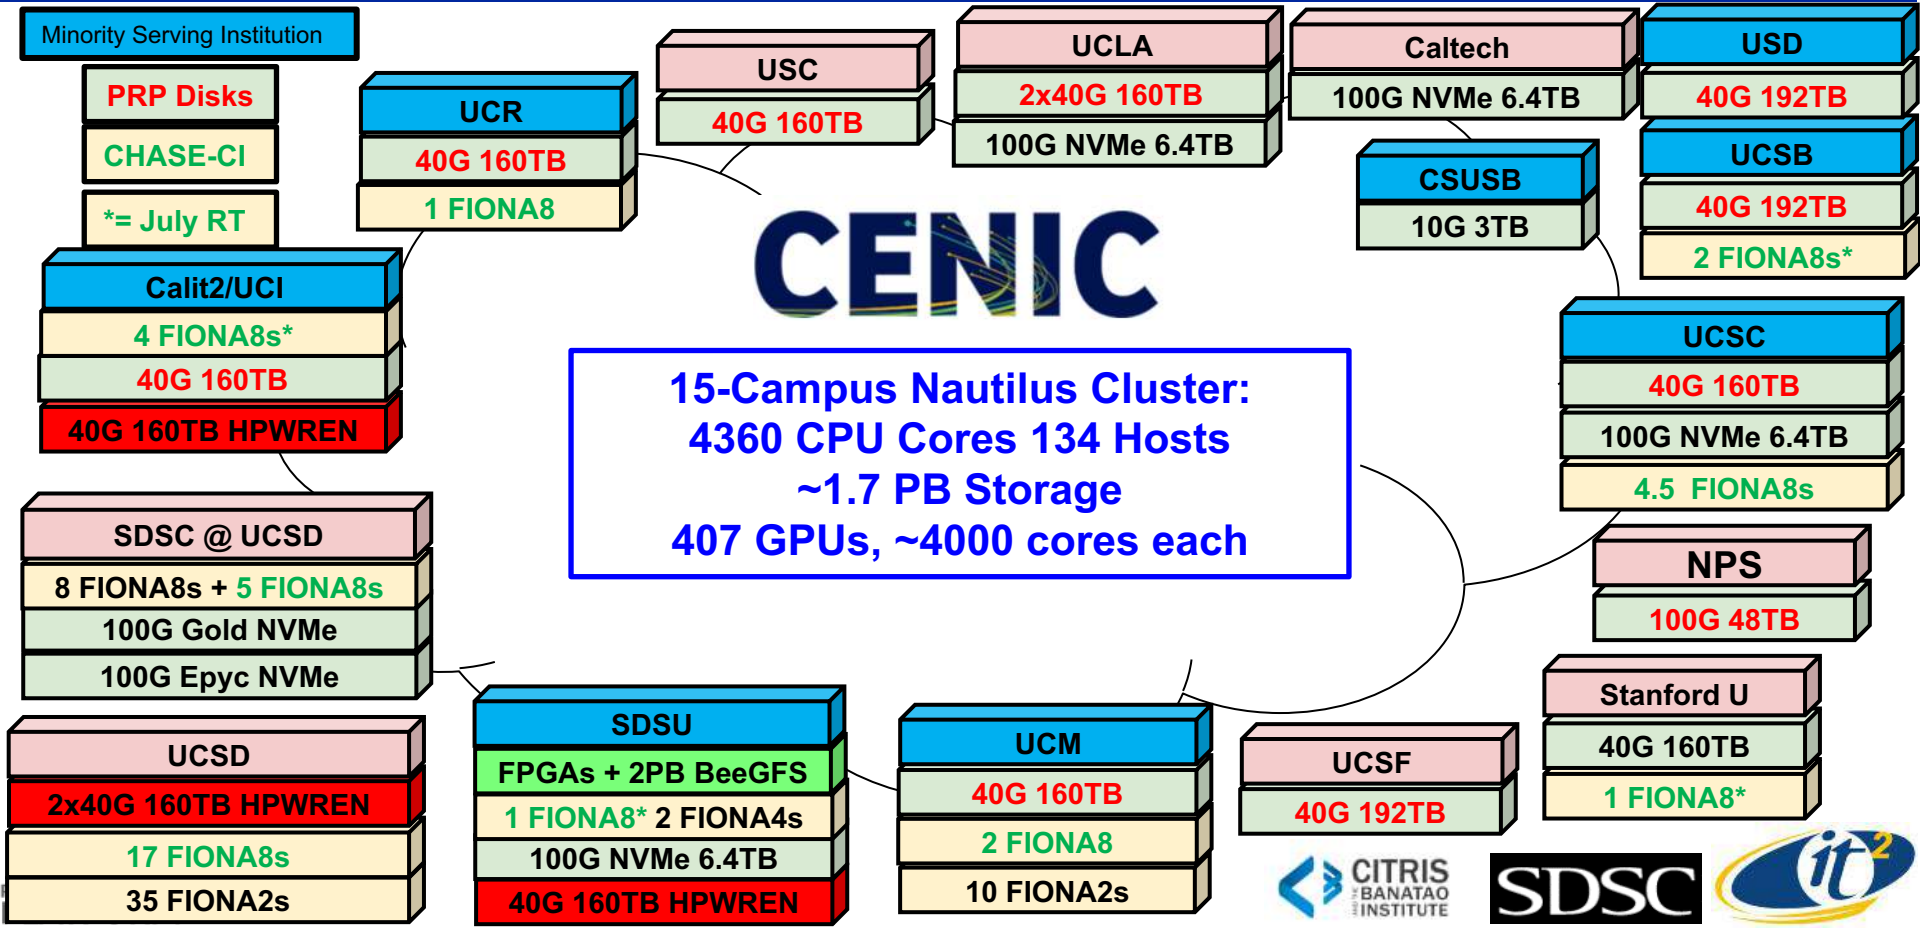

Grafana Showing State of Nautilus 9-10-19

2017-2020: CHASE-CI Adds Machine-Learning to the Data-Science Community Cyberinfrastructure

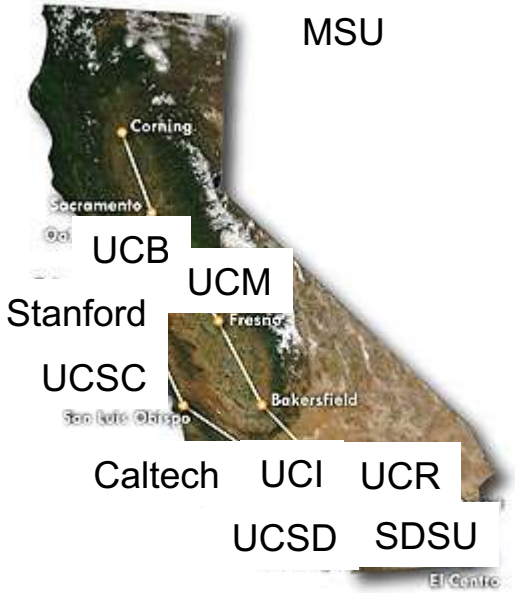

CI-New: Cognitive Hardware and Software Ecosystem Community Infrastructure (CHASE-CI)

For the Period September 1, 2017 – August 31, 2020

SUBMITTED – January 18, 2017

PI: Larry Smarr, Professor of Computer Science and Engineering, Director Calit2, UCSD
Co-PI: Tajana Rosing, Professor of Computer Science and Engineering, UCSD
Co-PI: Ken Kreutz-Delgado, Professor of Electrical and Computer Engineering, UCSD
Co-PI: Ilkay Altintas, Chief Data Science Officer, San Diego Supercomputer Center, UCSD
Co-PI: Tom DeFanti, Research Scientist, Calit2, UCSD

NSF Grant for 256 High Speed "Cloud" GPUs
For 32 ML Faculty & Their Students at 10 Campuses
To Train AI Algorithms on Big Data

PACIFIC RESEARCH PLATFORM

SDSC UC San Diego it²

Road Trip! Installing Community Shared Storage and GPUs in June, December & January at UC Merced, UC Santa Cruz, UC Riverside, and Stanford

160-192TB added to 14 Existing PRPv1 FIONAs

New FIONA8 at UCSC

PRP Engineers Designed and Built Several Generations of Optical-Fiber Big-Data Flash I/O Network Appliances (FIONAs)

UCSD-Designed FIONAs Solved the Disk-to-Disk Data Transfer Problem *at Near Full Speed* on Best-Effort 10G, 40G and 100G Networks

Two FIONA DTNs at UC Santa Cruz: 40G & 100G
Up to 192 TB Rotating Storage

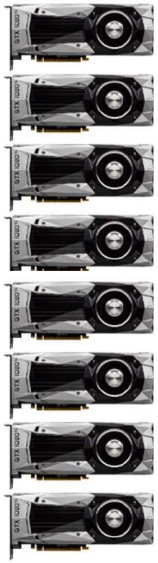

Add Up to 8 Nvidia GPUs Per 2U FIONA
To Add Machine Learning Capability

FIONAs Designed by UCSD's Phil Papadopoulos, John Graham, Joe Keefe, and Tom DeFanti

Top Nautilus GPU users August 2019					
PI	Campus	August 2019 GPU SU	FIONA8 Equivalent	August 2019 CPU SU	August 2019 Mem SU
Frank Wuerthwein	UCSD	80084	13.90	398124.41	8.13864E+14
Mark Alber	UCR	40761	7.08	37131.21	6.60061E+13
Hao Su	UCSD	16396	2.85	42547.91	2.78718E+14
Nuno Vasconcelos	UCSD	10991	1.91	11218.07	9.11693E+13
Jeff Krichmar	UCI	6587	1.14	6997.06	2.20582E+13
Falko Kuester	UCSD	6211	1.08	35404.91	5.68019E+14
Anshul Kundaje	Stanford	6063	1.05	1481.62	5.38638E+13
Ravi Ramamoorthi	UCSD	4822	0.84	6767.49	3.83436E+13
Larry Smarr	UCSD	4359	0.76	3171.25	2.20892E+13
Manmohan Chandraker	UCSD	3788	0.66	3304.47	1.02188E+14
Tom DeFanti	UCSD	3203	0.56	2040.4	8.82778E+12
Nuno Vasconcelos	UCSD	2293	0.40	3797.22	3.37342E+13
Kurt Schoenhoff	Australia	1921	0.33	4910.91	1.79054E+13
Nuno Vasconcelos	UCSD	1888	0.33	1017.46	1.67571E+13
Dinesh Bharadia	UCSD	1771	0.31	5724.15	2.71821E+13
Padhraic Smyth	UCI	1387	0.24	647.53	1.09787E+13
Jurgen Schulze	UCSD	1330	0.23	10.88	3.9717E+12
Larry Smarr	UCSD	1314	0.23	0.57	2.34185E+12
Jurgen Schulze	UCSD	1306	0.23	0.7	1.92583E+13
Nuno Vasconcelos	UCSD	1209	0.21	5984.29	1.33191E+13
Eric Shearer	UCI	1131	0.20	1308.7	3.85832E+12

Top Nautilus GPU Users in August 2019

FIONA8 equivalent: running an 8-GPU machine 24x7x30

Top User is IceCube in OSG background mode

Others are ML

PRP PACIFIC RESEARCH PLATFORM

it²

2017: PRP Connected 70 UCSD SunCAVE and 20 UCM WAVE 4K Screens to Share VR
2018: Added their 90 Game GPUs to PRP/OSG for Machine Learning Computations

UC Merced WAVE 20 Screens 20 GPUs

UCSD SunCAVE 70 Screens 70 GPUs

PRP PACIFIC RESEARCH PLATFORM

Leveraging UCM Campus Funds and NSF CNS-1456638 & CNS-1730158 at UCSD

it²

By Amble - Own work, CC BY-SA 3.0, https://commons.wikimedia.org/w/index.php?curid=8773726

- IceCube Neutrino Observatory has been using 120 Nautilus GPUs since March 8
- This would cost $2,880/day in a commercial cloud (at $1/hr) or ~$20,000/week
- An 8-GPU FIONA8 for Nautilus costs $20,000 to buy

GPU Simulations are Needed *to Improve Ice Model*.
=> Results in Significant Improvement in Pointing Resolution for Multi-Messenger Astrophysics

PRP PACIFIC RESEARCH PLATFORM

it²

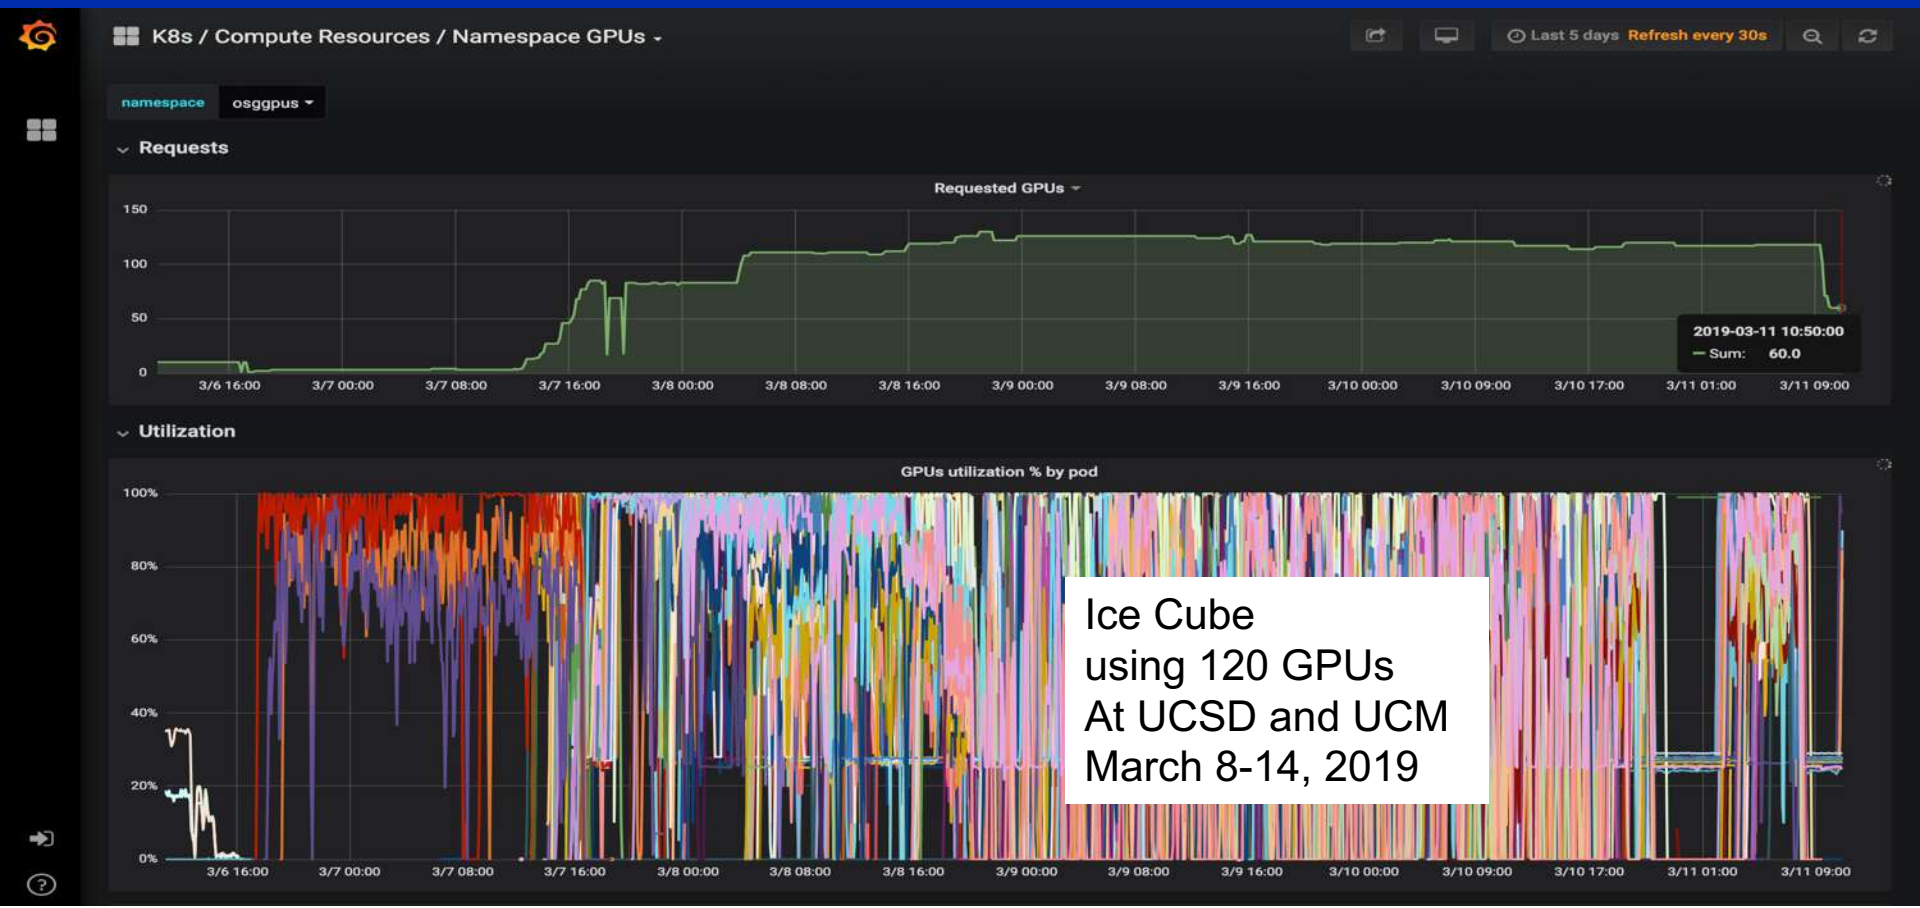

March 13, 2019: “This morning there was a big demo in the SunCAVE. The IceCube pods were kicked out automatically when the SunCAVE GPUs were in use, and restarted when the demo was over. **No admin intervention needed**.”—Igor Sfiligoi, SDSC

Very Cost-Effective for Academic Machine Learning and Data Sharing

- **Data science researchers need DTNs with lots of storage, encryption and lots of GPUS**
- **One UC spends $40,000 in cloud GPU per published grad student paper**
- **Another spends $20,000 for undergrad ML AWS access in just one course**
- **Instead, add to our Nautilus hypercluster (or clone it & federate):**
  - **UCSD ECE Department bought 4 FIONA8s, buying 4 more**
  - **UCSD Physics Department. bought 6 FIONA8s**
  - **UCSD CSE researchers bought 4 FIONA8s to add to Nautilus, buying 20 more**
  - **UCSD Instructional IT has 13 FIONA8s for Machine Learning/AI class labs**
- **Working Storage on Nautilus FIONAs is**
  - **very inexpensive (12TB drives are ~$430 each—16 per FIONA. FISA encrypted drives @ same cost)**
  - **and very high speed (most FIONAs are 40/100G and are located in ScienceDMZs)**

Clemson's Alex Feltus: "I cannot wait to add a node to the Nautilus compute fabric!" (He didn't wait)

PRP PACIFIC RESEARCH PLATFORM

it²

UCSD's Information Technology Services Adapted PRP FIONA8s To Support Data Science Courses

Instructional Data Science Machine Learning Platform:

Instead of Spending ~$20,000/Quarter/Course on Commercial Clouds:
97 Courses over 6 Quarters → $4M vs. $240K over 12 Quarters
At least 20,000 Students

- 104 GPUs
  - 80 GTX 1080Ti
  - 16 RTX 2080Ti
  - 1.05 Petaflops
- 28 nodes
  - 544 CPU cores
  - 6.5TB RAM
- 40TB Flash Storage
- 10G networking

Source: UCSD ITS

The Student GPUs
Have Supported Thousands of Students in Dozens of Courses

Source: UCSD ITS

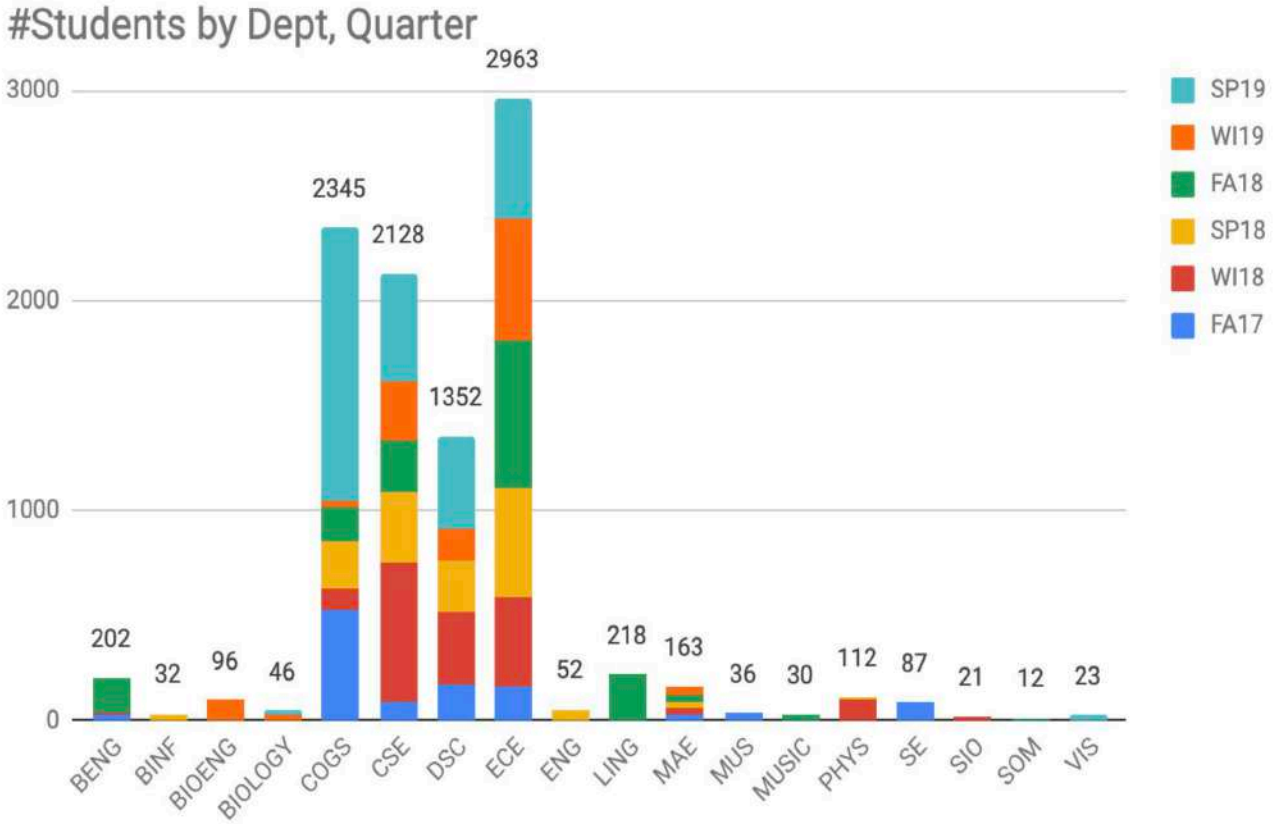

PACIFIC RESEARCH PLATFORM

it²

Student GPU Demand Is Variable Allowing for Other Student Uses

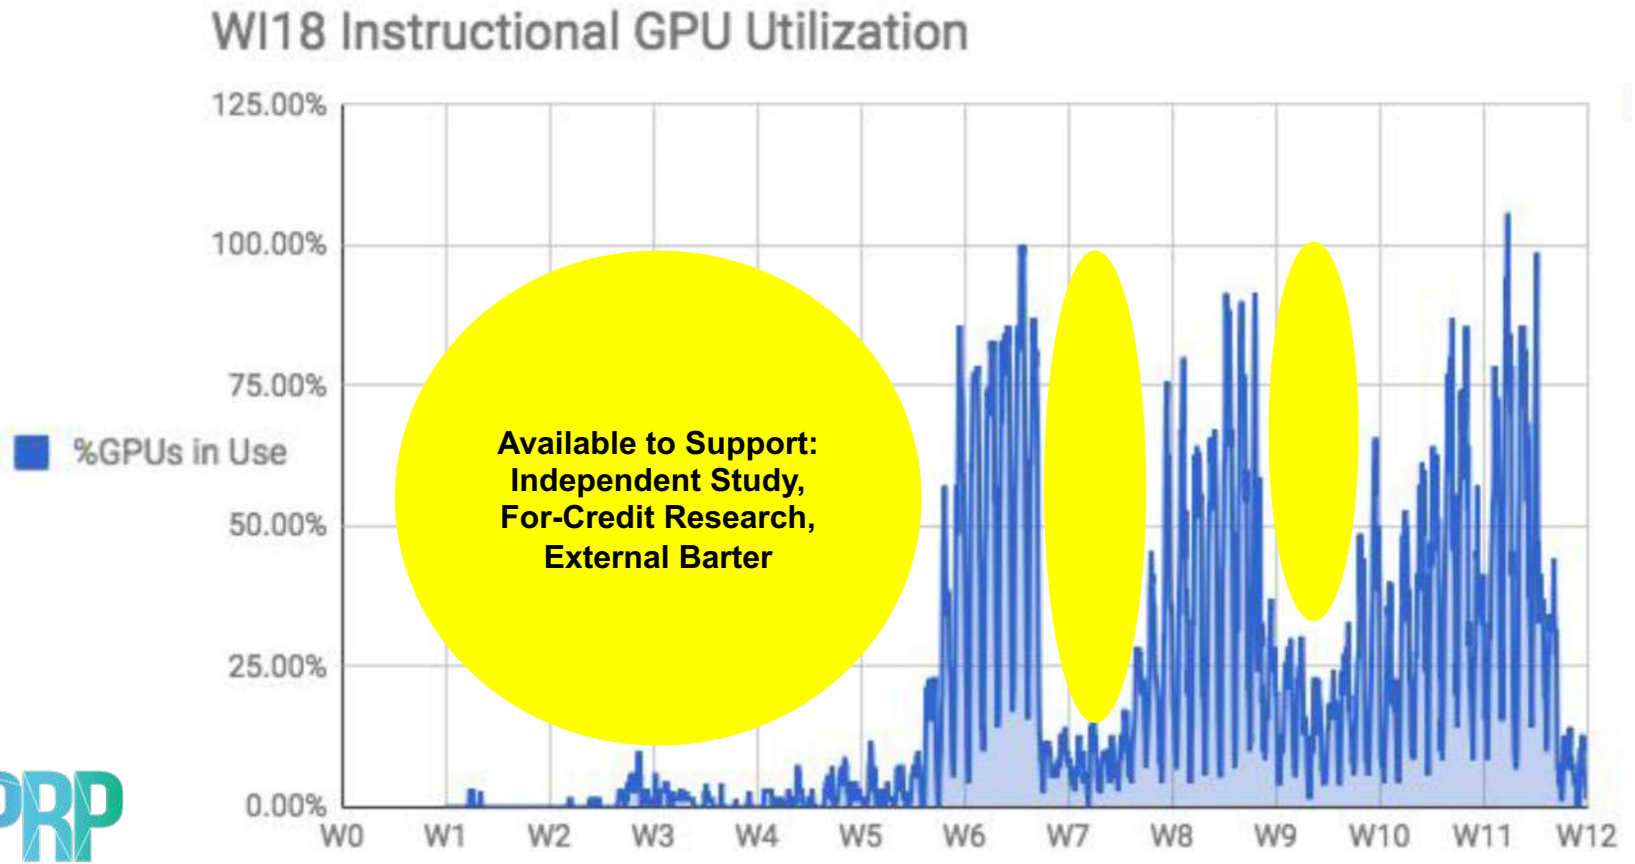

Source: UCSD ITS

405 Research GPUs in Nautilus 9-10-2019

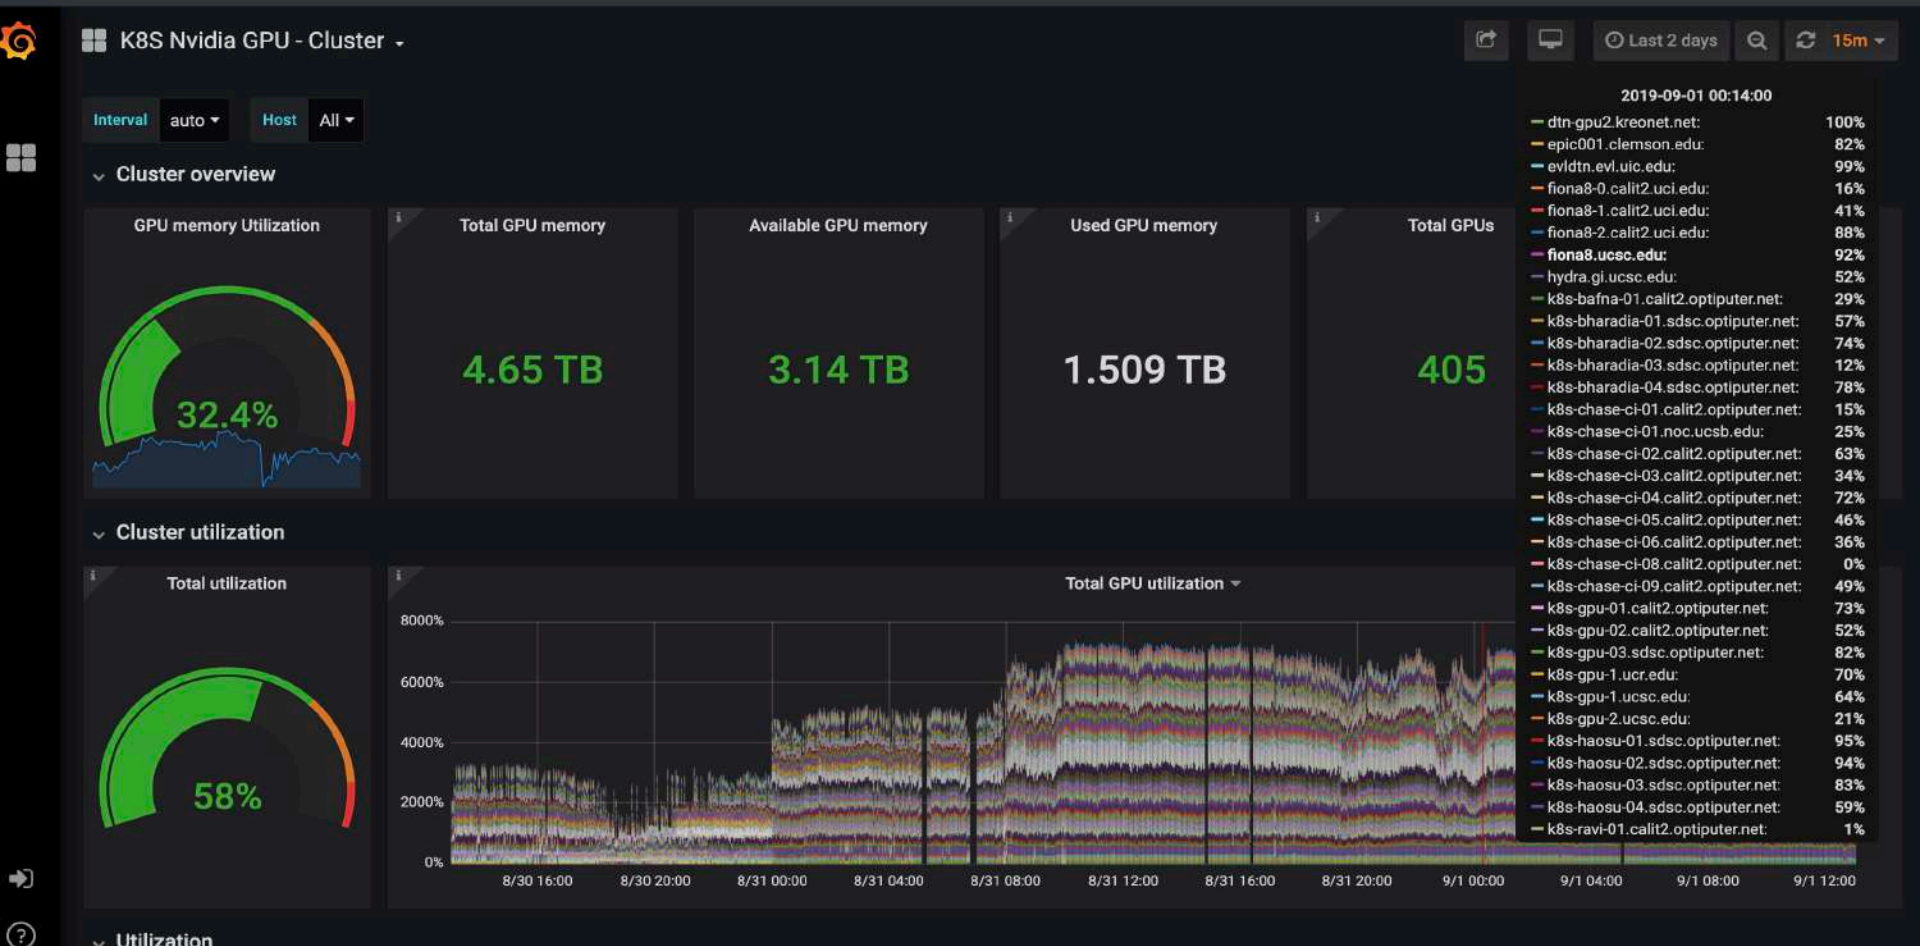

Community Participation Allowed PRP-Paid Nautilus Nodes to be Quadrupled—*Pot Luck Supercomputing*™

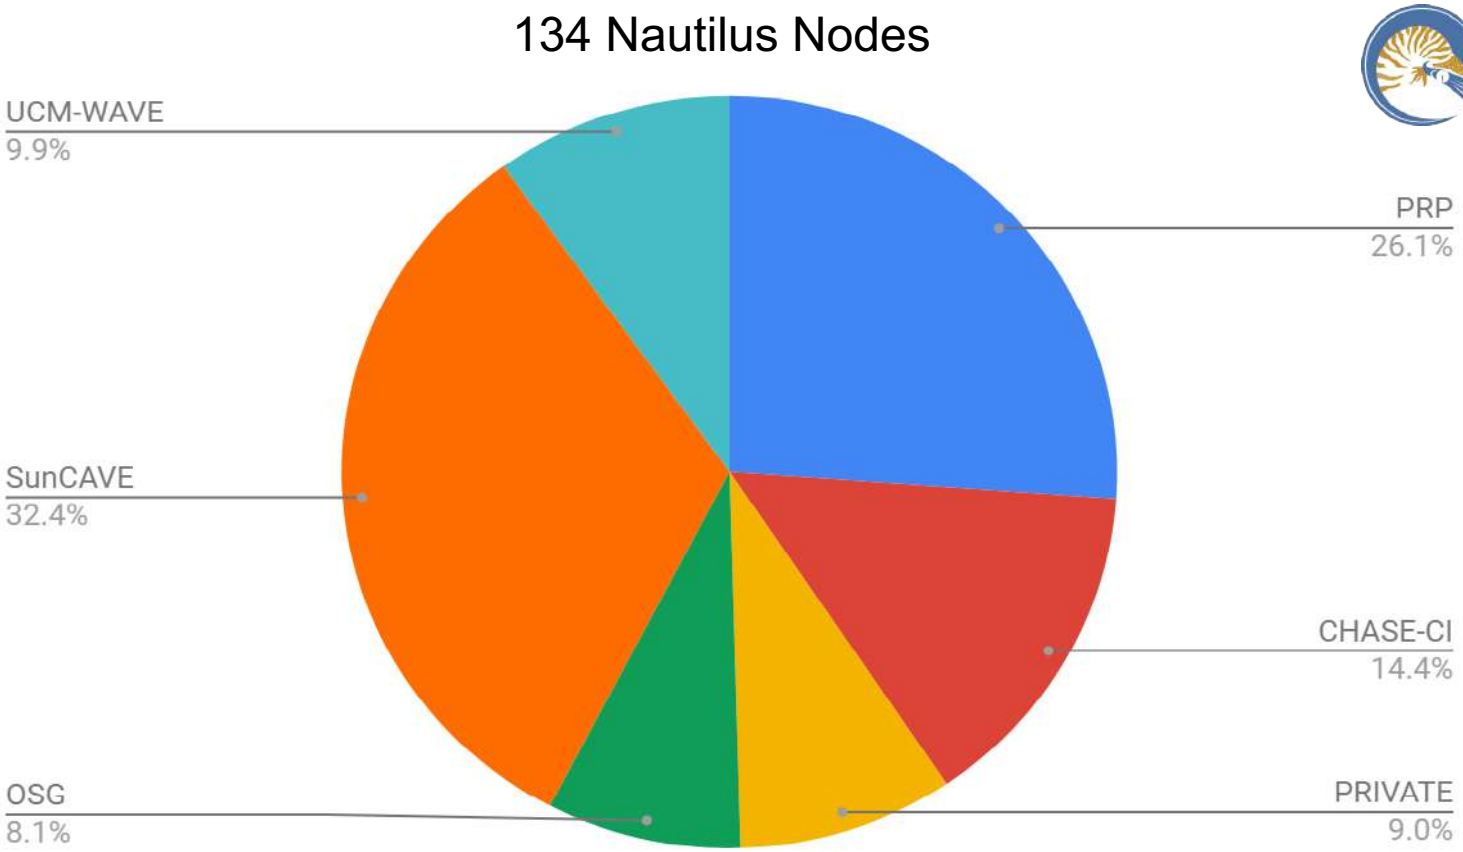

Why PRPv2 Adopted Kubernetes

- **PRP FIONAs Are Coupled by Kubernetes Into the "Nautilus Hypercluster"**
  - **Kubernetes "Pods" Encapsulate Application Container(s), Storage Resources, and Execution Options**
  - **Implements PRP Cooperative Research Groups Support with Policy-Based Scheduling by Use of CILogon and Kubernetes Namespaces—704 Users in Namespaces as of 7/15/19**
  - **Allows Cloud Native Storage Integration (e.g., Rook/Ceph/EdgeFS)**
  - **Enables Us to Update Overnight, without local assistance, a RP Scaling Necessity**
  - **Emerging Solutions for Sophisticated SDN Overlay Network, Firewall, and Network Policy Controls**
- **Allows Easy User Job Scaling to Heterogeneous Platforms:**
  - **Deskside, Rack-Mounted, Supercomputers, even IOT Gizmos like ML on Remote Cameras**
  - **Amazon Elastic Container Service for Kubernetes (Amazon EKS)**
  - **Google Kubernetes Engine (GKE) (TensorFlow)**
  - **Microsoft Azure Kubernetes Service (AKS)**
  - **Also Comet and other XSEDE assets**

"Kubernetes with Rook/Ceph Allows Us to **Manage** **Petabytes of Distributed Storage and GPUs for Data Science**,
While We Measure and Monitor Network Use."
--John Graham, Calit2/QI UC San Diego

PRP PACIFIC RESEARCH PLATFORM

it²

2018-2019: National-Scale Pilot - Using CENIC & Internet2 to Connect Quilt Regional R&E Networks

"Towards The NRP"
3-Year Grant
Funded by NSF
$2.5M
October 2018

PI Smarr
Co-PIs Altintas
Papadopoulos
Wuerthwein
Rosing
DeFanti

SDSC

PRP
PACIFIC RESEARCH PLATFORM

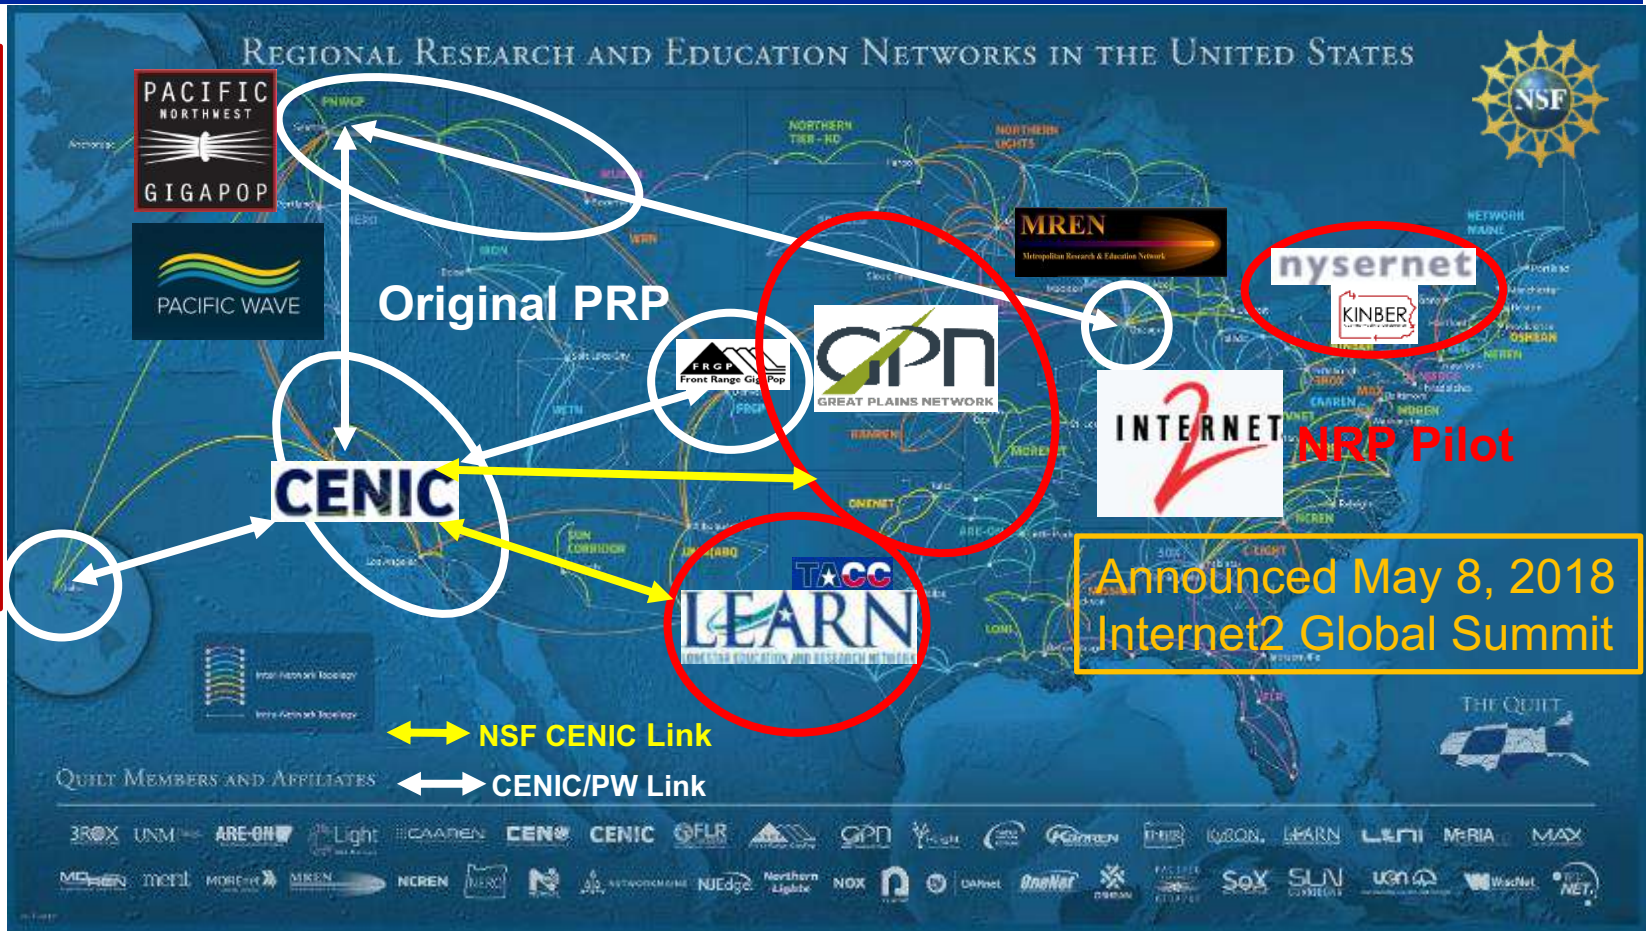

NRP_GridFTP - Throughput

Throughput >= 7500Mbps | Throughput < 7500Mbps | Throughput <= 5000Mbps | Unable to retrie

Found a total of 4 problems involving 3 hosts in the grid

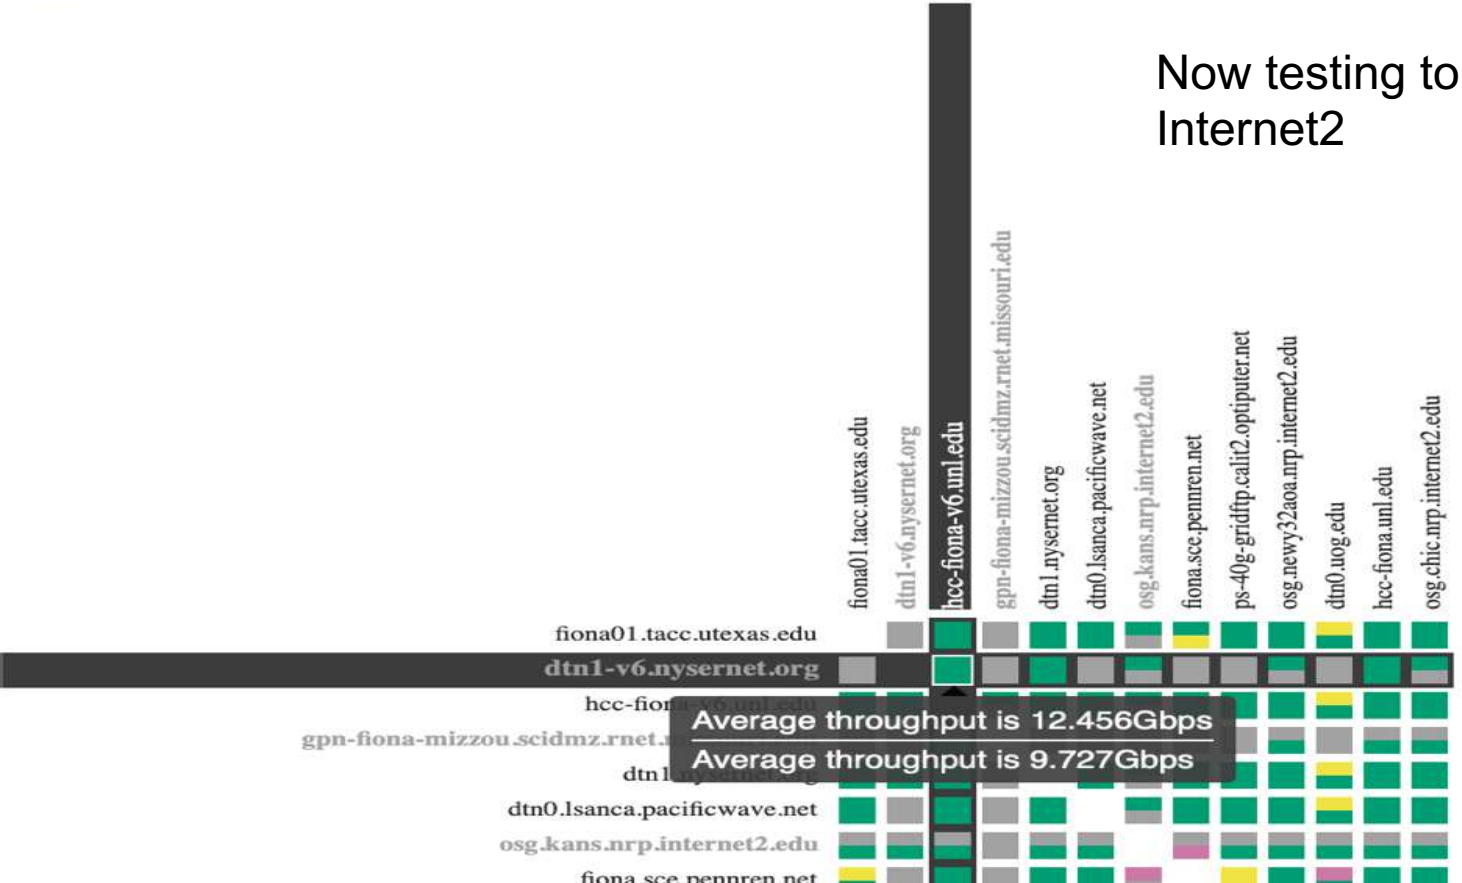

Now testing to AWS and Internet2

TNRP = PRP (CENIC, PNWGP, FRGP, HI, and MREN) + OSG + ESnet + Quilt + NRP Pilot (I2, KINBER, Learn, GPN, NYSERnet) + MCNC + NM Tribal + …

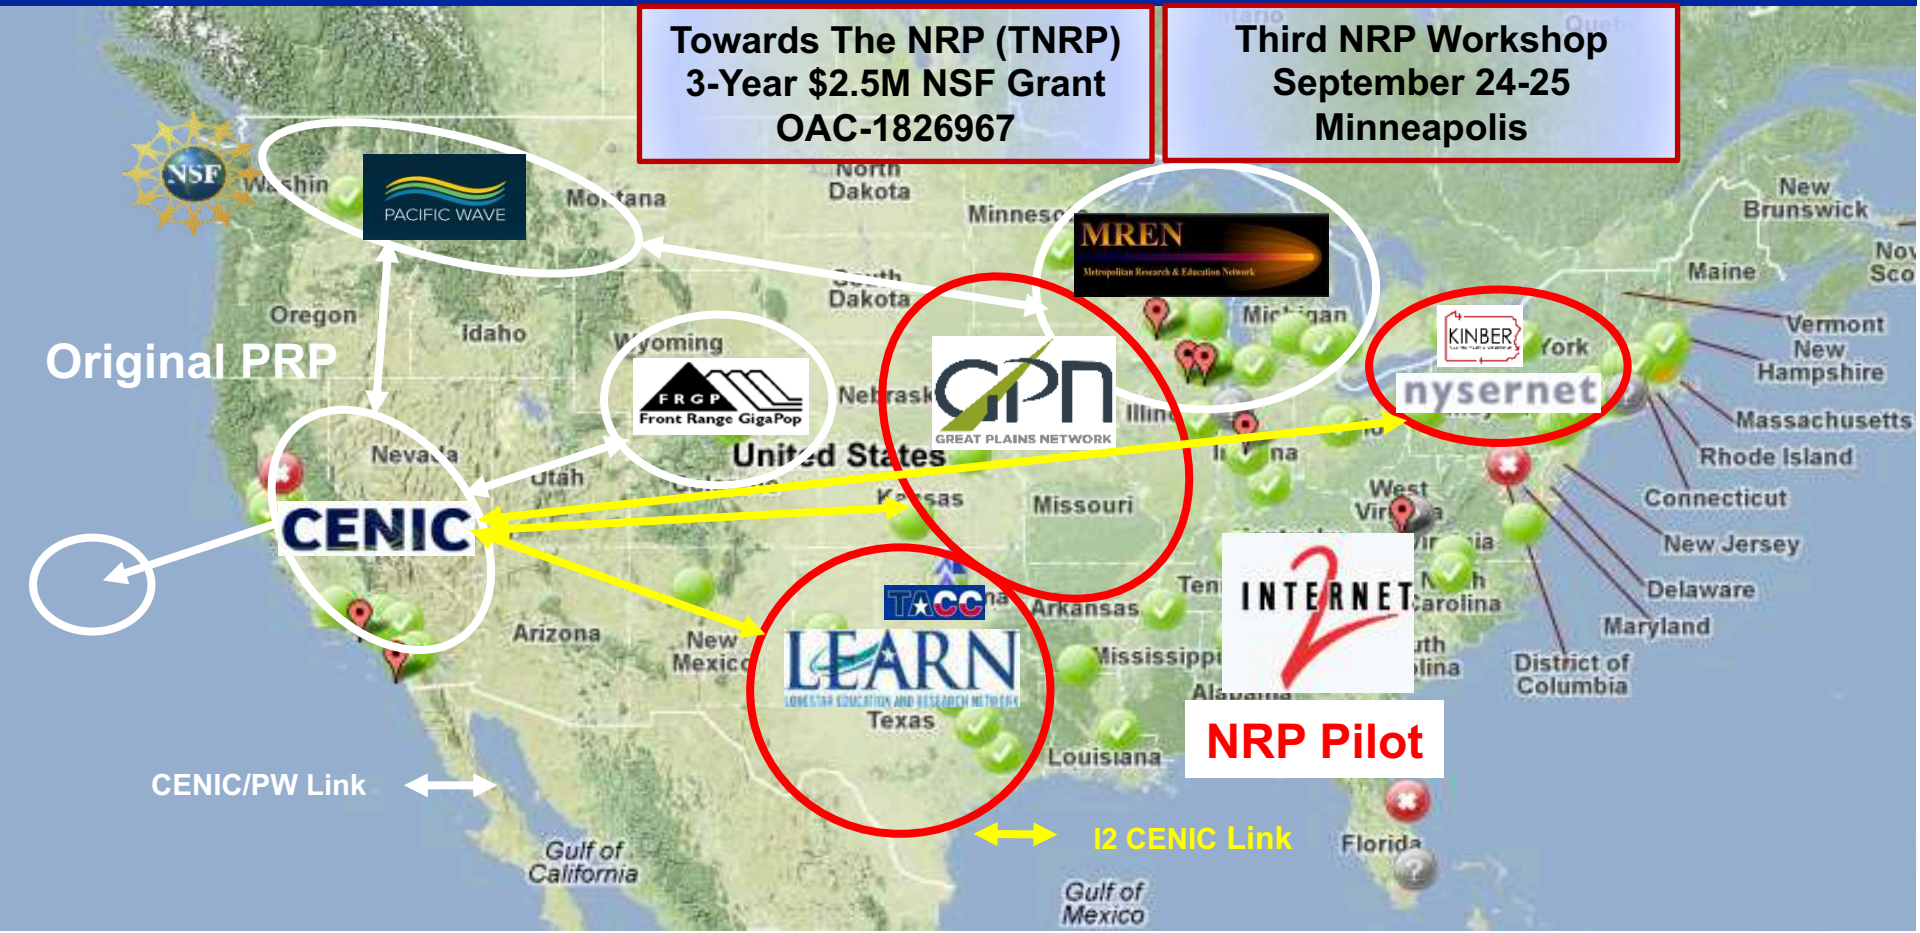

Nautilus Has International Nodes
The Global Research Platform is Emerging (1GRP Next Week Here!)

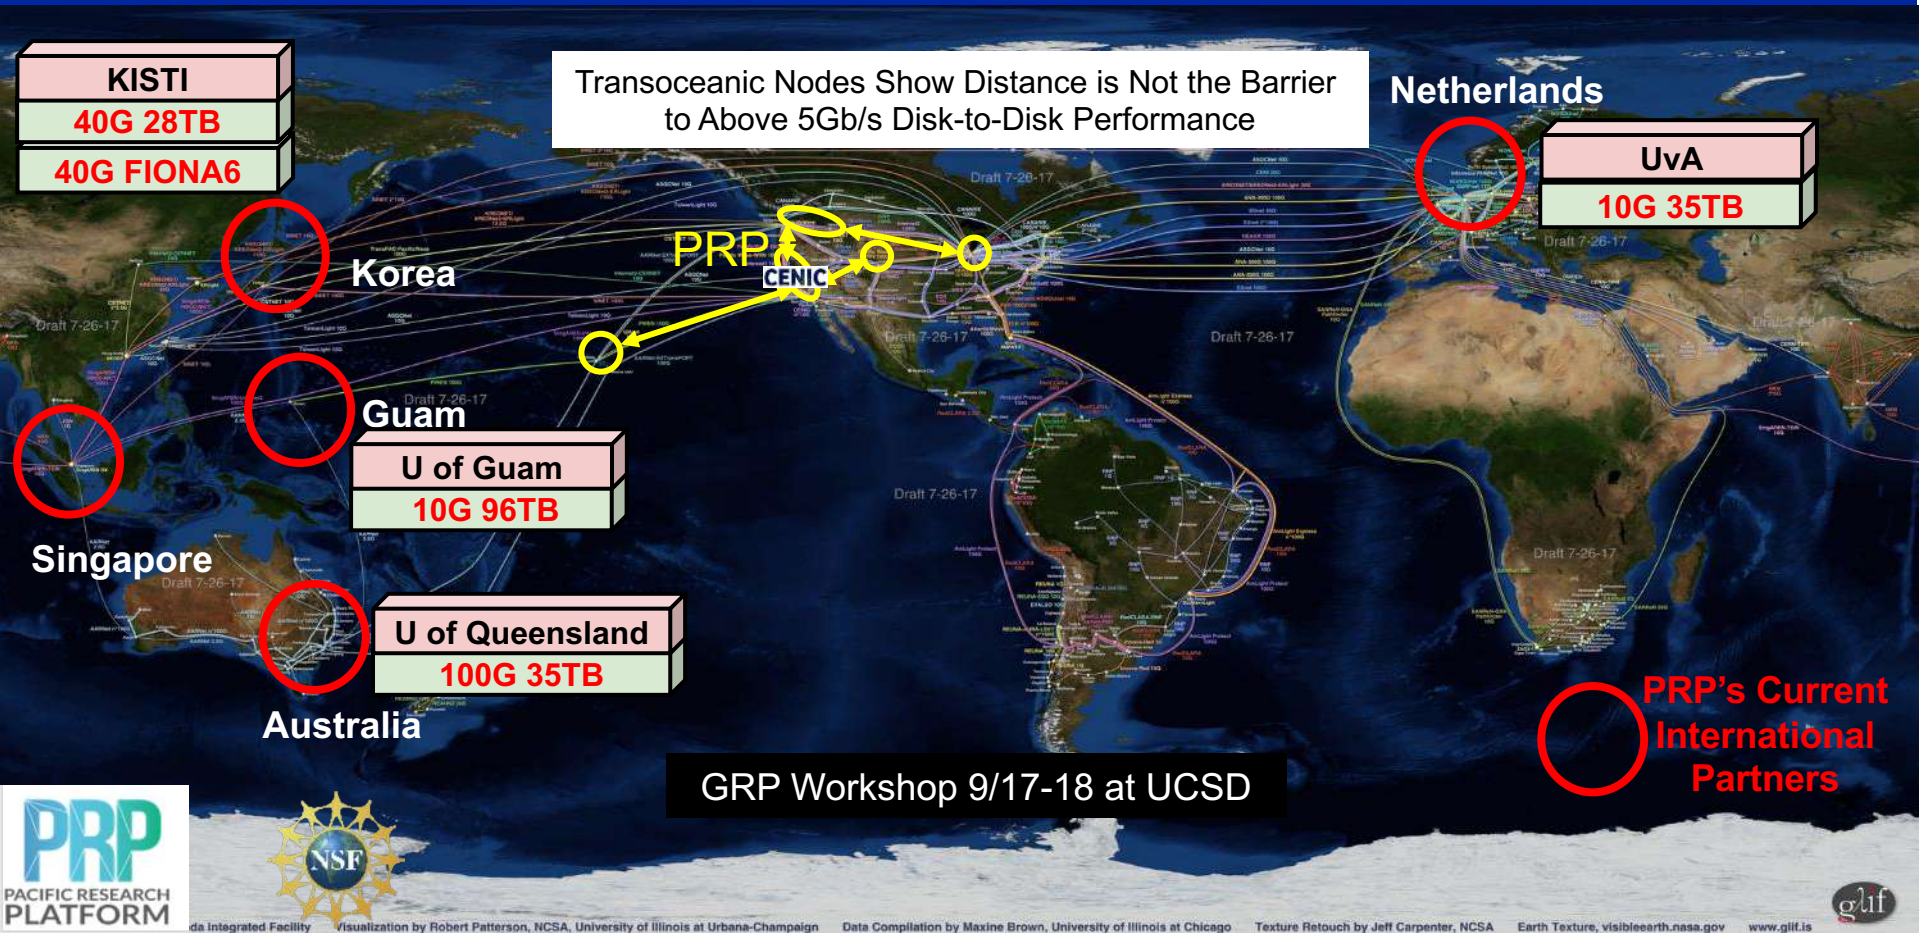

PRP PACIFIC RESEARCH PLATFORM

NSF

da Integrated Facility Visualization by Robert Patterson, NCSA, University of Illinois at Urbana-Champaign Data Compilation by Maxine Brown, University of Illinois at Chicago Texture Retouch by Jeff Carpenter, NCSA Earth Texture, visibleearth.nasa.gov www.glif.is

glif

Excellent Performance California to UQ (100G)

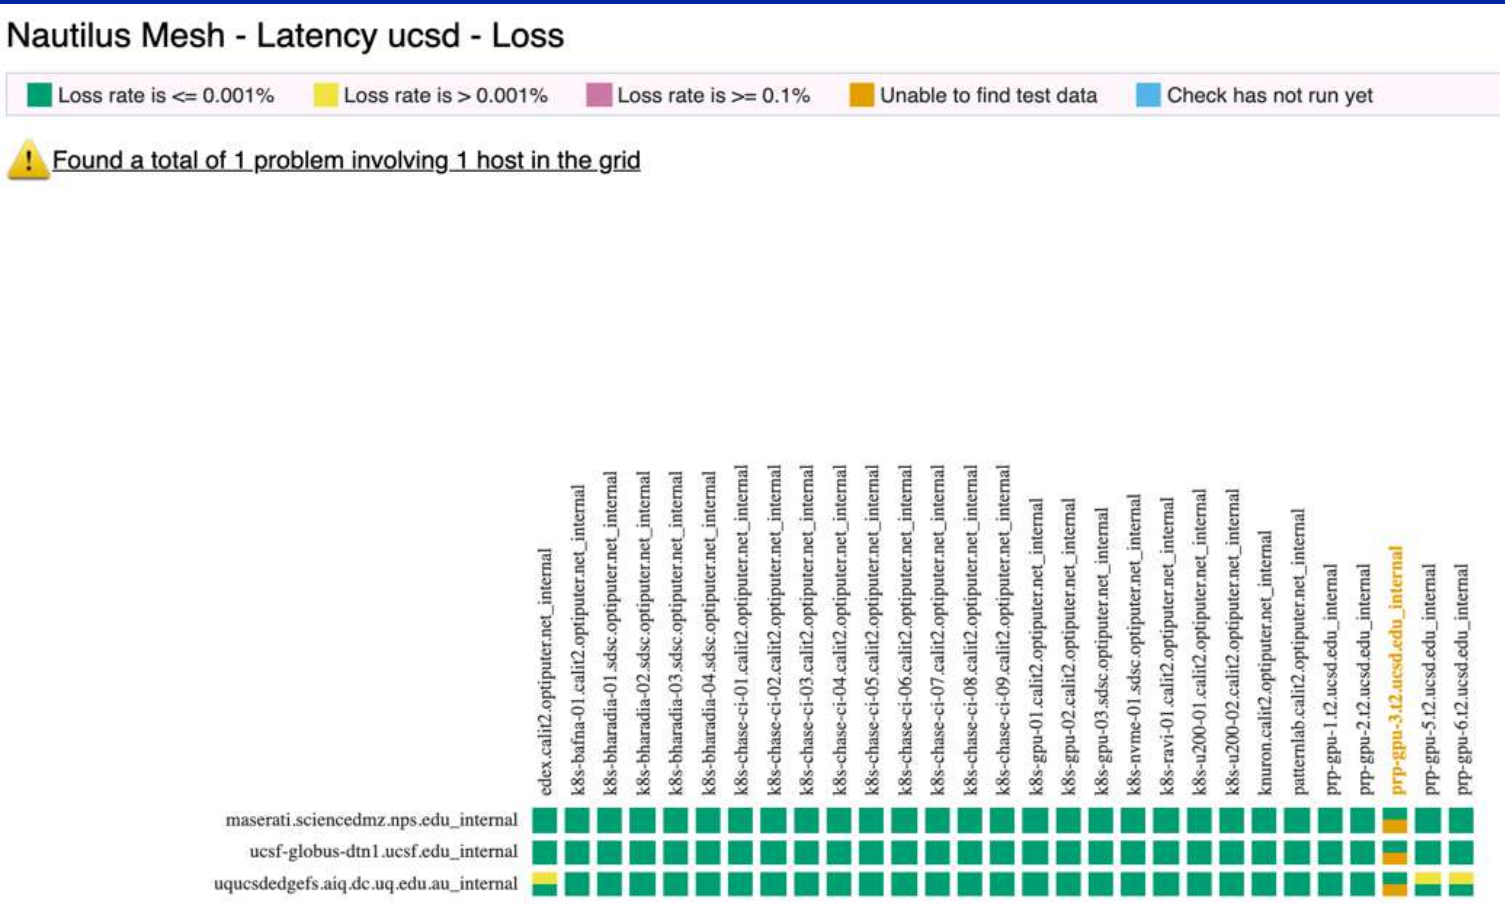

ElastiFlow: See Inter-cluster Campus-Level Traffic Flow Grouped by AS

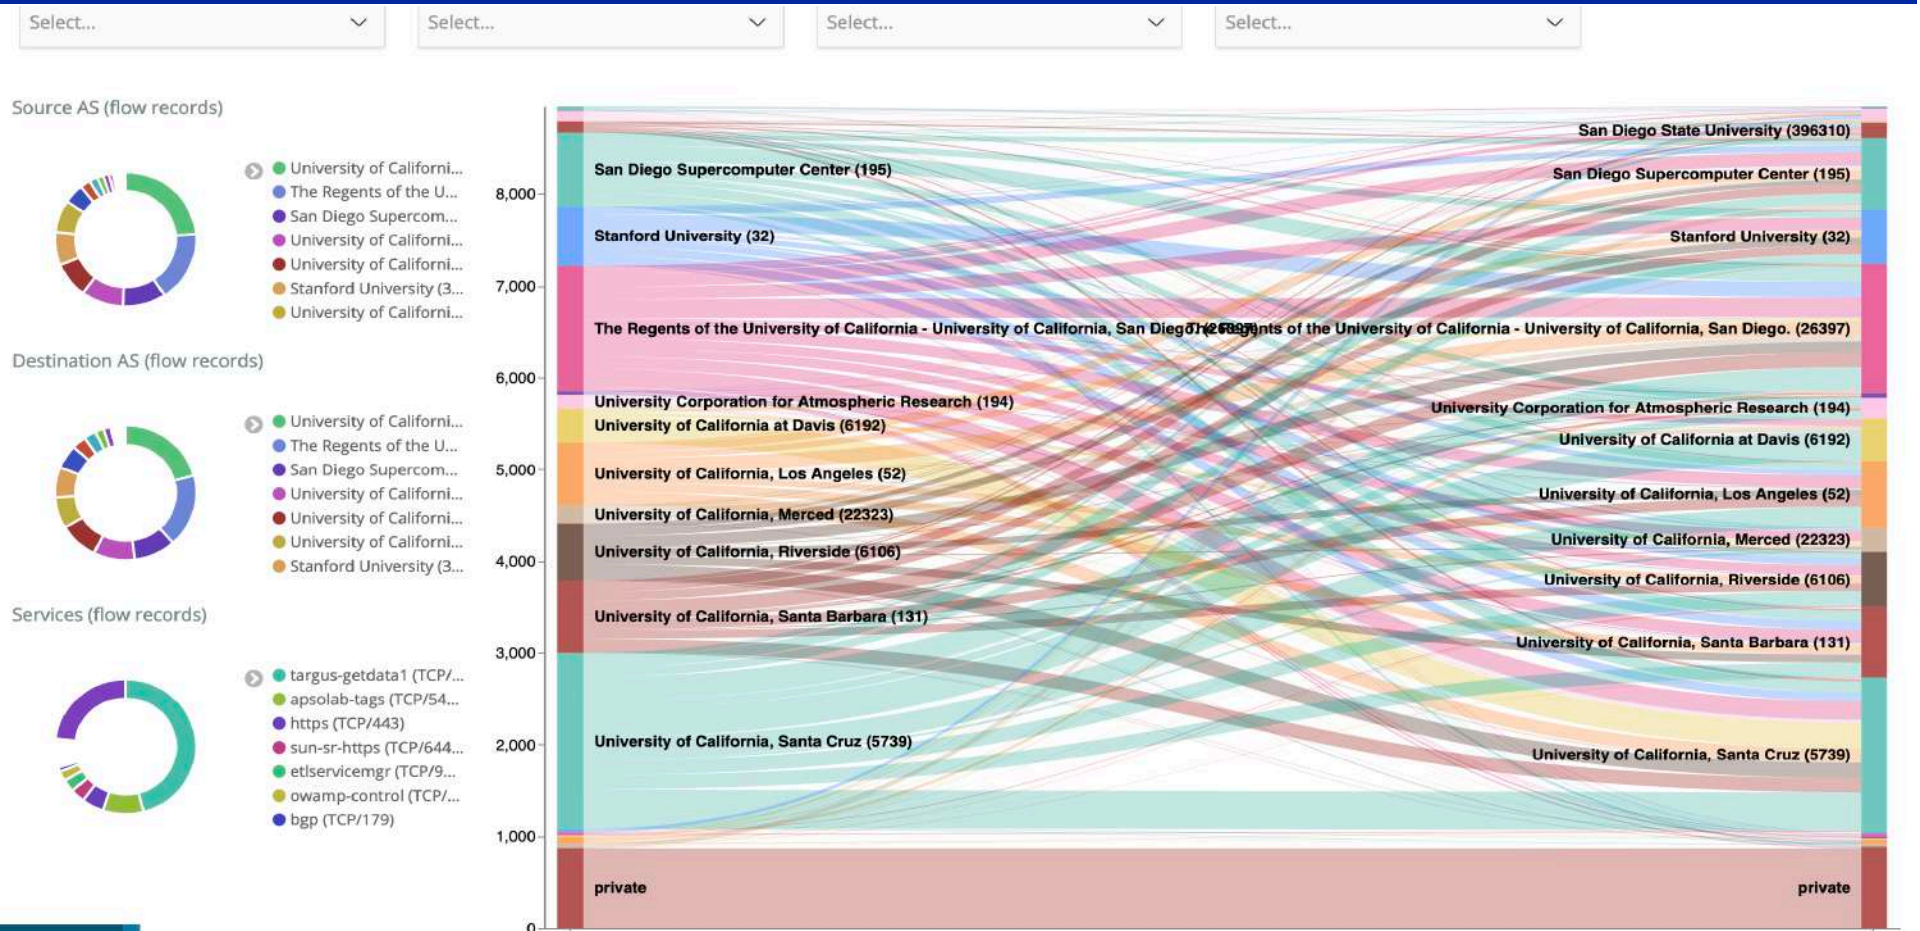

PRP Tech Coming

- Support users with IoT/Robotics/Augmented Reality needs
  - Nvidia Jetson Xaviers and Nanos
- Also FPGA data-center boards (Xilinx U200s, Micron SB-852)
  - Compute: application acceleration (e.g.,TensorFlow)
    - Climate/weather segmentation
    - Inferencing
    - Satellite imagery ortho-rectification (align w/wildfire maps)
  - 100G SDX P4 build out (SDSU, USC, NU, FIU, UCSD, Caltech)
- And Tensor Flow Cores and TPUs
  - Nvidia 2080-Ti cards: 544 Tensor Cores each, 4,352 per FIONA8
  - Our Nautilus Users can Access Google Cloud TPUs
  - Google Edge TPU Coral Development Boards and USB-C Edge TPU Accelerator/co-Processor

PACIFIC RESEARCH PLATFORM

P2PRP: Pacific to Pacific Rim Platform!

- **Top down Great Networking with 10-100Gbps Science DMZ Performance is a Necessary *but not Sufficient Condition* for Data-Driven Researchers**
  - **They need Science DMZs & DTNs with Lots of Low-Cost Storage, Encryption, Large RAM CPUs, GPUs, TPUs, FPGAs, and High-Availability Computing**
- **Measuring and Monitoring is Key to Better Usage and Security**
- **Compatibility with CloudBank, Google, Microsoft, and Amazon Clouds, and NSF/DOE Supercomputers Helps Ensure Scalability and Continuation**
- **Convergence with Open Science Grid/I2 Brings In Global Experience**

PRP PACIFIC RESEARCH PLATFORM

it2

PRP/TNRP/CHASE-CI Support and Community:

- US National Science Foundation (NSF) awards to UCSD, NU, and SDSC
  - CNS-1456638, CNS-1730158, ACI-1540112, ACI-1541349, & OAC-1826967
  - OAC 1450871 (NU) and OAC-1659169 (SDSU)
- UC Office of the President, Calit2 and Calit2's UCSD Qualcomm Institute
- San Diego Supercomputer Center and UCSD's Research IT and Instructional IT
- Partner Campuses: UCB, UCSC, UCI, UCR, UCLA, USC, UCD, UCSB, SDSU, Caltech, NU, UWash UChicago, UIC, UHM, CSUSB, HPWREN, UMo, MSU, NYU, UNeb, UNC,UIUC, UTA/Texas Advanced Computing Center, FIU, KISTI, UVA, AIST
- CENIC, Pacific Wave/PNWGP, StarLight/MREN, The Quilt, Kinber, Great Plains Network, NYSERNet, LEARN, Open Science Grid
- Internet2, DOE ESnet, NCAR/UCAR and Wyoming Supercomputing Center

PRP PACIFIC RESEARCH PLATFORM

it2