

## YAML2RPM – Another Attempt at Taming Scientific Software

Philip Papadopoulos, PhD.  
[ppapadop@uci.edu](mailto:ppapadop@uci.edu)

Director, UCI Research Cyberinfrastructure Center (RCIC)

Nadya Williams  
[npw@uci.edu](mailto:npw@uci.edu)  
RCIC



## What is The Research Cyberinfrastructure Center?

- Created at the recommendation of UCI Research Cyberinfrastructure Vision document
- Build and maintain computing/data infrastructure for all UCI researchers to improve competitiveness and operational efficiency
- Support grant proposals (entire process)
- Shared-funding model
  - Campus \$\$ **pay for baseline computing/data**
  - Researchers pay **cost recovery for additional cycles/storage**


# Key Resources and Functions

- Campus Research Storage Pool
  - Expandable, Reliable, On-line storage for research data
  - Two copies of all data at all times (2+ PB)
- Cluster Computing
  - HPC “Condo” (~10K cores). HPC<sup>3</sup> coming soon (4K+ cores)
- Parallel File System Storage
  - High-performance, lowest cost, single copy of data (2+ PB)
- Science Application Software Integration
  - More than 700 open-source domain-specific applications available
- Support Research Grant Process
  - Concept through implementation
  - Can include one or more of the following:
    - Hardware recommendation, Facilities, Data management, writing, fractional support of RCIC staff, Co-Investigator

# A Software Perspective...

- UCI runs a shared HPC cluster for campus researchers
- A very diverse community of researchers
- Care and feeding of an Applications “Zoo”:
  - Have built and must maintain/update:
    - 780+ Unique Scientific Software packages
    - 1200+ when considering multiple version
    - Note: CentOS7 builds/maintains ~9000 packages.
- Existing cluster is CentOS6
  - New cluster is CentOS7
  - Will need to be agile, CentOS8 is out and users will demand it in a few years
  - Need to rebuild applications for each revision.
- Application complexity is on par with a supercomputer center, but we are not the same size resource (or staff)

## Current (inherited) State



- How was Application X configured/compiled? (some have compile scripts)
- What's affected if you update application X? (can't easily answer this)
- When was something installed/updated? (partial answer: can look at file system snapshots/backups)
- What if users need different versions of the same software? (supported via environment modules)
- What if some applications should not be available on all nodes (shared FS makes this very problematic)

## Current state works, but ...



## Can we tame our software stack?

## Claim: Formally packaging software can address many of these issues

- Query what's installed ([manifest](#))
- Assure that no two software packages conflict in the file system ([reduce contention](#))
- Determine when something was added/removed/updated ([tracking](#))
- Query if the current state on disk is different than it is supposed to be. ([verification](#))
- If properly built, when a prerequisite software is removed, all dependent software is also removed ([consistency](#))
- Ability to update software version in place ([manageability](#))

[Management is improved dramatically](#) if you can use the packaging system that is [NATIVE](#) to the OS

- ➔ If using CentOS/RedHat use RPM format
- ➔ if using Debian/Ubuntu use DPKG format

## Why do many cluster admins/owners not build their own packages?

- “It’s too difficult to create a package for each piece of software, I have enough on my hands figuring out how to compile X”
- “Python/Perl/R/Anaconda/Brew manage their own packages, that’s most of my work”
- “I just manage one system, packaging is overkill”
- “My users want so much software that I just get it done as fast as possible. “
- “I don’t see anything wrong with compiling into a shared NFS area”





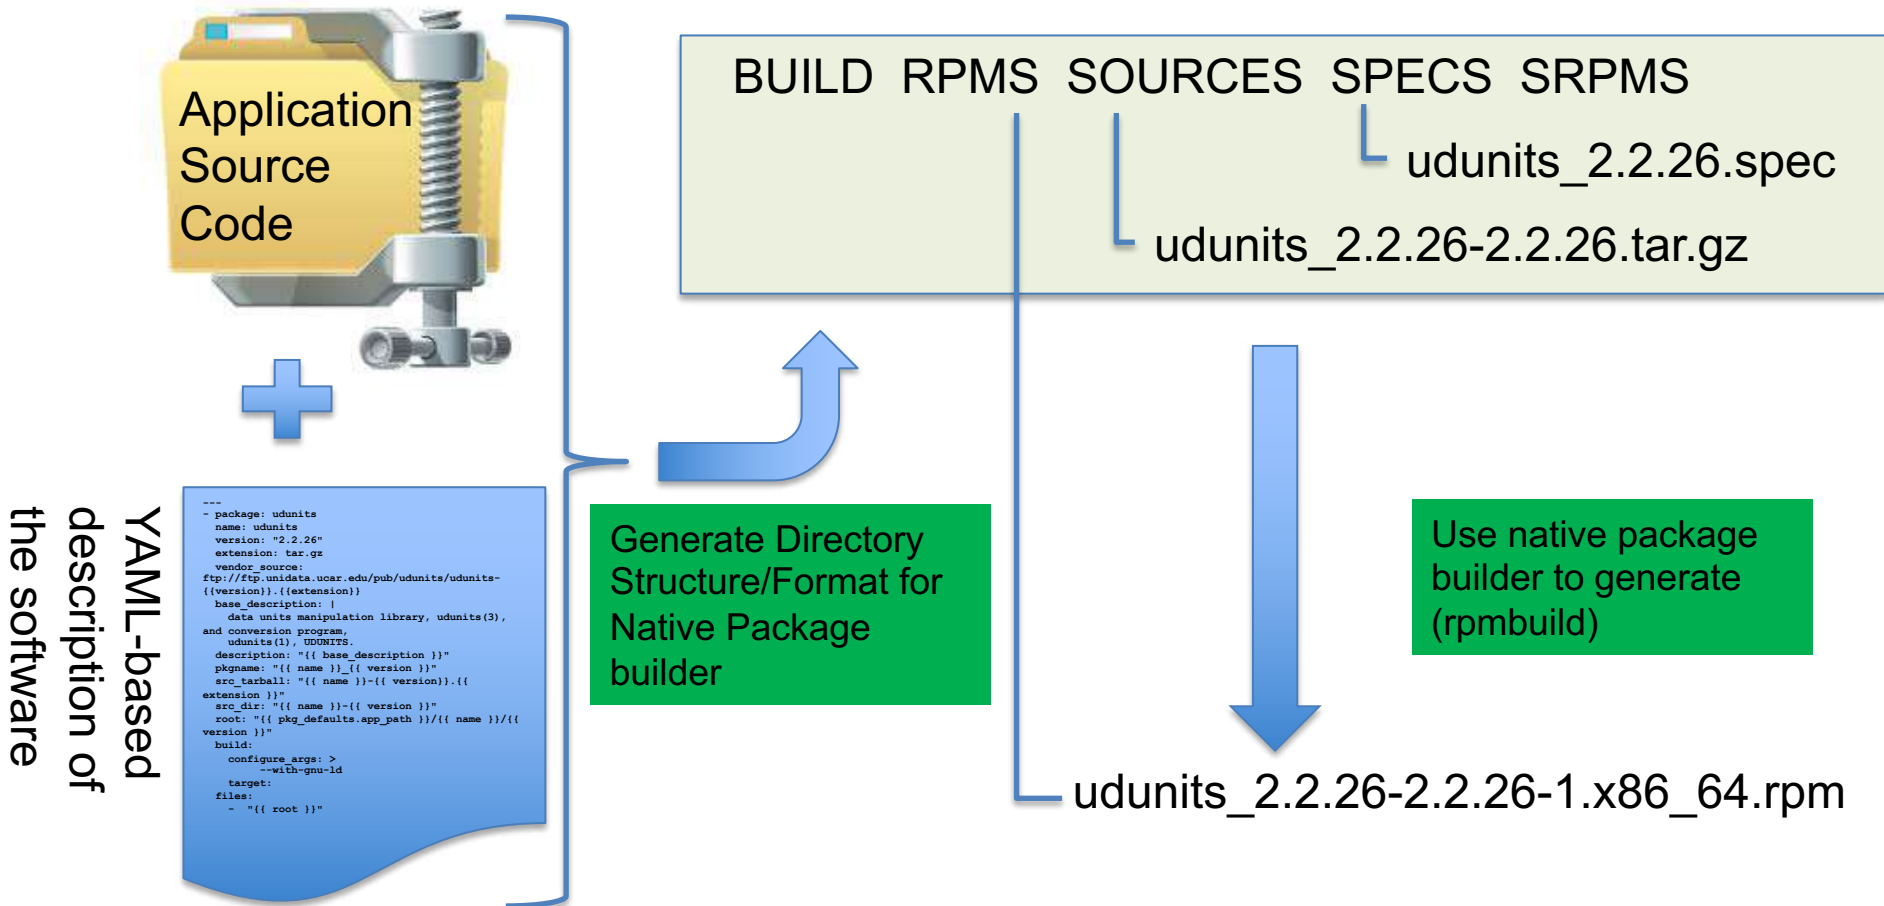
## Packaging isn't Glamorous..

- Can we build a system such that: *it is **practical** for a **a small team of people** to build an application stack that:*
  - Contains hundreds (to thousands) of discrete software components
  - Has consistent subsets of software to put into containers
  - Is repeatable/extensible
  - Works in concert with the OS-provided packaging system
- Not glamorous – but, walking this path can significantly improve stability/predictability.

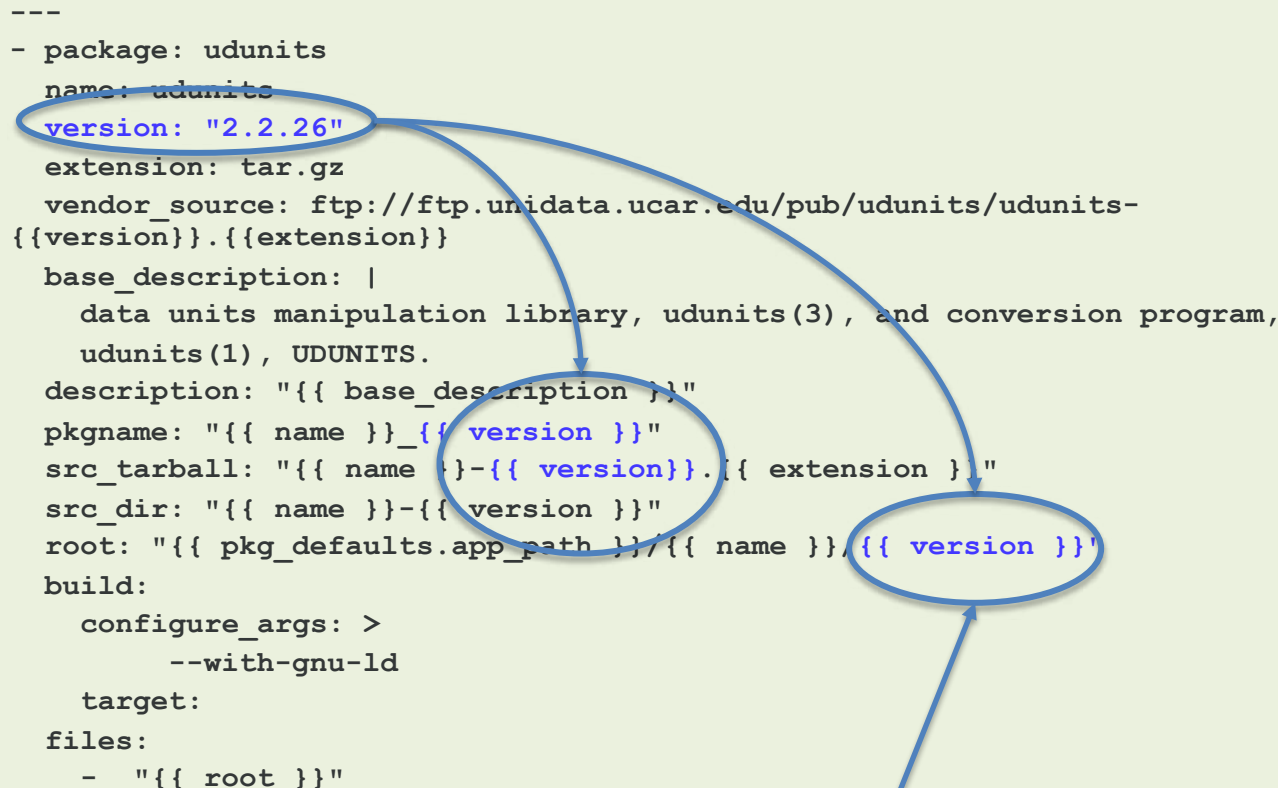
## An overriding goal

- If one can figure out how to compile package X (including its dependencies)
  - It should be straightforward to produce a binary package.
  - (often the most time-consuming step of scientific software builds is tracking down what prerequisites are needed)

YAML2RPM:  $f(\text{Source code, description file}) = \text{package}$



```
---
- package: udunits
  name: udunits
  version: "2.2.26"
  extension: tar.gz
  vendor_source: ftp://ftp.unidata.ucar.edu/pub/udunits/udunits-
  {{version}}.{{extension}}
  base_description: |
    data units manipulation library, udunits(3), and conversion program,
    udunits(1), UDUNITS.
  description: "{{ base_description }}"
  pkgname: "{{ name }}_{{ version }}"
  src_tarball: "{{ name }}-{{ version }}.{{ extension }}"
  src_dir: "{{ name }}-{{ version }}"
  root: "{{ pkg_defaults.app_path }}/{{ name }}/{{ version }}"
  build:
    configure_args: >
      --with-gnu-ld
    target:
  files:
    - "{{ root }}"
```



## Package info:

- Name
- Description
- Source

## How to build

- Any special build instructions
- Where to put it

Simple variable declaration and use  
(supports indirect definitions, too)

Note: There is nothing RPM-specific in this description. That's a feature!

```
# rpm -qip udunits_2.2.26-2.2.26-1.x86_64.rpm
```

```
Name       : udunits_2.2.26
Version    : 2.2.26
Release    : 1
Architecture: x86_64
Install Date: (not installed)
Group      : System Environment/Base
Size       : 631790
License    : University of California
Signature  : (none)
Source RPM : udunits_2.2.26-2.2.26-1.src.rpm
Build Date : Fri 06 Sep 2019 11:16:30 AM PDT
Build Host : admixdev-0.local
Relocations : (not relocatable)
Vendor     : RCIC @ UC Irvine
Summary    : udunits_2.2.26
```

**Description :**

**data units manipulation library, udunits(3), and conversion program, udunits(1), UDUNITS. . Configured with ./configure --prefix=/data/apps/udunits/2.2.26 --with-gnu-ld. Modules loaded for compilation: .**

```
# rpm -qp --provides udunits_2.2.26-2.2.26-1.x86_64.rpm
```

```
libudunits2.so.0() (64bit)
udunits_2.2.26 = 2.2.26-1
udunits_2.2.26(x86-64) = 2.2.26-1
```

- It is a real RPM.
- Can be queried for what it requires and what it provides
- Install in a package repository and use it to
  - Install on physical HW
  - Be part of a container
  - Verify in a test environment


# Drawing upon ideas from others...

- YAML format
  - Used by Kubernetes, Ansible (config mgmt.), and others
  - Already existing python-based parsers
  - Human-readable
- {{ var }}
- This is NOT part of any YAML spec, but Ansible uses this for declaring/resolving variables.
- We wrote our own custom (recursive) variable resolver
- Rocks Rolls
  - Extract the subdirectory and automated creation of RPM Spec files (directives to build packages) and required RPM structure
  - Re-uses this tooling without requiring build host to be a Rocks-defined system (all development on vanilla CentOS)

## How is this different from other efforts

- SPACK - <https://spack.io/>
    - Self-contained package format, dependency resolver. You must code in Python to generate a package
  - Anaconda (<https://www.anaconda.com>) – focused on datascience. Own packaging format/resolver/installer. Will happily install an OS-specific package on anything (resulting in broken installations)
  - Brew (<https://brew.sh/>) “missing” package manager for MacOS. Also works on Linux. You must code in Ruby.
- 
- These create **binary packages that might conflict** with OS-based packages (installers are unaware of each other)
  - YAML2RPM packages are **installed with the OS-based package manager (not with yet another installer)**
    - Kickstart, Docker Build, Singularity Recipe, Ansible, others all know how to install RPMS

## Don't let this happen to your cluster ...



Remember the Mars Climate Orbiter incident from 1999?

"RPM, Brew, SPACK, Whatever ..."





## Easy things should be straightforward ... complicated should be possible

- The simple udunits example follows a commonly-used build recipe
  1. `./configure <configure arguments>`
  2. `make`
  3. `make install`
- The build instructions in the YAML file are therefore very terse, with this pattern *implied*

```
build:  
  configure_args: >  
    --with-gnu-ld  
  target:
```

```
root: "{{ pkg_defaults.app_path }}/{{ name }}/{{ version }}"
```

# Complex builds are possible

- ✓ Custom build scripts
- ✓ Patching source code
- ✓ Locally-generated code
- ✓ Cmake-based
- ✓ Using alternative versions of support software (e.g., updated version of gcc, via environment modules)
- ✓ Unconventionally-named source archives and/or build trees
- ✓ environment-modules support
- ✓ Filtering of RPM auto-generated requires/provides to keep alternative versions (e.g. updated Python) well contained
- ✓ Reuse of build templates for generating multiple versions (e.g. gcc 6.5.0 and gcc 8.3.0)
- ✓ Simple bootstrapping when build of Software X requires Software Y to first be built and installed.
- ✓ Multiple YAML files in a single directory to build logical groups of packages

## Software Admix – Logical Package Groupings

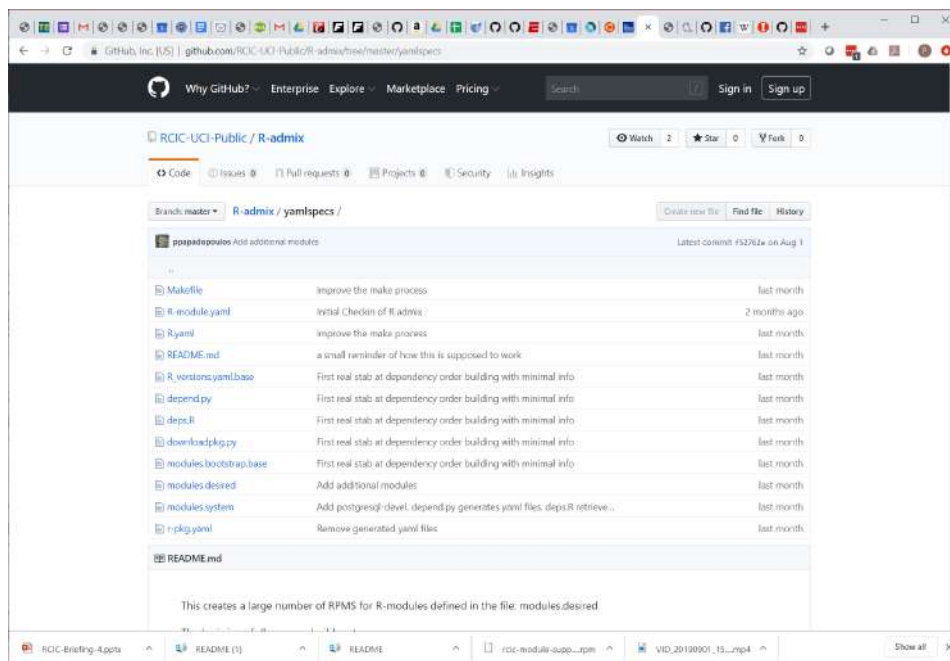
*admixture* (noun) ad·mix·ture | \ ad-'miks-  
chər - something added by mixing

Ones we are working on now:

- ❖ gcc compilers, biotools, fileformats, R, Perl, Python, chemistry, ...
- ❖ Admixes can generate 10s to 100s of related RPMs



# The R Admix– A more complicated Example



**Generates ~ 190 RPMs**

**Full Dependency resolution and recording. Can install a coherent subset of R modules. Ideal for containers.**

- 13 files in Repository
- 78 additional R modules deemed useful. Listed in a single file
- Builds R, and Environment Module (bootstrap)
- Uses CRAN (R online archive) to find all dependencies and encodes in the package
- Creates package yaml files for each module and builds in proper order

- Sample Admix: ganesha-admix
  - <https://github.com/RCIC-UCI-Public/ganesha-admix>
- YAML2RPM on GitHub
  - <https://github.com/RCIC-UCI-Public/yaml2rpm>

## README.md

## Building Ganesha (And other RPMS)

This is a yaml2rpm-based (<https://github.com/RCIC-UCI-Public/yaml2rpm>) repo for building RPMS with explicitly managing RPM spec files

0. If you have a vanilla CentOS7 system, you can prep it for building. See <https://github.com/RCIC-UCI-Public/development-RPMS>

1. make download

2. cd yamlspecs; make bootstrap; make

your ganesha rpm should be in RPMS/x86\_64

Install three  
“development”  
packages

## What's Next

- Still in development
  - Needs some documentation on how to handle “outlier” cases.
- Futures:
  - Checkout N admixes – automatically build in dependency order (need to define Admix-admix dependencies)
  - Generate most YAML files from a database (continue to [distill package specifics to bare minimum required](#))
  - Code backend generator for DPKG (Debian) format