

**Open Science Grid**

# THE OPEN SCIENCE GRID (OSG)

**And how it relates to PRAGMA**

by Igor Sfiligoi, UCSD/SDSC  
([isfiligoi@sdsc.edu](mailto:isfiligoi@sdsc.edu))

UC San Diego



# Premise – Not the Grid you were used to

- It is my understanding Grid has a negative connotation in PRAGMA
- While OSG has “Grid” in its name,  
I believe it is far from the “Grid” you experienced 15 years ago
- Please give me a chance to convince you of that

# OSG evolution

- 10 years ago, OSG provided an integrated software stack that organizations had to deploy in order for their researchers to benefit from OSG.
  - **We learned the hard way that this is too high a bar to jump across for most researchers and institutions.**
- **Today, we offer to operate services** to the groups we serve in order to minimize the threshold to entry for everybody.
  - Leveraging work performed by the likes of PRP/NRP-P/TNRP (e.g. Kubernetes)

# OSG Aspirational Goal

**A campus Research IT Organization  
should not have to learn anything "non-standard"  
in order to have their researchers benefit from OSG,  
or have their resources be available via OSG**

**Now .... Not there yet, but getting closer by the day...**

# OSG Mission

**Advancing Open Science  
through distributed High  
Throughput Computing**

# Open Science

- All of open science irrespective of discipline
- Include the maximum possible dynamic range of science, groups, and institutions
  - From **individual undergraduates** to **international collaborations** with thousands of members.
  - From **small colleges, museums, zoos,** to **national scale centers** of open science.



# What Sciences use OSG?

- Order here is from most active to least active.
  - Category is based on self-identification ... i.e. people say what they do.

High Energy Physics   Astrophysics   Medical Sciences   Biological Sciences  
Neuroscience   Particle Physics   Nuclear Physics   Training   Gravitational Physics  
Bioinformatics   Physics   Engineering   Physical Chemistry   Economics  
Biophysics   Evolutionary Sciences   Statistics   Physics and astronomy   Chemistry  
Community Grid   Biochemistry   Ecological and Environmental Sciences   Zoology  
Fluid Dynamics   Computational Condensed Matter Physics  
Biological and Critical Systems   Computer and Information Science and Engineering  
Chemical Engineering   Earth Sciences   Mathematical Sciences   Medical Imaging  
Geographic Information Science   Educational Psychology   Biomedical research  
Technology   Plant Biology   Education   Genomics   Materials Science  
Astronomy   Information, Robotics, and Intelligent Systems   Multi-Science Community

# Not all users are the same

- The **individual researchers** and small groups
- The **campus Research Support Organizations**
- **Multi-institutional Science Teams**
  - XENON, GlueX, SPT, Simons, and many many more
  - Collaborations between multiple campuses
- The 4 **“big science”** projects:
  - US-ATLAS, US-CMS, LIGO, IceCube



Used services  
change between  
different categories



# OSG Mission

**Advancing Open Science  
through distributed High  
Throughput Computing**

# distributed High Throughput Computing

(as opposed to e.g. HPC)

## DRIVING PRINCIPLES:

- Maximize total compute, do not optimize for any single application
- Allow for resources coming from anywhere in the world
- Embrace heterogeneity
- Federated policy control

# distributed High Throughput Computing

## DRIVING PRINCIPLES:

- Maximize total compute, do not optimize for any single application
- Allow for resources coming from anywhere in the world
- Embrace heterogeneity
- Federated policy control

## IMPLICATIONS:

- Expect **high network latencies**
- Expect **varying walltime** between identical runs
- Job **restarts** are the norm

(Although throughput should be high)

# Thriving in dHTC

- The path to successful dHTC involves two aspects:
  - Separate a big computing problem in many individually schedulable small problems.
  - Minimize your requirements in order to maximize the raw capacity that you can effectively use.
- The infrastructure will transparently take care of job restarts

**Ingenious  
Parallelism**



# dHTC is special

- dHTC **scales by definition perfectly.**
  - When a researcher understands how to partition their workflow into many individually schedulable compute problems they can **scale out** with ease **to seemingly arbitrary scales of computing.**
  - When we **integrate all IT resources** at Universities, National Labs, and the commercial cloud, **we arrive at a near infinite resource pool.**



# Stronger together

- For OSG VO, more than half of the compute comes from smaller sites (and this does not include any of the big sciences, that have their own pools)

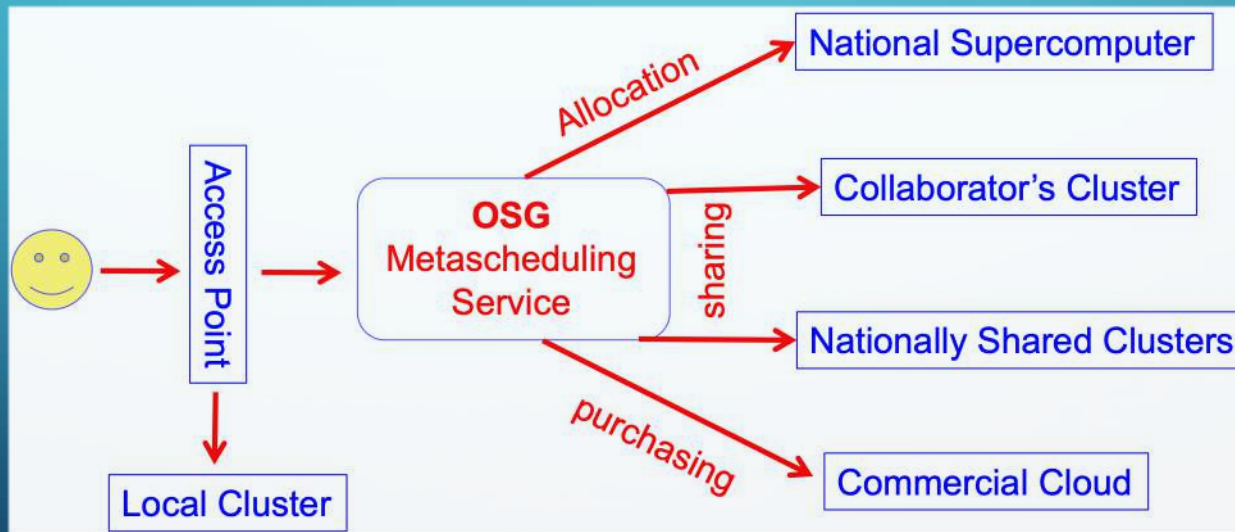


15k cores

# OSG IN DETAIL

# OSG Compute Services

- OSG provides the means for operating a compute federation
  - Mostly in the form of Metascheduling Service
  - But also operates one access point itself for the smaller players



# Federation i.e. Distributed control

- OSG works on three simple principles:
  - Resource Owners determine policy of use
    - This means that all policy of use is set locally by the clusters that join the federation.
  - Resource Consumers specify the types of resources they are willing to use.
    - How much RAM? How many cores per node? ...
  - OSG submits HTCondor batch system as (pilot) payload into all local batch systems that match requirements, creating an **overlay dHTC system**
- Jobs are submitted locally, queue centrally, and execute anywhere that matches requirements after resource becomes available.



# OSG Data Federation

- OSG operates a content delivery network (CDN)
  - Based on XRootD caching technology
- Data origins operated by Science communities
  - Data transparently published in our federated namespace.
  - Caches hide data access latencies and reduce network traffic.



| Directory                  | Working Set | Total Read |  |
|----------------------------|-------------|------------|--|
| /pnfs/fnal.gov/usr/dune    | 13.107GB    | 395.537TB  | 30k<br>1k<br>15<br>100<br>Depending on community,<br>files were read 10-30,000 times<br>during a two month period. |
| /pnfs/fnal.gov/usr/minerva | 255.266GB   | 270.994TB  |  |
| /gwdata/O1                 | 169.585GB   | 258.341TB  |  |
| /pnfs/fnal.gov/usr/des     | 193.57GB    | 120.993TB  |  |
| /user/ligo                 | 5.612TB     | 83.564TB   |  |
| /pnfs/fnal.gov/usr/nova    | 162.632GB   | 18.841TB   |  |



# OSG Software Distribution

- OSG can **host and distribute community software** on CVMFS
  - A CDN at heart, but based on Apache Squid technology
  - CVMFS provides a Read-Only POSIX interface
- OSG encourages the use of **containerization**
  - Encourages participating sites to provide Singularity
    - Which is optimized for use by unprivileged users (unlike Docker)
  - Containers can be distributed through CVMFS (for performance)

# Facilitation Services

- Proactive, personalized guidance and support for:
  - Institution-specific submit points
  - Sharing institutional resources via OSG
  - Data federation across OSG sites
  - Individual researchers using the OSG submit infrastructure
- We also offer
  - Local workshops
  - OSG-hosted education and training
  - Learning from the OSG Facilitation Community



# OSG AND PRAGMA

PRAGMA 37 - SEPT 13TH, 2019

# OSG AND PRAGMA

- OSG would be happy to facilitate any Open Science endeavor that has at least some of its members in the US
  - We are a NSF-funded project
- All OSG services are at your disposal
  - Compute federation
  - Data federation
  - Software distribution
  - Facilitation



# OSG and International Setups

- All our major users are international
  - CMS, ATLAS, LIGO, IceCube, XENON, GlueX
- But, alas, **we have been mostly collaborating with Europe**
- **PRAGMA could help us engage more Science communities in the Pacific Rim**



# Summary and Conclusion

- OSG's objective is to “Advance Open Science through distributed High Throughput Computing”
  - We are committed to help any Open Science community that has at least some members in the USA – Nice overlap with PRAGMA
- OSG can offer a wide variety of services
  - Both in terms of technology and facilitation
- We could use some help getting more involved in the Pacific Rim
- Contact us at: [help@opensciencegrid.org](mailto:help@opensciencegrid.org)



Open Science Grid