

Integrated Application and Performance Monitoring at the IoT Edge

Yutthana Boonpalit^a, Siwakorn Suwanjinda^a, Jason Haga^b, Shava Smallen^c, Prapaporn Rattanathamrong^a, Vahid Daneshmand^d, Kensworth Subratie^d
Thammasat University^a, National Institute of Advanced Industrial Science and Technology (AIST)^b,
University of California, San Diego^c, University of Florida^d
oatyootna@gmail.com^a, siwakorn4437@gmail.com^a, jh.haga@aist.go.jp^b, ssmallen@sdsc.edu^c, rattanat@gmail.com^a, vdaneshmand@acis.ufl.edu^d, kcratie@ufl.edu^d

Background

- Devices and protocols used within an IoT system may vary greatly in terms of power and connectivity requirements.
- The IoT system requires an IoT gateway to aggregate sensor data, translates between different protocols and process data before sending it to a cloud server.
- A smart home might collect sensor data and process it at a single gateway; on the other hand, a smart city management might have a gateway for each neighborhood in the entire city.
- Therefore, traditional ways of monitoring and troubleshooting used by system admins are no longer sufficient to give them the whole view of system status.

Aims & Challenges

- Build an integrated solution for IoT infrastructures and applications.
- The system must be flexible enough to work with various kinds of generally used IoT sensors and devices.
- The system must be scalable and power-efficient.

Methodology

- Lightweight agents at edge devices collect infrastructure and application related data periodically and send them to a centralized time-series data storage, which can later be queried and visualized for IoT system administrators.

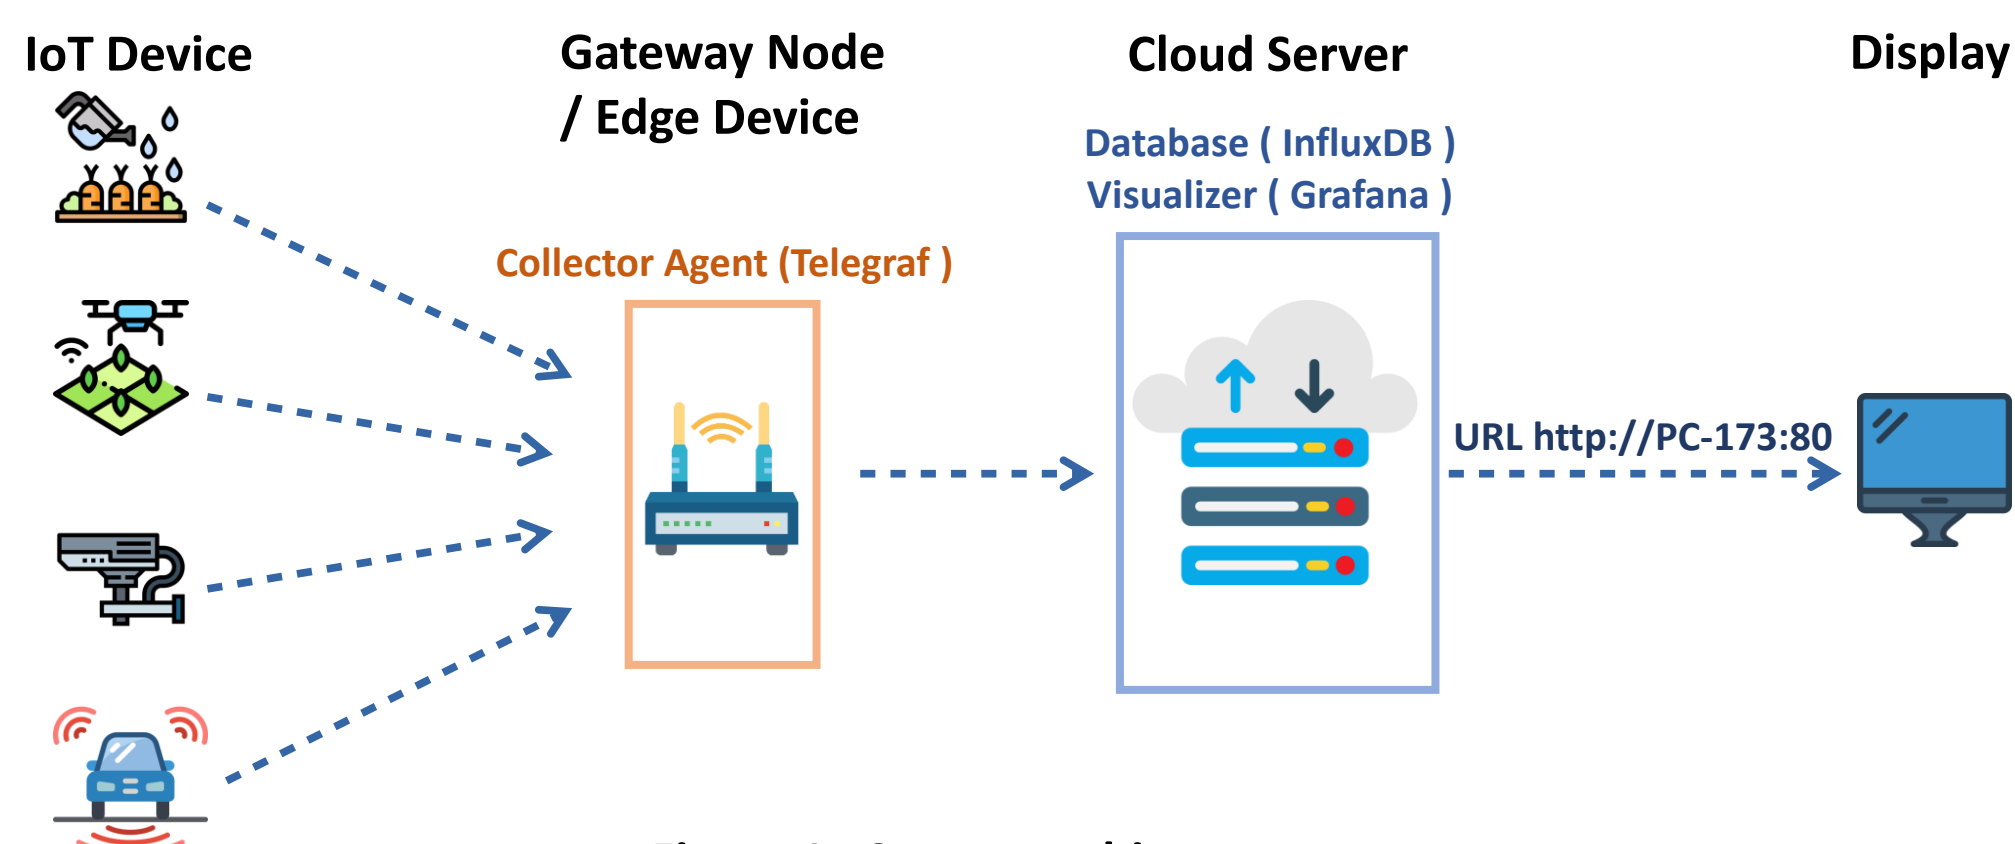


Figure 1: System architecture

- A telegraf container was deployed in a Raspberry Pi model B boards (Gateway Nodes) and InfluxDB and Grafana were deployed in a PRAGMA Cloud server.
- CPU, memory and disk I/O were collected as infrastructure-related metrics.
- We simulated application data on the IoT gateways in two ways: (1) logging data from a temperature sensor, and (2) replaying logged data from the Smart Reservoir project's gateway from its GitHub repository.




Figure 2: IoT gateway with temperature sensor

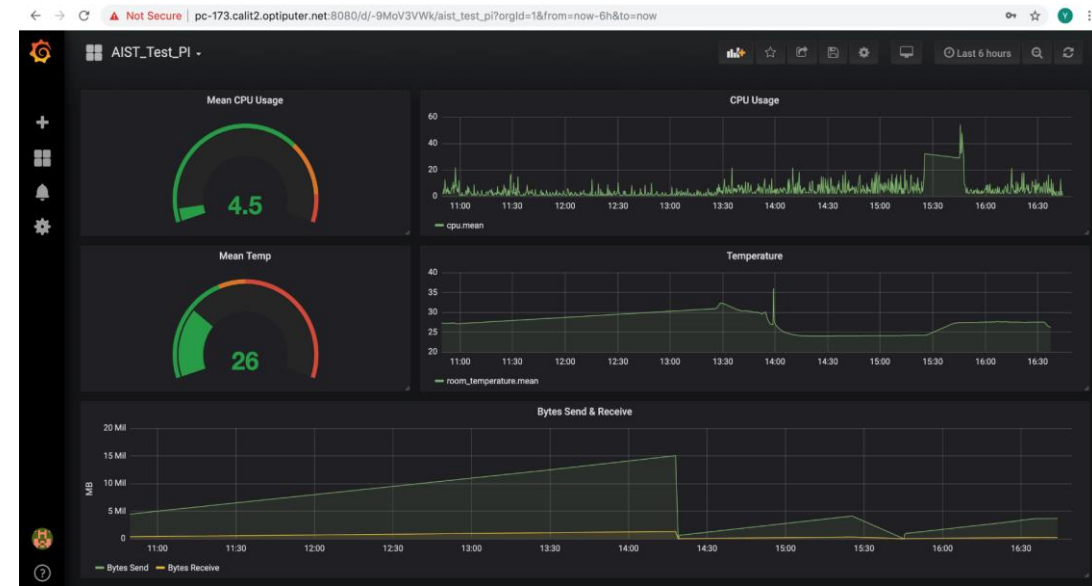


Figure 3: Grafana dashboard

- We conducted experiments to measure the overhead and power consumption of Telegraf in our setup compared to our baseline set up with no Telegraf installed. We varied the frequencies of data collection and the number of monitored metrics.

Experiment Results

Impact of Collected Metrics

- Memory usage is about 18% when Telegraf is running.
- When the number of metrics is increased from 1 to 10, Telegraf's use of CPU, memory, and disk I/O is negligible
- Telegraf's power consumption is negligible.




Figure 4: Telegraf system usage




Figure 5: Telegraf power consumption

Impact of Collection Frequencies

- Data collection rates at 100 Hz (0.01s) caused Telegraf to use about at least 247.21% increase in CPU, 5.42% increase in memory and 60.60% increase of power consumption, when compared to the lower rates.
- Telegraf's use of disk I/O is negligible.

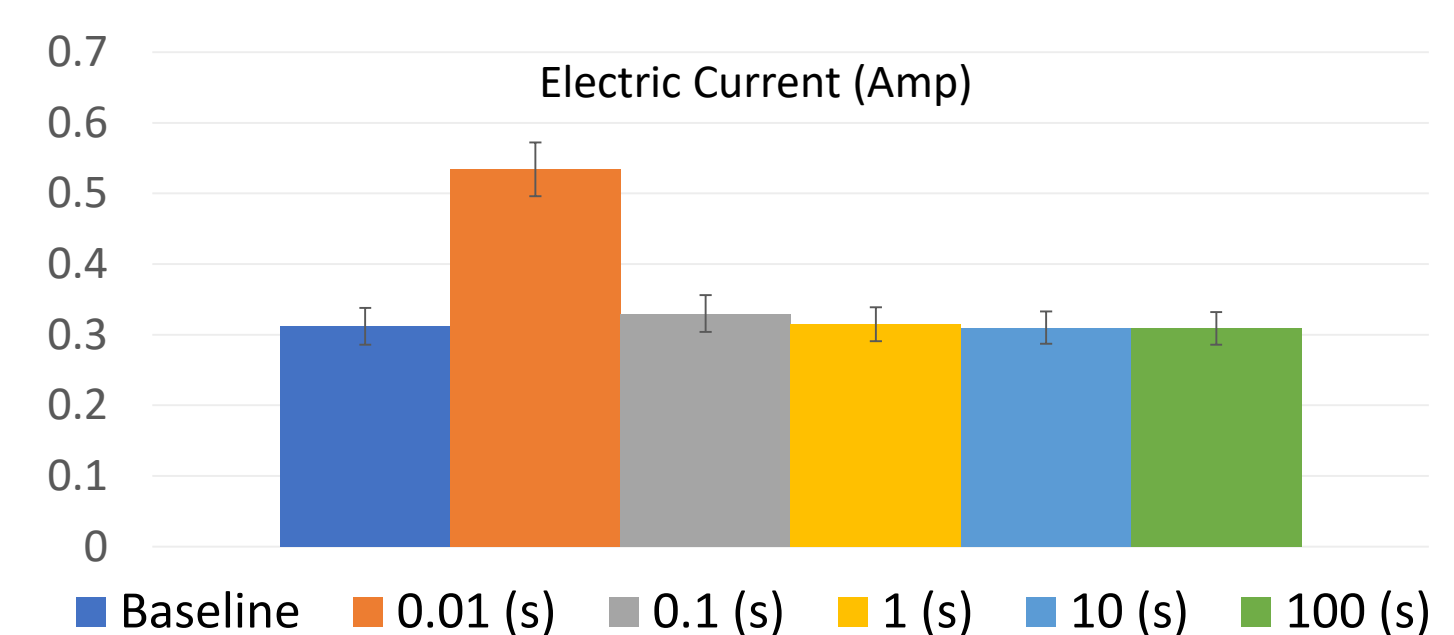


Figure 6: Telegraf power consumption




Figure 7: Telegraf system usage

Conclusions and Ongoing Work

- Telegraf's scalability and low power consumption make it ideal for deployment on Raspberry Pi.
- InfluxDB is suitable for storing time-series data.
- Grafana also scales well for data visualization.
- Data collection rates above 10 Hz can induce some system overhead, that may require Raspberry Pi deployments to use AC power.
- Investigate other approaches for Telegraf to collect application data.
- Investigate the impact of other approaches on system performance.